

# A Comprehensive Guide to Heaps and Tries

Ever wondered how your phone autocompletes your texts so fast, or how a navigation app can recalculate the "best" route on the fly? The magic often lies in specialized data structures. For my assignment, I will tackle two of these workhorses: Heaps and Tries. I'll break down the **Heap**, which acts like a VIP line or an emergency room triage system, always prioritizing the most "important" item. We'll see how this is key for algorithms that need to constantly make priority-based decisions. Next, we'll explore the **Trie**, the secret sauce behind dictionary searches and autocomplete. Think of it as organizing a massive library not just by book title, but by every single letter, allowing you to find all books starting with "Th" in a blink. This tutorial aims to be a practical journey, showing how to implement these structures from scratch and apply them to problems you'd actually encounter in an interview or on the job.

## TABLE OF CONTENTS

Chapter 1: Introduction.....	5
1.1 Welcome & Tutorial Goals.....	5
1.2 What are Data Structures and Why Do They Matter?.....	5
1.3 A Quick Look at Heaps and Tries.....	5
Heaps: The Priority Pass  .....	6
Tries: The Word Navigator  .....	7
Chapter 2: Heap Fundamentals.....	9
2.1 What is a Heap? The "Tournament Bracket" Analogy.....	9
The Two Golden Rules of Heaps.....	9
The Heap Property in Code.....	10
2.2 The Two Flavors: Min-Heaps vs. Max-Heaps.....	10
Max-Heaps: Finding the Champion  .....	10
Min-Heaps: Finding the First in Line  .....	11
Choosing Your Flavor in C++.....	12
2.3 Key Property: The Complete Binary Tree.....	13
The "Filling a Bus" Analogy  .....	13
The Power of Array Representation in C+.....	14
2.4 Why Choose a Heap? Use Cases at a Glance.....	15
Key Use Cases.....	15
1. Priority Queues.....	15
2. "Top K" Problems.....	15
3. Graph Algorithms (Shortest Path).....	16
C++ Code: Finding the "Top K" Elements.....	16
Chapter 3: Implementing a Heap.....	18
3.1 The Array-Based Representation: A Perfect Fit.....	18
Why an Array is a Perfect Fit.....	18
The C++ Representation.....	19
3.2 Core Operations (with Code & Complexity Analysis).....	20
3.2.1 peek(): A Glimpse at the Top ( $O(1)$ ).....	20
3.2.2 insert(): The "Bubble-Up" Process.....	20
3.2.3 extractMax(): The "Sink-Down" Process.....	21
Complete C++ Implementation.....	22
3.3 Building a Heap: The Smart Way ( $O(n)$ ).....	24
The "Middle Management" Analogy.....	24
The Algorithm.....	25
C++ Implementation.....	25
Chapter 4: Advanced Heap Concepts & Applications.....	28
4.1 Heap Sort: In-Place Sorting with $O(1)$ Space.....	28

The "King of the Hill" Analogy 	28
The Heap Sort Algorithm.....	28
C++ Implementation.....	29
<b>4.2 Finding the K-th Largest/Smallest Element.....</b>	<b>31</b>
The "VIP Lounge Bouncer" Analogy 	31
The Algorithm (for K-th Largest).....	31
C++ Implementation.....	32
<b>4.3 Application Spotlight: Dijkstra's &amp; Prim's Algorithms.....</b>	<b>33</b>
Dijkstra's Algorithm: Finding the Shortest Path.....	33
Prim's Algorithm: Building a Minimum Spanning Tree.....	33
C++ Code Spotlight: Dijkstra's Algorithm.....	34
<b>4.4 Advanced Variants (A Brief Overview).....</b>	<b>36</b>
4.4.1 d-ary Heaps.....	36
4.4.2 Binomial & Fibonacci Heaps.....	37
<b>Chapter 5: Heap Problem-Solving Workshop.....</b>	<b>40</b>
5.1 Problem 1: Find Median from a Data Stream.....	40
The Problem: Find Median from a Data Stream.....	40
The "Two Halves of a Bridge" Analogy 	40
The Algorithm.....	40
C++ Implementation.....	41
5.2 Problem 2: Merge K Sorted Lists.....	43
The Problem: Merge K Sorted Lists.....	43
The "Race Finish Line" Analogy 	43
The Algorithm.....	43
C++ Implementation.....	44
<b>Chapter 6: Trie Fundamentals.....</b>	<b>47</b>
6.1 What is a Trie? The "Autocomplete" Analogy.....	47
The Structure of a Trie.....	47
The C++ Building Block: The TrieNode.....	47
6.2 The Trie Node: Building Blocks of Words.....	49
The "Signpost at a Crossroads" Analogy 	49
The Anatomy of a TrieNode.....	49
The C++ TrieNode Struct.....	51
6.3 Why Choose a Trie over a Hash Table?.....	52
The "Encyclopedia vs. Filing Cabinet" Analogy.....	52
Key Differences.....	52
The C++ Code Test: A Prefix Search.....	53
<b>Chapter 7: Implementing a Trie.....</b>	<b>56</b>
7.1 The Node Structure in Code.....	56
The "Lego Brick" Analogy 	56
The TrieNode in C++.....	56

The C++ Code.....	57
7.2 Core Operations (with Code & Complexity Analysis).....	58
7.2.1 insert(): Adding a Word to the Lexicon.....	58
7.2.2 search(): Finding an Exact Match.....	59
7.2.3 startsWith(): Checking for a Prefix.....	59
7.3 The Tricky Task of delete().....	62
Chapter 8: Advanced Trie Concepts & Applications.....	66
8.1 Saving Space: Compressed Tries (Radix Trees).....	66
The "Express Train" Analogy  .....	66
The C++ Node Structure.....	67
8.2 Application Spotlight 1: Spell Checkers & Dictionaries.....	68
The "Path in a Dictionary" Analogy  .....	68
How Tries Power These Features.....	68
C++ Code: Implementing a Suggestion Feature.....	70
8.3 Application Spotlight 2: IP Routing Tables.....	72
The Router's Dilemma.....	72
The "Postal ZIP Code" Analogy  .....	72
C++ Code: Simulating Longest Prefix Match.....	73
8.4 Advanced Variant: Ternary Search Trees (A Brief Overview).....	75
The "20 Questions" Analogy  .....	75
The Structure of a TST Node.....	76
C++ Code: The TSTNode.....	77
Chapter 9: Trie Problem-Solving Workshop.....	79
9.1 Problem 1: Add and Search Word.....	79
9.2 Problem 2: Word Search II.....	81
Chapter 10: Conclusion.....	86
10.1 Head-to-Head: Heap vs. Trie Comparison Table.....	86
10.2 Key Takeaways & Final Thoughts.....	86
10.3 Further Reading & Resources.....	87

# Chapter 1: Introduction

## 1.1 Welcome & Tutorial Goals

Welcome! If you've ever wondered how your favorite app can instantly suggest a contact as you type, or how a video game always knows which enemy to attack first, you've come to the right place. You're about to dive into two of the cleverest tools in a programmer's toolbox: **Heaps** and **Tries**.

Think of data structures like specialized tools. You wouldn't use a hammer to turn a screw. In the same way, we use different data structures to solve different problems efficiently.

In this tutorial, we'll explore:

- **Heaps:** Imagine an emergency room's triage system. It doesn't matter who arrived first; the most critical patient is always treated next. A Heap works just like that, acting as a "priority queue" that always keeps the most important item (whether highest or lowest value) ready at the top for instant access.
- **Tries:** Picture your phone's contact list. The moment you type 'J', it suggests 'Jake', 'Jane', and 'John'. A Trie is the magic behind this, a tree built for lightning-fast text searches and autocomplete suggestions based on prefixes.

Throughout this guide, we'll use simple diagrams to visualize how these structures are built and how data flows through them.

We'll provide clean, commented code snippets along the way to connect the concepts to a practical implementation.

## 1.2 What are Data Structures and Why Do They Matter?

A **data structure** is a way of organizing data to perform specific tasks efficiently. Think of a library: you can arrange books on a single shelf (like an **array**) for fast lookups, or link them with notes (like a **linked list**) for easy insertions. Neither is better; they're just good at different things. This choice between speed and flexibility is the most important concept in data structures.

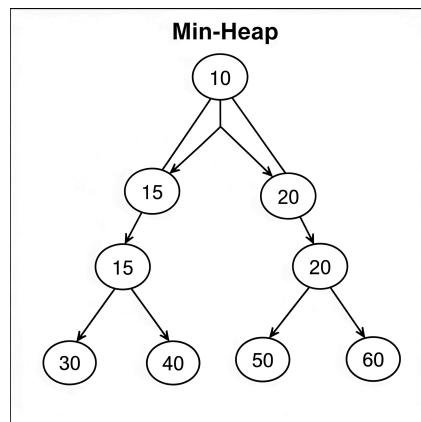
## 1.3 A Quick Look at Heaps and Tries

Let's briefly introduce our two main topics. This is a high-level preview—we'll dive into the "how" and "why" in the upcoming chapters.

## Heaps: The Priority Pass 🏆

A **Heap** is all about one thing: **priority**. It's a data structure that ensures you can find the element with the highest (or lowest) value in an instant.

- **Real-World Analogy:** Think of the "Now Serving" board at a deli or a VIP line at a concert. The system doesn't care about the whole line; it only cares about who is next. A Heap provides this "highest priority" item immediately.



- **Usage in C++:** C++ has a built-in Heap implementation called `std::priority_queue`. Here's how you could use it to find the top 3 scores from a list. Notice how it naturally keeps the largest items at the top.

```
#include <iostream>
#include <vector>
#include <queue> // For priority_queue

int main() {
    std::vector<int> scores = {88, 95, 72, 100, 99, 85};
    std::priority_queue<int> top_scores;

    // Insert all scores into the heap
    for (int score : scores) {
        top_scores.push(score);
    }

    // The highest scores are now at the top
    std::cout << "Top 3 scores:\n";
    for (int i = 0; i < 3; ++i) {
        std::cout << i + 1 << ". " << top_scores.top() << std::endl;
    }
}
```

```

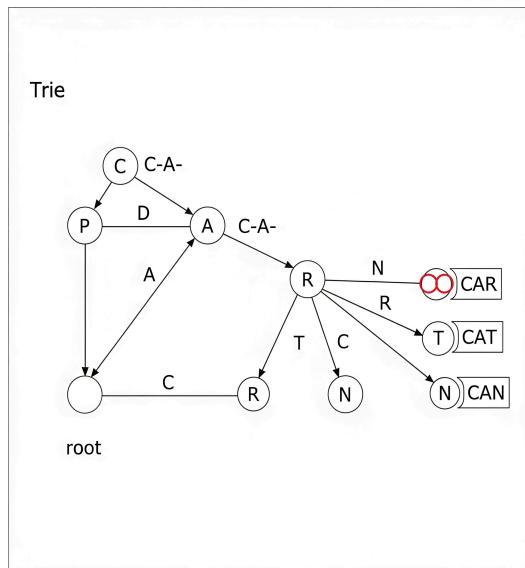
        top_scores.pop(); // Remove the top element
    }
    return 0;
}

```

## Tries: The Word Navigator

A **Trie** (pronounced "try"), or **Prefix Tree**, is a specialized structure for handling strings. It's designed for lightning-fast prefix-based lookups.

- **Real-World Analogy:** A Trie is the engine behind **autocomplete**. When you type "comp" into a search bar, it instantly suggests "computer," "compiler," and "component." A Trie stores strings in a way that makes finding all words with a common prefix incredibly efficient.



- **Conceptual C++ Code:** We'll build our own Trie class later, but its interface would look something like this. The operations are intuitive and directly reflect its purpose.

```

// This is just a conceptual interface. We'll implement it later!
class Trie {
public:
    // Adds a word to the Trie
    void insert(const std::string& word);

    // Returns true if the exact word exists
    bool search(const std::string& word);

    // Returns true if there is any word that starts with the given prefix
}

```

```
    bool startsWith(const std::string& prefix);  
};
```

This quick tour gives you a sense of their distinct roles. Heaps are for priority; Tries are for prefixes.

## Part I: Heaps

# Chapter 2: Heap Fundamentals

## 2.1 What is a Heap? The "Tournament Bracket" Analogy

A **Heap** is a specialized tree-based data structure that efficiently manages and retrieves items based on their priority.

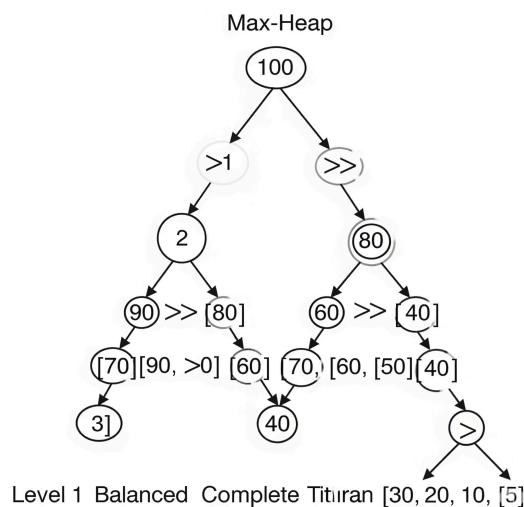
The easiest way to understand it is through a **tournament bracket**. In any tournament, after each match, a winner advances. This continues until a single champion remains at the top. The key rule is simple: the winner of any given match is always "better" (stronger, higher score, etc.) than the participants below them in that branch of the bracket.

A **Max-Heap** works exactly like this. The value at any node is always greater than or equal to the values of its children. This guarantees that the largest element—the overall champion—is always at the root of the tree, ready to be picked in an instant.

### The Two Golden Rules of Heaps

For a tree to be officially called a Heap, it must follow two strict rules:

1. **It must be a Complete Binary Tree.** This means the tree is filled on all levels, with the possible exception of the last level, which must be filled from **left to right**. There are no "gaps" in the tree.
2. **It must satisfy the Heap Property.** For a **Max-Heap**, every parent's value must be greater than or equal to its children's values. For a **Min-Heap**, every parent's value must be less than or equal to its children's values.



## The Heap Property in Code

We don't need to write a full implementation just yet, but the core logic of a Max-Heap boils down to a simple check. If we have the indices of a parent and its children in an array (which we'll cover in the next section), the check is straightforward.

Conceptually, for any given parent node in the heap, the following must be true:

```
// This is a conceptual check, not a full program.
// 'heapArray' is the array storing our heap.
// 'parentIndex' is the index of the node we are checking.

bool isHeapPropertySatisfied(int parentIndex) {
    int leftChildIndex = 2 * parentIndex + 1;
    int rightChildIndex = 2 * parentIndex + 2;

    // Check if left child exists and violates the property
    if (leftChildIndex < heapArray.size() && heapArray[parentIndex] <
        heapArray[leftChildIndex]) {
        return false; // Parent is smaller than its left child!
    }

    // Check if right child exists and violates the property
    if (rightChildIndex < heapArray.size() && heapArray[parentIndex] <
        heapArray[rightChildIndex]) {
        return false; // Parent is smaller than its right child!
    }

    // If we passed both checks, the property holds for this node.
    return true;
}
```

This simple, powerful rule is the foundation of everything a Heap can do. In the next sections, we'll see how to enforce this rule when we add or remove elements.

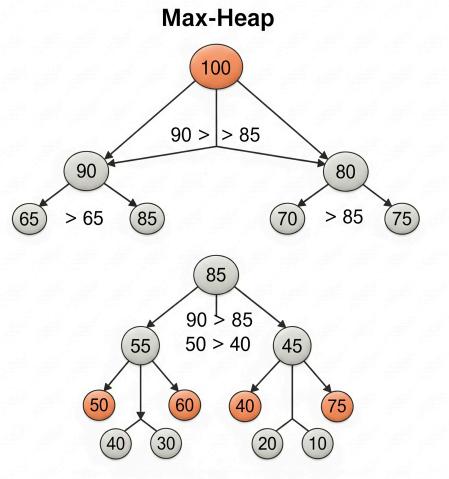
## 2.2 The Two Flavors: Min-Heaps vs. Max-Heaps

A heap's "flavor" just depends on what it considers high priority: the largest value or the smallest value. This gives us two types: Max-Heaps and Min-Heaps.

Max-Heaps: Finding the Champion 🥇

A **Max-Heap** is structured to always keep the **largest** element at the top.

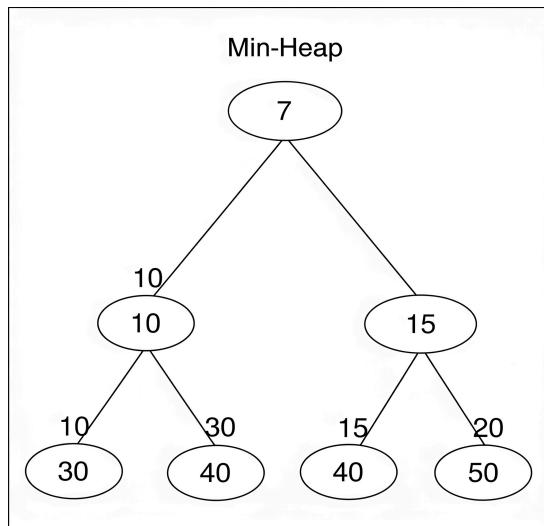
- **Real-World Analogy:** Think of a video game's AI trying to target the enemy with the most health, or an auction system that always displays the current highest bid. The system is built to instantly identify the maximum value.
- **The Rule:** The value of a parent node must be **greater than or equal to** the values of its children.



## Min-Heaps: Finding the First in Line

A **Min-Heap** is the exact opposite. It's structured to always keep the **smallest** element at the top.

- **Real-World Analogy:** Imagine an operating system's scheduler that needs to run the task requiring the least memory first, or a system finding the lowest-priced flight from a list of search results. The system is built to instantly identify the minimum value.



- **The Rule:** The value of a parent node must be **less than or equal to** the values of its children.

## Choosing Your Flavor in C++

In C++, you can easily create both using the `std::priority_queue`. By default, it's a Max-Heap. To get a Min-Heap, you just provide a different comparison function.

```
#include <iostream>
#include <vector>
#include <queue> // For priority_queue

int main() {
    std::vector<int> numbers = {30, 100, 25, 90};

    // 1. Max-Heap (Default behavior)
    // Keeps the Largest element at the top.
    std::priority_queue<int> maxHeap;
    for (int n : numbers) {
        maxHeap.push(n);
    }
    std::cout << "Max-Heap top element: " << maxHeap.top() << std::endl; // Output: 100

    // 2. Min-Heap
    // To get a Min-Heap, we provide a container and a comparator.
    // std::greater<int> tells the queue to keep the "greater" elements at the bottom.
    std::priority_queue<int, std::vector<int>, std::greater<int>> minHeap;
    for (int n : numbers) {
        minHeap.push(n);
    }
    std::cout << "Min-Heap top element: " << minHeap.top() << std::endl; // Output: 25

    return 0;
}
```

The choice is simple: if you need fast access to the biggest item, use a Max-Heap. If you need the smallest, use a Min-Heap.

## 2.3 Key Property: The Complete Binary Tree

A heap has a strict structural requirement: it must be a **complete binary tree**. This property is the secret to a heap's efficiency and its elegant implementation.

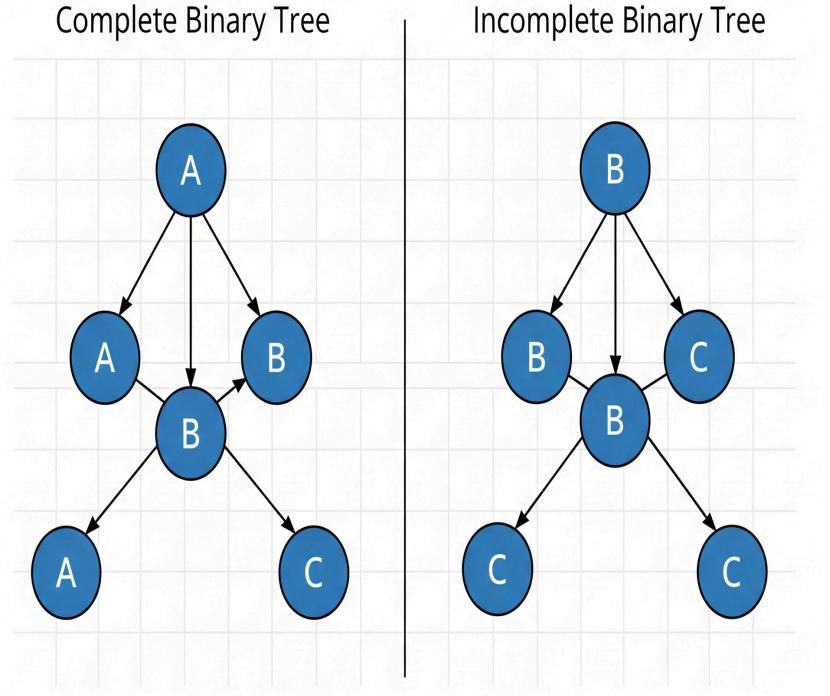
A tree is considered "complete" if all its levels are fully filled, except possibly for the last level. The last level must be filled from **left to right**, with no gaps.

### The "Filling a Bus" Analogy

Think about how people fill seats on a bus. They fill the first row completely, then the second, and so on. No one starts a new row until the previous one is full. On the last occupied row, people take the seats starting from the left, not leaving random empty seats in the middle. A complete binary tree fills its nodes in exactly the same way.

This structure ensures two things:

1. **Minimum Height:** The tree is as compact as possible, which keeps its height at  $O(\log n)$ . This is vital for fast insertions and deletions.
2. **No Wasted Space:** It allows us to perfectly map the tree to an array, which is a simple and highly efficient way to store it in memory.



## The Power of Array Representation in C++

Because a heap is a complete binary tree, we can confidently store it in an array and use simple formulas to navigate between parents and children without needing pointers. This is a huge advantage.

Given a node at index  $i$  in an array:

- Its parent is at index  $(i - 1) / 2$
- Its left child is at index  $2 * i + 1$
- Its right child is at index  $2 * i + 2$

This C++ code demonstrates how we can map a heap's structure to an array and use these formulas to traverse it.

```
#include <iostream>
#include <vector>
#include <string>

// Helper function to print a node and its family
void print_node_info(const std::vector<int>& heap, int i) {
    if (i >= heap.size()) {
        std::cout << "Index " << i << " is out of bounds.\n";
        return;
    }

    std::cout << "Node at index " << i << ": Value = " << heap[i] << std::endl;

    // Calculate parent index and value
    if (i > 0) {
        int p_idx = (i - 1) / 2;
        std::cout << " - Parent at index " << p_idx << ": Value = " <<
heap[p_idx] << std::endl;
    } else {
        std::cout << " - This is the root node.\n";
    }

    // Calculate left child index and value
    int l_idx = 2 * i + 1;
    if (l_idx < heap.size()) {
        std::cout << " - Left child at index " << l_idx << ": Value = " <<
heap[l_idx] << std::endl;
    }
}
```

```

// Calculate right child index and value
int r_idx = 2 * i + 2;
if (r_idx < heap.size()) {
    std::cout << " - Right child at index " << r_idx << ": Value = " <<
heap[r_idx] << std::endl;
}
std::cout << "-----\n";
}

int main() {
// This array represents a complete binary tree (and also a Max-Heap)
std::vector<int> heap_array = {100, 90, 80, 70, 60, 50, 40};

std::cout << "Analyzing Node at index 1 (Value 90):\n";
print_node_info(heap_array, 1);

std::cout << "Analyzing Node at index 3 (Value 70):\n";
print_node_info(heap_array, 3);
}

```

This predictable, gap-free structure is the backbone of a heap's performance.

## 2.4 Why Choose a Heap? Use Cases at a Glance

You should choose a heap when your application needs to **repeatedly and efficiently find the highest or lowest priority item** from a collection that changes over time. If you only need the min/max once, a simple loop is fine. But if you need it constantly, a heap is the champion.

### Key Use Cases

#### 1. Priority Queues

This is the most direct application of a heap. A priority queue is a data structure where each element has a "priority," and elements with higher priority are served before elements with lower priority.

- **Real-World Analogy** : An Operating System's process scheduler. Your computer runs dozens of tasks at once (e.g., your mouse, a web browser, a background update). The OS uses a priority queue (implemented as a heap) to decide which process gets CPU time next. An urgent user input will have a higher priority than a routine background check.

#### 2. "Top K" Problems

Heaps excel at finding the "Top K" items (e.g., top 10 scores, 5 most frequent words) from a large dataset without storing the whole dataset in sorted order.

- **Real-World Analogy** 🎮: Imagine maintaining a live leaderboard for the top 10 scores in an online game with millions of players. You can use a **Min-Heap of size 10**. When a new score comes in, you compare it to the smallest score on your leaderboard (the heap's root). If the new score is higher, you remove the smallest and insert the new one. This is far more efficient than constantly re-sorting all million scores.
- **Diagram:** The diagram below shows this "Top K" logic. A stream of numbers flows towards a Min-Heap of size 3. The heap only keeps the three largest elements seen so far by ejecting its smallest element whenever a larger number arrives.

### 3. Graph Algorithms (Shortest Path)

Heaps are crucial for optimizing famous graph algorithms like Dijkstra's, which finds the shortest path between two points.

- **Real-World Analogy** 🌎: A **GPS navigation app** finding the fastest route. The app explores possible paths and needs to constantly answer the question: "Of all the unvisited locations, which one is currently the closest to my start point?" A Min-Heap is perfect for managing these locations and efficiently retrieving the one with the minimum distance.

### C++ Code: Finding the "Top K" Elements

This code snippet solves the classic "Kth Largest Element" problem. It uses a Min-Heap to keep track of the top `k` elements seen so far in a much more efficient way than sorting the entire array.

```
#include <iostream>
#include <vector>
#include <queue>

// This function finds the Kth Largest element in an array.
// As a result, the heap will contain the top K Largest elements.
int findKthLargest(const std::vector<int>& nums, int k) {
    // Create a Min-Heap. We use std::greater to make the priority_queue
    // keep the smallest element at the top.
    std::priority_queue<int, std::vector<int>, std::greater<int>> minHeap;

    for (int num : nums) {
        if (minHeap.size() < k) {
            minHeap.push(num);
        } else if (num > minHeap.top()) {
            // If the current number is larger than the smallest in our heap,
            // replace the smallest with the current number.
        }
    }
}
```

```
        minHeap.pop();
        minHeap.push(num);
    }
}

// The root of the heap is the Kth Largest element.
return minHeap.top();
}

int main() {
    std::vector<int> scores = {3, 2, 1, 5, 6, 4};
    int k = 2;
    int kth_largest = findKthLargest(scores, k);

    // After the function runs, the heap contains {5, 6}.
    // The top element of the min-heap is 5, which is the 2nd Largest.
    std::cout << "The " << k << "nd largest element is: " << kth_largest <<
    std::endl; // Output: 5

    return 0;
}
```

# Chapter 3: Implementing a Heap

### 3.1 The Array-Based Representation: A Perfect Fit

A heap's strict **complete binary tree** structure makes a simple array the perfect way to store it. This approach avoids the memory overhead of pointers and is highly efficient.

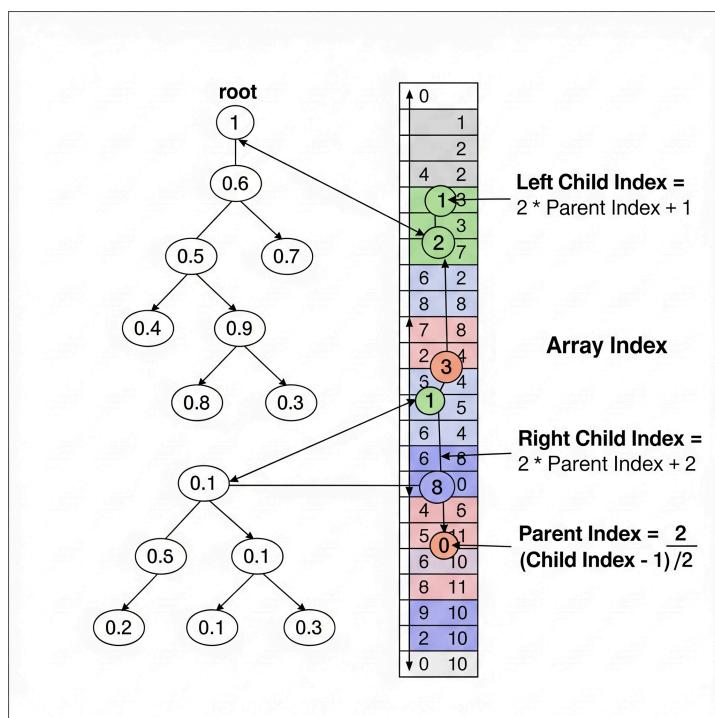
# Why an Array is a Perfect Fit

Because a heap has no gaps, we can store its nodes level by level in an array. The element at index **0** is the root. Its children are at indices **1** and **2**. Their children are at **3**, **4**, **5**, and **6**, and so on. This dense packing means:

- **No Wasted Space:** Every slot in the array corresponds to a node.
  - **No Pointers Needed:** We don't need `Node` objects with `left` and `right` pointers.
  - **Simple Math for Navigation:** We can find any node's parent or children with simple arithmetic.

The navigation formulas for a node at index  $i$  are:

- Parent:  $(i - 1) / 2$
  - Left Child:  $2 * i + 1$
  - Right Child:  $2 * i + 2$



## The C++ Representation

Our C++ heap implementation will be built around a `std::vector`, which acts as a dynamic array. The core logic is just these simple navigation functions.

```
#include <iostream>
#include <vector>

// This class will be the foundation of our heap.
// For now, it just holds the data and has navigation methods.
class Heap {
public:
    std::vector<int> nodes;

    // --- Navigation ---
    int parent_idx(int i) {
        if (i == 0) return -1; // Root has no parent
        return (i - 1) / 2;
    }

    int left_child_idx(int i) {
        return 2 * i + 1;
    }

    int right_child_idx(int i) {
        return 2 * i + 2;
    }
};

int main() {
    Heap myHeap;
    myHeap.nodes = {100, 95, 80, 70, 60, 50}; // A valid Max-Heap

    int idx = 2; // Node with value 80
    std::cout << "Node at index " << idx << " has value " << myHeap.nodes[idx]
    << std::endl;

    int parentIdx = myHeap.parent_idx(idx);
    std::cout << " -> Its parent is at index " << parentIdx
        << " with value " << myHeap.nodes[parentIdx] << std::endl;

    int leftChildIdx = myHeap.left_child_idx(idx);
    std::cout << " -> Its left child is at index " << leftChildIdx
        << " with value " << myHeap.nodes[leftChildIdx] << std::endl;
```

```
    return 0;  
}
```

This elegant mapping is the foundation upon which we'll build all other heap operations.

## 3.2 Core Operations (with Code & Complexity Analysis)

Now that we have our heap stored in an array, let's implement the three essential operations that make it so powerful. We'll use a **Max-Heap** for our examples.

### 3.2.1 `peek()`: A Glimpse at the Top (O(1))

This is the simplest operation. It lets you see the highest-priority item without removing it.

- **Analogy:** `peek()` is like looking at the "Now Serving" number at a deli counter. You get the information instantly without affecting the queue.
- **Implementation:** Since the highest-priority element is always at the root in our array-based representation, we just return the element at index `0`.
- **Complexity:** This is an O(1) operation because it's a direct array lookup, making it incredibly fast.

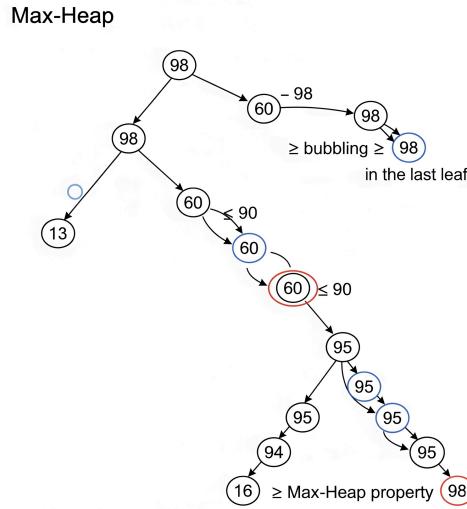
### 3.2.2 `insert()`: The "Bubble-Up" Process

When we add a new element, we must maintain both the **complete tree** property and the **heap** property.

- **Analogy:** A talented new player joins a seeded tournament. To maintain fairness, they start at the very bottom. They then challenge the player above them. If they win (i.e., their value is higher), they swap places. They keep "bubbling up" the bracket until they meet a player they can't beat, finding their correct rank.

The process is two steps:

1. **Add to End:** Place the new element at the end of the array. This maintains the complete tree structure.
2. **Bubble Up (Heapify-Up):** The new element might be larger than its parent, violating the heap property. To fix this, we repeatedly compare it to its parent and swap if it's larger, continuing until it's in the correct spot or it becomes the new root.



- **Complexity:** The height of a complete binary tree is  $\log_2 n$ . In the worst case, a new element might have to bubble all the way up to the root. Therefore, the complexity is  $O(\log n)$ .

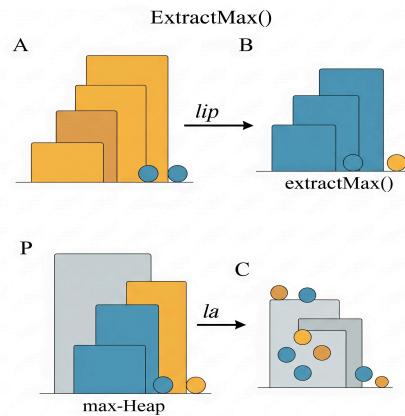
### 3.2.3 extractMax(): The "Sink-Down" Process

Removing the highest-priority element also requires a careful two-step process to maintain the heap's integrity.

- **Analogy:** The tournament champion retires. To fill the top spot, we temporarily promote the very last-ranked player to champion. This creates a mismatch. This new "champion" then has to defend their spot against their stronger children, "sinking down" the bracket by losing matches and swapping places until they reach a level where they are stronger than the players below them.

The process is:

1. **Replace Root:** Take the value from the last element in the array and place it at the root (index 0). Remove the last element. This keeps the tree complete.
2. **Sink Down (Heapify-Down):** The new root is likely smaller than its children, violating the heap property. We fix this by swapping it with its **largest** child. We repeat this



"sinking" process until the element is no longer smaller than its children, restoring the heap property.

- **Complexity:** Just like insertion, the element might have to sink all the way from the root to a leaf. The number of swaps is limited by the tree's height, so the complexity is  $O(\log n)$ .

## Complete C++ Implementation

Here is a full implementation of a `MaxHeap` class in C++ with all the core operations.

```
#include <iostream>
#include <vector>
#include <stdexcept>
#include <algorithm> // For std::swap

class MaxHeap {
private:
    std::vector<int> nodes;

    // Helper to move an element up to its correct position
    void heapifyUp(int index) {
        if (index == 0) return;
        int parentIndex = (index - 1) / 2;
        if (nodes[index] > nodes[parentIndex]) {
            std::swap(nodes[index], nodes[parentIndex]);
            heapifyUp(parentIndex); // Recurse
        }
    }

    // Helper to move an element down to its correct position
    void heapifyDown(int index) {
        int leftChildIndex = 2 * index + 1;
        int rightChildIndex = 2 * index + 2;
        int largestIndex = index;

        // Find which is largest: the node, its left child, or its right child
        if (leftChildIndex < nodes.size() && nodes[leftChildIndex] >
nodes[largestIndex]) {
            largestIndex = leftChildIndex;
        }
        if (rightChildIndex < nodes.size() && nodes[rightChildIndex] >
nodes[largestIndex]) {
            largestIndex = rightChildIndex;
        }
    }
}
```

```

    // If the Largest is not the current node, swap and continue sinking
    down
    if (largestIndex != index) {
        std::swap(nodes[index], nodes[largestIndex]);
        heapifyDown(largestIndex); // Recurse
    }
}

public:
// Check if the heap is empty
bool isEmpty() const {
    return nodes.empty();
}

// Get the max element without removing it
int peek() const {
    if (isEmpty()) {
        throw std::out_of_range("Heap is empty");
    }
    return nodes[0];
}

// Add a new element
void insert(int value) {
    nodes.push_back(value);
    heapifyUp(nodes.size() - 1);
}

// Remove and return the max element
int extractMax() {
    if (isEmpty()) {
        throw std::out_of_range("Heap is empty");
    }
    int max_value = nodes[0];
    nodes[0] = nodes.back(); // Move Last element to root
    nodes.pop_back(); // Erase Last element

    if (!isEmpty()) {
        heapifyDown(0); // Restore heap property
    }
    return max_value;
}
};

int main() {

```

```

MaxHeap heap;
heap.insert(50);
heap.insert(95);
heap.insert(80);
heap.insert(100);
heap.insert(60);

// Expected peek: 100
std::cout << "Highest priority item is: " << heap.peek() << std::endl;

std::cout << "Extracting elements:" << std::endl;
while (!heap.isEmpty()) {
    std::cout << heap.extractMax() << " "; // Expected: 100 95 80 60 50
}
std::cout << std::endl;

return 0;
}

```

### 3.3 Building a Heap: The Smart Way ( $O(n)$ )

While you can build a heap by inserting elements one by one, that takes  $O(n \log n)$  time. If you have all the elements upfront in an array, there's a much faster, more clever way to do it in just  $O(n)$  time.

#### The "Middle Management" Analogy

Imagine you need to structure a new company department.

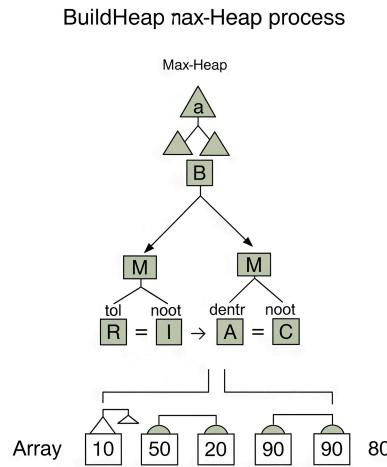
- **The Naive Way ( $O(n \log n)$ ):** Hire every employee one by one. Each new hire might cause a chain reaction of promotions and changes all the way to the top. This is slow.
- **The Smart Way ( $O(n)$ ):** Put everyone in their initial seats. Then, starting with the lowest-level managers (the last non-leaf nodes), you tell them to organize their immediate team. You work your way up the hierarchy. By the time you get to the CEO, most of the company is already in order, and only a few major adjustments are needed. This bottom-up approach is far more efficient.

This "smart way" is the `buildHeap` algorithm. We ignore the leaves (as they are already valid tiny heaps of one) and start by running our "sink-down" (`heapifyDown`) process on the last non-leaf node, working our way backward to the root.

## The Algorithm

1. Find the index of the last non-leaf node:  $(n / 2) - 1$ , where  $n$  is the number of elements.
2. Starting from this index, iterate backward to the root (index 0).
3. In each iteration, call `heapifyDown` on the current node to ensure it and its children satisfy the heap property.

**Why is it  $O(n)$ ?** While a full proof is complex, the intuition is that `heapifyDown` is cheap for nodes near the bottom of the tree. Since about half the nodes are leaves (costing 0) and three-quarters are in the bottom two levels (costing very little), the expensive operations on nodes near the root are rare. The total work averages out to be linear.



## C++ Implementation

We can add a new constructor to our `MaxHeap` class that takes a vector and efficiently converts it into a heap.

```
#include <iostream>
#include <vector>
#include <stdexcept>
#include <algorithm> // For std::swap

class MaxHeap {
private:
    std::vector<int> nodes;

    // sink-down process (from previous section)
    void heapifyDown(int index) {
        int leftChildIndex = 2 * index + 1;
        int rightChildIndex = 2 * index + 2;
        int largestIndex = index;
```

```

        if (leftChildIndex < nodes.size() && nodes[leftChildIndex] >
nodes[largestIndex]) {
            largestIndex = leftChildIndex;
        }
        if (rightChildIndex < nodes.size() && nodes[rightChildIndex] >
nodes[largestIndex]) {
            largestIndex = rightChildIndex;
        }

        if (largestIndex != index) {
            std::swap(nodes[index], nodes[largestIndex]);
            heapifyDown(largestIndex);
        }
    }

public:
// Default constructor
MaxHeap() = default;

// The EFFICIENT buildHeap constructor
MaxHeap(const std::vector<int>& array) {
    nodes = array; // Copy the array
    // Start from the last non-leaf node and heapify down
    for (int i = (nodes.size() / 2) - 1; i >= 0; --i) {
        heapifyDown(i);
    }
}

// (Other methods like insert, extractMax, peek, isEmpty go here...)

int peek() const { /* ... */ }
bool isEmpty() const { /* ... */ }
int extractMax() { /* ... */ }
};

int main() {
    std::vector<int> unordered_data = {15, 60, 5, 90, 30, 80, 40};

    // Build the heap in O(n) time using the new constructor
    MaxHeap heap(unordered_data);

    // The heap is now correctly structured. The top element should be 90.
    std::cout << "Heap built from array. The max element is: " << heap.peek()
    << std::endl;
}

```

```
    std::cout << "Extracting all elements to show heap property:" << std::endl;
    while (!heap.isEmpty()) {
        std::cout << heap.extractMax() << " "; // Expected: 90 80 60 40 30 15 5
    }
    std::cout << std::endl;

    return 0;
}
```

# Chapter 4: Advanced Heap Concepts & Applications

## 4.1 Heap Sort: In-Place Sorting with O(1) Space

Heap Sort is an efficient, in-place comparison-based sorting algorithm. It leverages the power of a heap data structure to achieve  $O(n \log n)$  time complexity while using only  $O(1)$  constant extra space.

### The "King of the Hill" Analogy

Imagine a game of "King of the Hill" to sort people by strength.

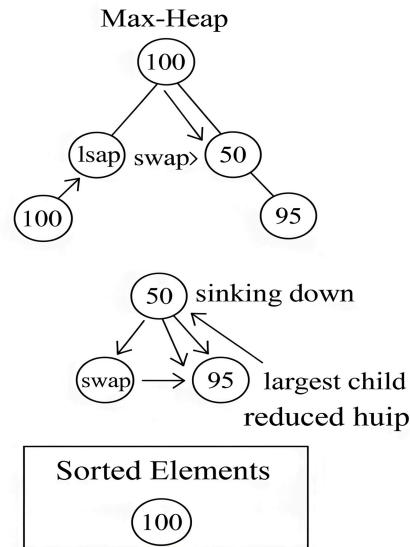
1. **Phase 1: Find the King.** Everyone scrambles up the hill. After a quick organization (the  $O(n)$  `buildHeap` process), the strongest person emerges at the very peak. This is our "king."
2. **Phase 2: Build the Winners' Circle.** The king is declared the first winner. We move them off the hill to the start of a "winners' circle" at the bottom. This spot is now sorted and final. A new player is put at the peak, and the remaining players quickly re-organize to find the *new* strongest person. This new king is then moved to the second spot in the winners' circle.

This process repeats—finding the king of the remaining group and moving them to the next spot in the sorted line—until no one is left on the hill. The winners' circle is now a perfectly sorted line from weakest to strongest.

### The Heap Sort Algorithm

The algorithm consists of two main phases, performed directly on the input array:

1. **Build Max-Heap:** First, convert the unsorted array into a **Max-Heap**. We use the efficient `buildHeap` method from the previous chapter, which takes  $O(n)$  time. After this step, the largest element in the array is at the root (index `0`).
2. **Repeatedly Extract Max:**
  - Swap the element at the root (`A[0]`) with the last element of the current heap (`A[i]`).
  - The largest element is now in its final sorted position at the end of the array. We decrease the considered size of the heap by one.
  - The new root may violate the heap property. Call `heapifyDown` on the root of the reduced heap to find the next largest element.
  - Repeat this process until the heap size is 1. The array is now sorted.



## C++ Implementation

Here is a complete, in-place Heap Sort implementation in C++. Notice it doesn't require any extra arrays for sorting.

```
#include <iostream>
#include <vector>
#include <algorithm> // for std::swap

// Helper function to sink an element down in a heap of a specific size
void heapifyDown(std::vector<int>& arr, int n, int index) {
    int largest = index;
    int left = 2 * index + 1;
    int right = 2 * index + 2;

    // Check if left child is larger than root
    if (left < n && arr[left] > arr[largest]) {
        largest = left;
    }
    // Check if right child is larger than Largest so far
    if (right < n && arr[right] > arr[largest]) {
        largest = right;
    }
    // If Largest is not root, swap and continue heapifying
    if (largest != index) {
        std::swap(arr[index], arr[largest]);
        heapifyDown(arr, n, largest);
    }
}
```

```

    }

}

// The main function to perform Heap Sort
void heapSort(std::vector<int>& arr) {
    int n = arr.size();

    // --- Phase 1: Build Max-Heap ---
    // Start from the Last non-Leaf node
    for (int i = n / 2 - 1; i >= 0; i--) {
        heapifyDown(arr, n, i);
    }

    // --- Phase 2: Repeatedly Extract Max ---
    // One by one, extract an element from the heap
    for (int i = n - 1; i > 0; i--) {
        // Move current root (max element) to the end
        std::swap(arr[0], arr[i]);
        // Call heapifyDown on the reduced heap of size 'i'
        heapifyDown(arr, i, 0);
    }
}

void printArray(const std::vector<int>& arr) {
    for (int val : arr) {
        std::cout << val << " ";
    }
    std::cout << std::endl;
}

int main() {
    std::vector<int> data = {12, 11, 13, 5, 6, 7};
    std::cout << "Unsorted array: ";
    printArray(data);

    heapSort(data);

    std::cout << "Sorted array: ";
    printArray(data); // Expected: 5 6 7 11 12 13

    return 0;
}

```

## 4.2 Finding the K-th Largest/Smallest Element

A common and practical problem is finding the k-th largest or smallest element in a collection without having to sort the entire collection. Heaps provide a highly efficient solution.

### The "VIP Lounge Bouncer" Analogy 🧑

Imagine a VIP lounge with a strict capacity of **k** guests. The bouncer's job is to ensure that at all times, the **k** most important people are inside.

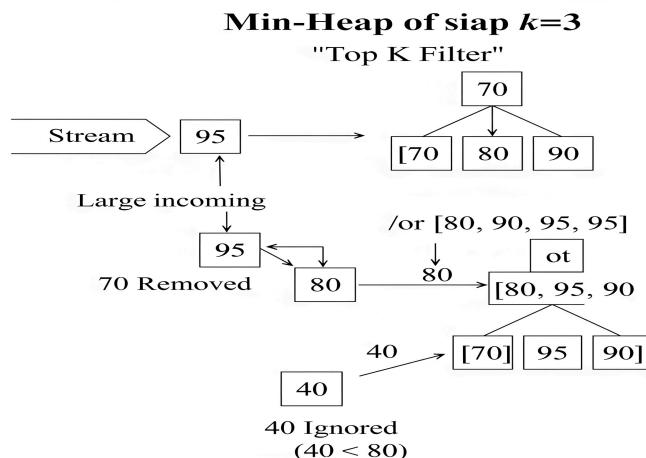
- The bouncer doesn't need to rank everyone waiting outside. They only need to know one thing: **who is the least important person currently inside the lounge?**
- When a new guest arrives, the bouncer compares them to this least important person.
- If the new guest is more important, the bouncer asks the least important person to leave and lets the new guest in. Otherwise, the new guest is turned away.

In this analogy, the VIP lounge is a **Min-Heap of size k**. The bouncer's logic is our algorithm, and the "least important person" is always at the root of the Min-Heap, available in O(1) time.

### The Algorithm (for K-th Largest)

1. Create a **Min-Heap**.
2. Iterate through the first **k** elements of the input array and push them into the heap. The heap now holds the top **k** elements seen so far.
3. For the rest of the elements in the array, compare the current element with the root of the heap (`heap.top()`).
4. If the current element is **larger** than the root, it means it deserves a spot in the "Top K". We `pop()` the smallest element (the root) and `push()` the current, larger element.
5. After iterating through the entire array, the heap contains the **k** largest elements. The root of the heap is the **k**-th largest element.

This approach is much faster than sorting the whole array (O( $n\log n$ )), especially when **k** is small compared to **n**.



## C++ Implementation

Here is a C++ function that finds the k-th largest element using a Min-Heap (`std::priority_queue` with a `std::greater` comparator).

```
#include <iostream>
#include <vector>
#include <queue>

// Finds the k-th Largest element in an array using a Min-Heap.
int findKthLargest(const std::vector<int>& nums, int k) {
    // Create a Min-Heap using a priority_queue.
    std::priority_queue<int, std::vector<int>, std::greater<int>> minHeap;

    for (int num : nums) {
        if (minHeap.size() < k) {
            minHeap.push(num);
        } else if (num > minHeap.top()) {
            // If the current number is bigger than the smallest in the heap,
            // it belongs in our "Top K".
            minHeap.pop(); // Remove the smallest.
            minHeap.push(num); // Add the new, larger number.
        }
    }

    // The root of the heap is now the k-th Largest element.
    return minHeap.top();
}

int main() {
    std::vector<int> data = {3, 2, 1, 5, 6, 4};
    int k = 2;
    int result = findKthLargest(data, k);
    std::cout << "The " << k << "nd largest element is: " << result <<
    std::endl; // Expected: 5

    std::vector<int> data2 = {3, 2, 3, 1, 2, 4, 5, 5, 6};
    k = 4;
    result = findKthLargest(data2, k);
    std::cout << "The " << k << "th largest element is: " << result <<
    std::endl; // Expected: 4

    return 0;
}
```

To find the **k-th smallest element**, the logic is identical, but you would use a **Max-Heap** to keep track of the **k** smallest elements.

## 4.3 Application Spotlight: Dijkstra's & Prim's Algorithms

Heaps, in the form of Priority Queues, are the key to making several classic graph algorithms efficient. They are used to constantly and quickly answer the question: "Out of all available options, which one is the cheapest or closest?" We'll spotlight two of the most famous: Dijkstra's and Prim's.

### Dijkstra's Algorithm: Finding the Shortest Path

Dijkstra's algorithm finds the shortest path from a starting node to all other nodes in a weighted graph, like a GPS finding the fastest route.

- **Analogy: The Spreading Fire**  Imagine starting a fire at one point in a grassy field. The fire doesn't spread equally; it spreads fastest along the driest paths (edges with the lowest weight). Dijkstra's algorithm explores the graph just like this. The "fire front" is all the nodes adjacent to the already "burnt" area. A Min-Heap acts as the brain of the fire, always knowing which blade of grass on the entire fire front is the easiest to ignite next (the node with the minimum distance from the start).
- **How the Heap Helps:** The algorithm needs to repeatedly find the unvisited node with the smallest known distance from the source. A naive search would be slow ( $O(V)$ ). A Min-Heap stores the unvisited nodes prioritized by distance, making this selection a super-fast  $O(\log V)$  operation. This optimization is crucial, changing the algorithm's complexity from a slow  $O(V^2)$  to a much faster  $O(E\log V)$ .

### Prim's Algorithm: Building a Minimum Spanning Tree

Prim's algorithm finds a Minimum Spanning Tree (MST)—a way to connect all nodes in a graph using the least possible total edge weight.

- **Analogy: Building a Road Network**  You need to build roads to connect a set of towns for the minimum possible cost. You start at one town. At each step, you look at all the potential roads that lead from your current, connected road network to a new, unconnected town. You always pick the absolute cheapest road to build next. A Min-Heap keeps track of all these potential "next roads," prioritized by cost, allowing you to instantly select the cheapest one to expand your network.

- **How the Heap Helps:** Just like Dijkstra's, Prim's needs to find the minimum-weight edge connecting the growing MST to a node outside of it. The Min-Heap stores these "fringe" edges, making the selection process fast and efficient.

## C++ Code Spotlight: Dijkstra's Algorithm

Here is a C++ implementation of Dijkstra's algorithm. Pay close attention to the `std::priority_queue` (our Min-Heap). It stores pairs of `{distance, vertex}` and drives the entire logic of the algorithm by ensuring we always process the closest node next.

```
#include <iostream>
#include <vector>
#include <queue>
#include <limits>

using namespace std;

const int INF = numeric_limits<int>::max();

// Represents a weighted edge in the graph
struct Edge {
    int to;
    int weight;
};

// Represents the graph using an adjacency list
using Graph = vector<vector<Edge>>;

void dijkstra(const Graph& graph, int start_node) {
    int num_vertices = graph.size();
    vector<int> distances(num_vertices, INF);

    // Min-Heap: stores {distance, vertex}.
    // `greater` makes it a min-heap based on the first element (distance).
    priority_queue<pair<int, int>, vector<pair<int, int>>, greater<pair<int, int>>> pq;

    // Start at the source node
    distances[start_node] = 0;
    pq.push({0, start_node}); // {distance, vertex}

    while (!pq.empty()) {
        // 1. Get the closest unvisited node from the heap
        int u = pq.top().second;
        int d = pq.top().first;
```

```

    pq.pop();

    // If we've found a shorter path already, skip
    if (d > distances[u]) {
        continue;
    }

    // 2. Explore its neighbors
    for (const auto& edge : graph[u]) {
        int v = edge.to;
        int weight = edge.weight;

        // 3. If a shorter path is found, update it and push to the heap
        if (distances[u] + weight < distances[v]) {
            distances[v] = distances[u] + weight;
            pq.push({distances[v], v});
        }
    }
}

// Print the shortest distances
cout << "Shortest distances from node " << start_node << ":" << endl;
for (int i = 0; i < num_vertices; ++i) {
    if (distances[i] == INF) {
        cout << "Node " << i << ": infinity" << endl;
    } else {
        cout << "Node " << i << ": " << distances[i] << endl;
    }
}
}

int main() {
    int V = 5;
    Graph graph(V);

    graph[0].push_back({1, 9});
    graph[0].push_back({2, 6});
    graph[0].push_back({3, 5});
    graph[1].push_back({4, 3});
    graph[2].push_back({1, 2});
    graph[2].push_back({3, 4});

    dijkstra(graph, 0);
    return 0;
}

```

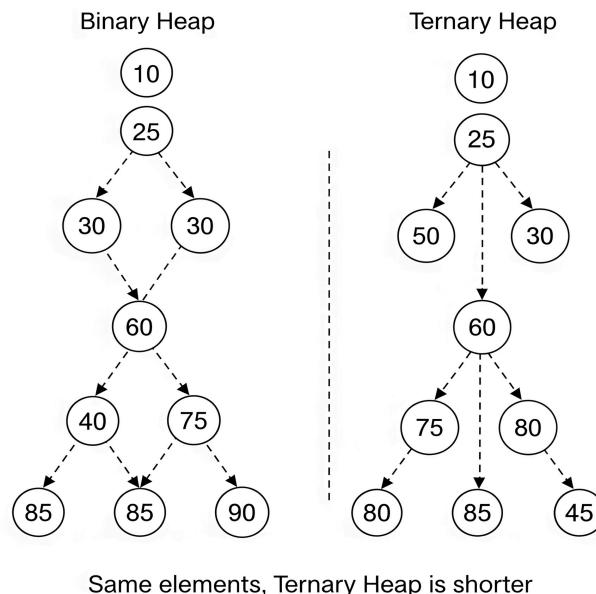
## 4.4 Advanced Variants (A Brief Overview)

This section provides a high-level look at more advanced heap structures. A full implementation is beyond the scope of this tutorial, but understanding their purpose shows a deeper command of data structures.

### 4.4.1 d-ary Heaps

A **d-ary heap** is a generalization of a binary heap where each node has **d** children instead of just two.

- **Analogy: Flatter Management Structure**  A binary heap is like a company with a strict rule: every manager has exactly two subordinates. A d-ary heap is like a company that allows a manager to have 'd' subordinates (e.g., 3, 4, or 5). This makes the company hierarchy (the tree) flatter and wider. The benefit is a shorter chain of command (reduced tree height), but the downside is that each manager has more people to compare when finding the best performer (more comparisons in a `heapifyDown` operation).
- **When to use it:** A d-ary heap can be faster than a binary heap when the `decreaseKey` operation is more frequent than `extractMin`, or on hardware where branching is expensive. The reduced height ( $O(\log_{dn})$ ) can lead to fewer swaps.



- **C++ Code (Navigation Formulas):** The implementation is very similar to a binary heap; only the navigation formulas change.

```
#include <iostream>

// Navigation for a d-ary heap
class D_aryHeap {
private:
    int d; // Number of children per node

public:
    D_aryHeap(int num_children) : d(num_children) {}

    int parent_idx(int i) {
        return (i - 1) / d;
    }

    // Get the k-th child of node i (where k is from 0 to d-1)
    int child_idx(int i, int k) {
        return d * i + k + 1;
    }
};

int main() {
    D_aryHeap ternary_heap(3);
    // Find the parent of node at index 10
    std::cout << "Parent of node 10 in a 3-ary heap is: "
        << ternary_heap.parent_idx(10) << std::endl; // Output: 3

    // Find the 2nd child (k=1) of node 3
    std::cout << "2nd child of node 3 in a 3-ary heap is: "
        << ternary_heap.child_idx(3, 1) << std::endl; // Output: 11
}
```

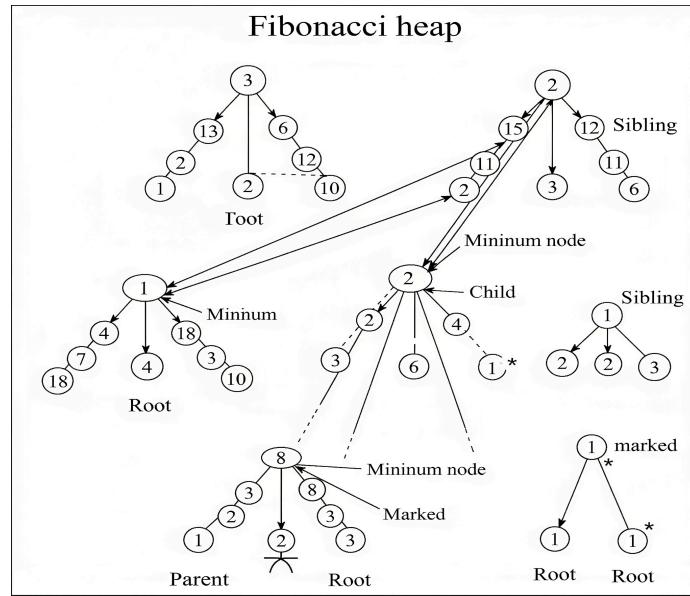
#### 4.4.2 Binomial & Fibonacci Heaps

These are even more advanced priority queues often used in high-performance algorithms. They are not single trees but a **collection (or forest) of heap-ordered trees**.

- Analogy: The Lazy Procrastinator Manager 😴  
A Fibonacci Heap is the ultimate lazy manager. Instead of carefully structuring its team (insert), it just throws new employees into a room and says "you're hired." It postpones all organizational work until it's absolutely forced to (extractMin). When a task's priority needs to be updated (decreaseKey), it's incredibly fast because the structure is so loose.

This "lazy" or "amortized" approach makes them one of the fastest priority queues for algorithms like Dijkstra's on dense graphs. Binomial Heaps are similar but slightly more structured.

- **When to use them:** Their main advantage is an extremely fast (amortized O(1)) `decreaseKey` operation and a fast `merge` operation. They are the preferred choice in network routing algorithms and advanced implementations of Dijkstra's and Prim's algorithms.



- **C++ Code (Conceptual Interface):** A full implementation of a Fibonacci heap is exceptionally complex. What's important is to understand the powerful interface it provides.

```
// This is a conceptual C++ interface, not a full implementation.
// It highlights the key operations that make these heaps powerful.
```

```
template <typename T>
class AdvancedPriorityQueue {
public:
    // A handle or pointer to a node, needed for decreaseKey
    using NodeHandle = Node<T>*;

    // Standard insert
    NodeHandle insert(const T& value);

    // Standard extractMin
    T extractMin();
```

```
// The key operation: O(1) amortized time!
// This is much faster than the O(log n) of a binary heap.
void decreaseKey(NodeHandle node, const T& new_value);

// Fast merging of two heaps
void merge(AdvancedPriorityQueue& other_heap);
};
```

These advanced variants show the trade-offs involved in data structure design, where relaxing some structural rules (like being a single, complete tree) can lead to huge performance gains for specific operations.

# Chapter 5: Heap Problem-Solving Workshop

## 5.1 Problem 1: Find Median from a Data Stream

Let's put our heap knowledge to the test with a classic interview problem. This challenge beautifully demonstrates how two heaps can work together to solve a seemingly complex task.

### The Problem: Find Median from a Data Stream

You are asked to design a data structure that supports two operations:

1. `addNum(int num)`: Add an integer from a data stream to the data structure.
2. `findMedian()`: Return the median of all elements seen so far.

The **median** is the middle value in an ordered list of numbers. If the list size is even, the median is the average of the two middle values.

### The "Two Halves of a Bridge" Analogy

A naive approach of sorting the list every time is too slow ( $O(n\log n)$ ). We need a way to access the "middle" elements instantly.

Imagine you're building a bridge across a river, and people (numbers) are arriving one by one. You want to split them into two groups on opposite banks to find the median person's height.

- **The Left Bank:** Holds the shorter half of the people. We'll use a **Max-Heap** for this group, so we can always instantly see the **tallest person on the short side**.
- **The Right Bank:** Holds the taller half of the people. We'll use a **Min-Heap** for this group, so we can always instantly see the **shortest person on the tall side**.

These two people at the river's edge (the tops of our heaps) are the middle elements. The median is either one of them or their average. Our main job is to keep the two groups balanced in size.

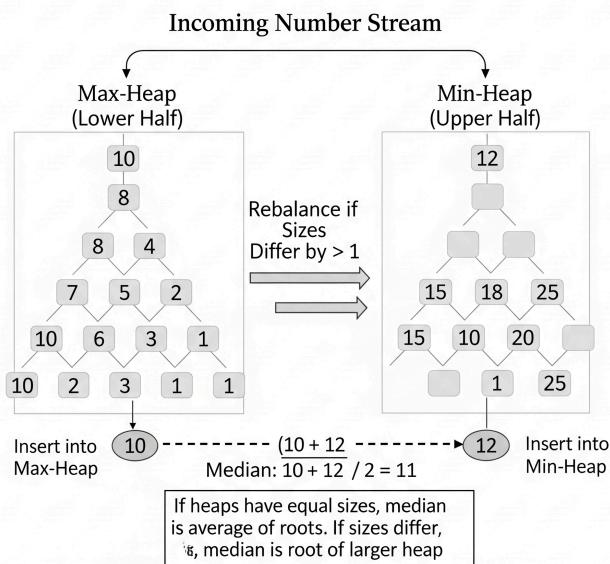
### The Algorithm

1. Maintain two heaps: a `smallers` Max-Heap and a `largers` Min-Heap.
2. **Balancing Act:** We'll enforce a rule: the `smallers` heap can have at most one more element than the `largers` heap.
3. `addNum(num):`
  - Add the new number to the `smallers` (Max-Heap).
  - To maintain the partition, take the largest element from `smallers` (`smallers.top()`) and move it to `largers`.
  - If `largers` now has more elements than `smallers`, move its smallest element (`largers.top()`) back to `smallers`. This rebalances the sizes.

4. `findMedian()`:

- If the heaps have different sizes, the median is the top of the larger heap (`smallers.top()`).
- If they have the same size, the median is the average of the two tops:  $(\text{smallers.top()} + \text{largers.top()}) / 2.0$ .

## Median of a Number Stream using Two Heaps



## C++ Implementation

Here is a full `MedianFinder` class in C++ that implements this logic.

```
#include <iostream>
#include <vector>
#include <queue>

class MedianFinder {
private:
    // smallers stores the smaller half of numbers, but as a Max-Heap
    // so we can quickly get the largest of the small numbers.
    priority_queue<int> smallers;

    // largers stores the larger half of numbers, as a Min-Heap
    // so we can quickly get the smallest of the large numbers.
    priority_queue<int, vector<int>, greater<int>> largers;
```

```

public:
    MedianFinder() {}

    void addNum(int num) {
        // 1. Add to the 'smallers' heap first.
        smallers.push(num);

        // 2. Balance the heaps' values:
        // Ensure every number in 'smallers' is <= every number in 'Lagers'.
        // To do this, move the largest from 'smallers' to 'Lagers'.
        largers.push(smaller.top());
        smaller.pop();

        // 3. Balance the heaps' sizes:
        // If 'Lagers' has more elements, it means the balance is off.
        // Move the smallest from 'Lagers' back to 'smallers'.
        if (largers.size() > smaller.size()) {
            smaller.push(lagers.top());
            largers.pop();
        }
    }

    double findMedian() {
        if (smallers.size() > largers.size()) {
            // If total count is odd, the median is the top of the larger heap.
            return smaller.top();
        } else {
            // If total count is even, it's the average of the two tops.
            return (smaller.top() + largers.top()) / 2.0;
        }
    }
};

int main() {
    MedianFinder mf;

    mf.addNum(1);
    // smaller: {1}, Lagers: {} -> Median is 1
    cout << "Median is: " << mf.findMedian() << endl;

    mf.addNum(2);
    // smaller: {1}, Lagers: {2} -> Median is (1+2)/2 = 1.5
    cout << "Median is: " << mf.findMedian() << endl;

    mf.addNum(3);
}

```

```

// smallers: {2, 1}, Largers: {3} -> Median is 2
cout << "Median is: " << mf.findMedian() << endl;

mf.addNum(4);
// smallers: {2, 1}, Largers: {3, 4} -> Median is (2+3)/2 = 2.5
cout << "Median is: " << mf.findMedian() << endl;

return 0;
}

```

## 5.2 Problem 2: Merge K Sorted Lists

This next problem is another interview favorite. It demonstrates how a heap can be used to efficiently manage multiple sorted sources of data.

### The Problem: Merge K Sorted Lists

You are given an array of **k** linked lists, where each linked list is sorted in ascending order. Your task is to merge all of them into one single sorted linked list.

### The "Race Finish Line" Analogy 🏁

Trying to solve this by concatenating all lists and then sorting is inefficient because it ignores that the lists are already sorted. A better approach is needed.

Imagine you are a judge at a race with **k** lanes. In each lane, the runners are already sorted by their qualifying time (fastest at the front). Your job is to find the overall ranking.

You don't need to look at every runner at once. You only need to focus on the **front-runner in each of the k lanes**. To find the next person to cross the finish line, you just compare these **k** front-runners and pick the fastest.

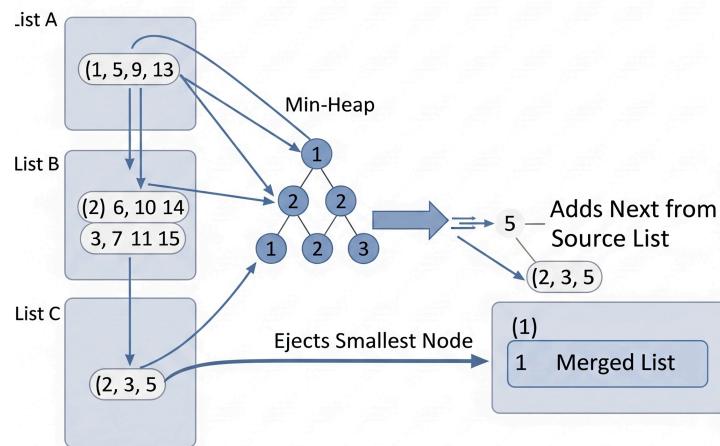
A **Min-Heap** is the perfect assistant for this. It keeps track of the **k** front-runners and instantly tells you which one is the fastest overall (the one with the smallest value). Once you record that runner, the *next* runner from that *same lane* takes their place at the front, and your assistant updates their list. You repeat this until every lane is empty.

### The Algorithm

1. Create a **Min-Heap** that will store pointers to the linked list nodes (`ListNode*`). It will order them by the node's value.
2. Initialize the heap by pushing the head node of each of the **k** lists into it.
3. Create a **dummy head** for your result list and a **tail** pointer to build upon it.

4. Loop while the heap is not empty:
  - a. Get Min: pop() the node with the smallest value from the heap. This is our next node in the sorted list.
  - b. Append: Attach this node to the tail of the result list and advance the tail.
  - c. Add Next: If the node you just popped has a next element in its original list, push that next element into the heap.
5. Return the list starting from `dummy->next`.

### K-way merge Sort



### C++ Implementation

For this problem, we need to tell our `priority_queue` how to compare `ListNode*` pointers. We do this by creating a custom comparator struct.

```
#include <iostream>
#include <vector>
#include <queue>

using namespace std;

// Definition for singly-linked list.
struct ListNode {
    int val;
    ListNode *next;
}
```

```

    ListNode(int x) : val(x), next(nullptr) {}

};

// Custom comparator for the priority queue.
// It tells the heap to compare nodes based on their 'val'.
struct CompareNodes {
    bool operator()(const ListNode* a, const ListNode* b) {
        // We want a Min-Heap, so we use > to order
        // the smaller values to have higher priority.
        return a->val > b->val;
    }
};

class Solution {
public:
    ListNode* mergeKLists(vector<ListNode*>& lists) {
        // A Min-Heap of ListNode pointers, using our custom comparator.
        priority_queue<ListNode*, vector<ListNode*>, CompareNodes> minHeap;

        // 1. Push the head of each non-empty list into the heap.
        for (ListNode* head : lists) {
            if (head != nullptr) {
                minHeap.push(head);
            }
        }

        // 2. Create a dummy head for the result list.
        ListNode* dummy = new ListNode(0);
        ListNode* tail = dummy;

        // 3. Loop while the heap has nodes.
        while (!minHeap.empty()) {
            // 4. Get the node with the smallest value.
            ListNode* smallestNode = minHeap.top();
            minHeap.pop();

            // 5. Append it to our result list.
            tail->next = smallestNode;
            tail = tail->next;

            // 6. If there's a next node in that list, add it to the heap.
            if (smallestNode->next != nullptr) {
                minHeap.push(smallestNode->next);
            }
        }
    }
}

```

```

        ListNode* result_head = dummy->next;
        delete dummy; // Clean up the dummy node.
        return result_head;
    }
};

// Helper function to print a list
void printList(ListNode* head) {
    while (head != nullptr) {
        cout << head->val << " -> ";
        head = head->next;
    }
    cout << "NULL" << endl;
}

int main() {
    Solution sol;

    // Create 3 sorted lists
    ListNode* l1 = new ListNode(1);
    l1->next = new ListNode(4);
    l1->next->next = new ListNode(5);

    ListNode* l2 = new ListNode(1);
    l2->next = new ListNode(3);
    l2->next->next = new ListNode(4);

    ListNode* l3 = new ListNode(2);
    l3->next = new ListNode(6);

    vector<ListNode*> lists = {l1, l2, l3};

    ListNode* merged_head = sol.mergeKLists(lists);

    cout << "Merged list: ";
    printList(merged_head); // Expected: 1 -> 1 -> 2 -> 3 -> 4 -> 4 -> 5 -> 6
-> NULL

    return 0;
}

```

## Part II: Tries

# Chapter 6: Trie Fundamentals

Welcome to the second half of our journey! We're leaving the world of priority and numbers to enter the realm of characters and words. Let's meet the Trie, a data structure purpose-built for string operations.

## 6.1 What is a Trie? The "Autocomplete" Analogy

A **Trie** (pronounced "try," from the word "retrieval"), also known as a **Prefix Tree**, is a special tree used to store and search for strings efficiently.<sup>1</sup> Unlike other trees that store whole keys in a node, a Trie's structure is defined by the characters of the keys themselves.<sup>2</sup>

The best real-world analogy is the **autocomplete** feature in a search engine or your phone's keyboard.<sup>3</sup>

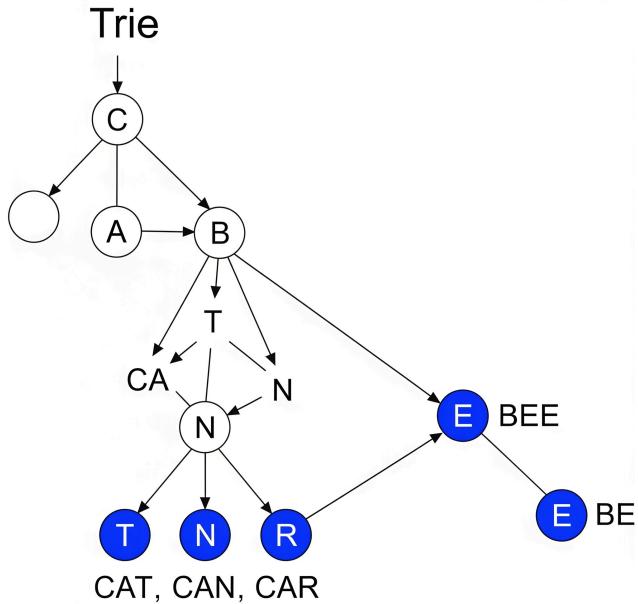
- The Analogy: Autocomplete When you start typing "comp" into a search bar, it doesn't scan a massive list of every word in the dictionary. Instead, it navigates a Trie. It follows the path c -> o -> m -> p. The node it lands on represents the prefix "comp," and from that point, it can quickly find all possible descendant paths to suggest words like "computer," "compiler," and "component." Each character you type takes you one level deeper into the tree.

## The Structure of a Trie

A Trie is made of nodes where:

- The **root** node is empty.<sup>4</sup>
- Each **node** represents a single character.<sup>5</sup>
- A **path** from the root to a node represents a prefix.<sup>6</sup>
- A special **flag** on a node marks if it's the end of a valid word.<sup>7</sup> This is crucial to distinguish a word like "car" from a prefix like "car" in "careful."

## The C++ Building Block: The TrieNode



Before we build the whole data structure, let's look at its fundamental component: the `TrieNode`. Each node must contain two things: pointers to its potential children and a flag to mark the end of a word.

For the English alphabet (a-z), a simple array of 26 pointers is a common and efficient way to store children.

```
#include <iostream>

// The number of possible characters (a-z)
const int ALPHABET_SIZE = 26;

// The fundamental building block of a Trie
struct TrieNode {
    // An array of pointers to child nodes
    TrieNode* children[ALPHABET_SIZE];

    // isEndOfWord is true if the node represents the end of a word
    bool isEndOfWord;

    // Constructor to initialize a new node
    TrieNode() {
        isEndOfWord = false;
        for (int i = 0; i < ALPHABET_SIZE; i++) {
            children[i] = nullptr;
        }
    }
}
```

```

        }
    }

};

int main() {
    std::cout << "Creating a new Trie root node..." << std::endl;

    // The root of our Trie will be a single TrieNode
    TrieNode* root = new TrieNode();

    // In the next sections, we will use this root
    // to build our Trie class and add operations.
    std::cout << "Root node created successfully." << std::endl;

    // In a real program, you'd need to manage memory and delete nodes.
    delete root;

    return 0;
}

```

This simple `TrieNode` structure is the foundation for all the powerful prefix-based operations we'll implement next.

## 6.2 The Trie Node: Building Blocks of Words

A `TrieNode` is the fundamental building block of a Trie. While the entire Trie represents a dictionary, each node is like a single letter-routing station.

### The "Signpost at a Crossroads" Analogy

Think of each `TrieNode` as a **signpost at a crossroads**.

- It doesn't know the full destination (the complete word).
- It has up to 26 different signposts pointing down different roads, one for each letter of the alphabet. This is its array of **children pointers**.
- It also has a special marker that says, "This crossroads itself is a valid destination!" This is the **isEndOfWord flag**.

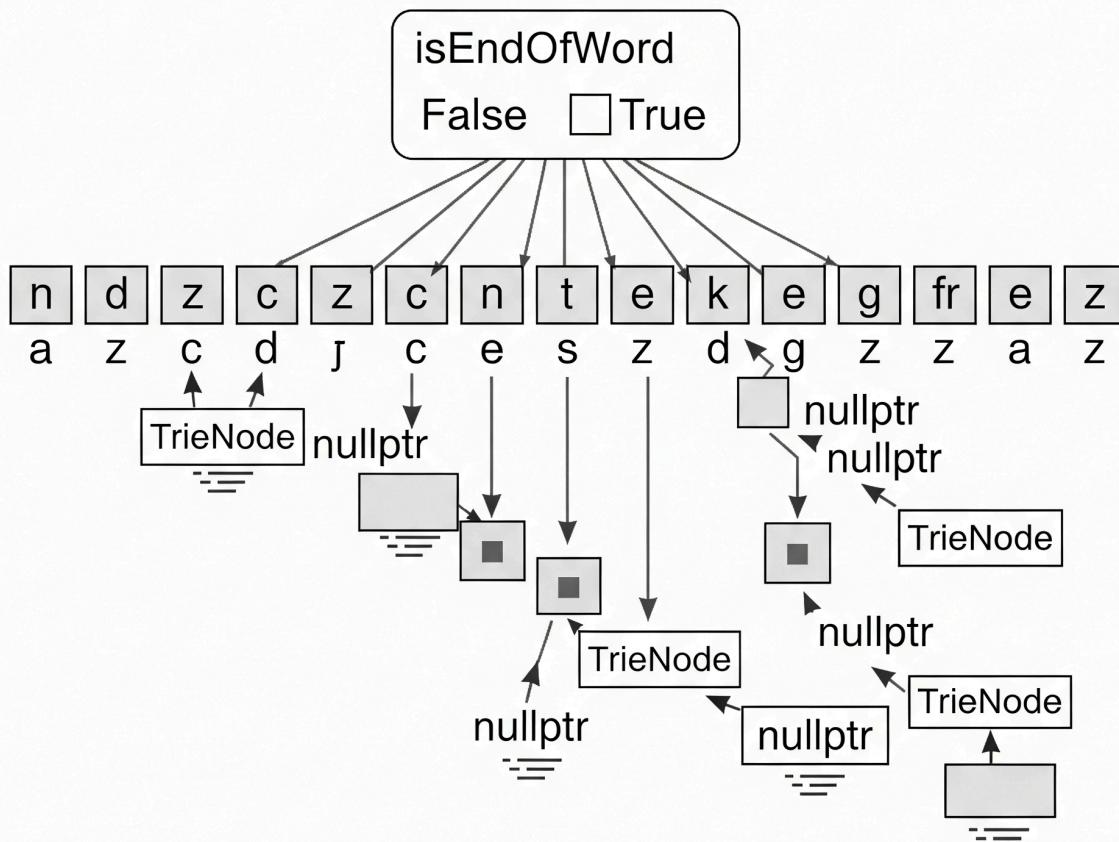
For example, when spelling "CAR," you start at the root, follow the sign for 'C' to the next crossroads. From there, you follow the sign for 'A', and then 'R'. The node for 'R' will have its "valid destination" flag turned on.

### The Anatomy of a `TrieNode`

Each node needs to store two critical pieces of information:

1. **Pointers to Children:** A way to reference the next node in the path for each possible character. For lowercase English letters (a-z), a simple array of 26 pointers is perfect. `children[0]` would point to the 'a' node, `children[1]` to the 'b' node, and so on.
2. **isEndOfWord Flag:** A boolean that tells us if a path ending at this specific node constitutes a complete word in our dictionary. This is what distinguishes the prefix "bee" from the actual word "bee".

## TrieNode



## The C++ TrieNode Struct

Here is the C++ struct for our `TrieNode`. The code is simple, but it's the heart of the entire Trie data structure. We've added a small demonstration of how a character like 'c' maps to an index in the `children` array.

```
#include <iostream>

const int ALPHABET_SIZE = 26;

// A single node in the Trie data structure.
struct TrieNode {
    // children[0] corresponds to 'a', children[1] to 'b', and so on.
    // If a pointer is nullptr, it means that path does not exist.
    TrieNode* children[ALPHABET_SIZE];

    // This flag is true if the node represents the end of a complete word.
    bool isEndOfWord;

    // Constructor to initialize a new node with null children.
    TrieNode() {
        isEndOfWord = false;
        for (int i = 0; i < ALPHABET_SIZE; i++) {
            children[i] = nullptr;
        }
    }
};

int main() {
    TrieNode* root = new TrieNode();

    // Let's imagine we want to add the word "cat".
    // The first letter is 'c'.
    char first_char = 'c';
    int index = first_char - 'a'; // This gives 2

    // Conceptually, we would link the root to a new node for 'c'.
    if (root->children[index] == nullptr) {
        std::cout << "Path for '" << first_char << "' does not exist. Creating it." << std::endl;
        root->children[index] = new TrieNode();
    }

    // Our full insertion logic in the next chapter will repeat this process
    // for every character in a word.
}
```

```

    // Clean up memory
    delete root->children[index];
    delete root;

    return 0;
}

```

## 6.3 Why Choose a Trie over a Hash Table?

Both Tries and Hash Tables can store and search for strings, so when should you choose one over the other? The answer lies in what kind of questions you need to ask your data.

### The "Encyclopedia vs. Filing Cabinet" Analogy

- A **Hash Table** is like a high-tech **filing cabinet**. You give it a document (a string), it computes a code (a hash), and stores the document in a specific, seemingly random drawer. It's incredibly fast ( $O(1)$  average) for two things: putting a document in and retrieving a specific document if you know its exact title. But if you ask it for "all documents starting with 'Project X,'" it's useless—you'd have to open every single drawer.
- A **Trie** is like the **index of an encyclopedia**. All topics are arranged alphabetically. Finding a specific topic is easy. More importantly, finding all topics that start with "Comp" is trivial—you just go to that section. You can also read the index from start to finish to get a sorted list of all entries.

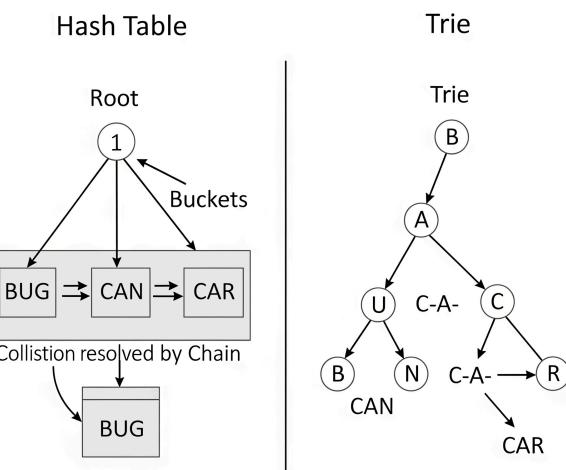
**The Bottom Line:** Use a Hash Table for lightning-fast *exact* matches. Choose a Trie when you need to perform operations related to prefixes or alphabetical order.

### Key Differences

Feature	Hash Table	Trie
<b>Exact Search</b> (search)	<b>Excellent</b> ( $O(1)$ avg)	Good ( $O(L)$ )
<b>Prefix Search</b> (startsWith)	<b>Inefficient</b> ( $\$O(N^*L\$)$ )	<b>Excellent</b> ( $O(L)$ )

Feature	Hash Table	Trie
Ordered Traversal	Not Possible	Possible
Space Usage	Can be high; stores whole strings	Efficient if many prefixes are shared

$L$  = length of the string,  $N$  = number of strings in the set.



## The C++ Code Test: A Prefix Search

Let's see what it would take to find all words starting with "ca" using both structures. This example makes the architectural difference crystal clear.

```
#include <iostream>
#include <vector>
#include <string>
#include <unordered_set>

using namespace std;

// --- Using a Hash Table (unordered_set) ---
```

```

void find_prefix_in_hash_table(const unordered_set<string>& dict, const string& prefix) {
    cout << "Searching in Hash Table for prefix '" << prefix << ":" << endl;
    int checks = 0;
    for (const string& word : dict) {
        checks++;
        // We must check every word to see if it starts with the prefix.
        if (word.find(prefix, 0) == 0) {
            cout << " - Found: " << word << endl;
        }
    }
    cout << " (Total checks performed: " << checks << ")" << endl;
}

// --- Conceptual Trie Usage ---
// In the next chapter, we'll implement a Trie class that can do this
efficiently.
// The code would look clean and direct, like this:
/*
void find_prefix_in_trie(const Trie& trie, const string& prefix) {
    cout << "Searching in Trie for prefix '" << prefix << ":" << endl;
    vector<string> results = trie.findAllWithPrefix(prefix); // This is an
efficient operation
    for (const string& word : results) {
        cout << " - Found: " << word << endl;
    }
    cout << " (No need to check every word in the dictionary!)" << endl;
}
*/

int main() {
    unordered_set<string> dictionary = {"cat", "car", "cart", "dog", "dove"};
    string prefix_to_find = "car";

    // With a hash table, we have no choice but to iterate through everything.
    find_prefix_in_hash_table(dictionary, prefix_to_find);

    // With a Trie, the search would be a direct, efficient traversal,
    // avoiding the need to look at "dog" or "dove" at all.

    return 0;
}

```

As the code shows, the hash table has to check every single word. A Trie, by its very design, would navigate directly to the "car" prefix and only explore from there, making it the clear winner for prefix-based applications like autocomplete.

# Chapter 7: Implementing a Trie

## 7.1 The Node Structure in Code

It's time to get our hands dirty and build a Trie from the ground up. The first step in implementing any complex data structure is to perfect its fundamental component. For a Trie, that's the `TrieNode`.

### The "Lego Brick" Analogy

Think of a `TrieNode` as a single, specialized **Lego brick**. On its own, it's simple. But it's designed with specific connection points that allow it to combine with other bricks to create an elaborate structure.

- The **stud**s on top are like the `isEndOfWord` flag, marking a potential completion point.
- The **tube**s on the bottom are like the `children` array, providing 26 possible connection points for the next brick (or letter) in the sequence.

Our job is to design this one perfect brick. Once we have that, building the entire Trie is just a matter of connecting these bricks correctly.

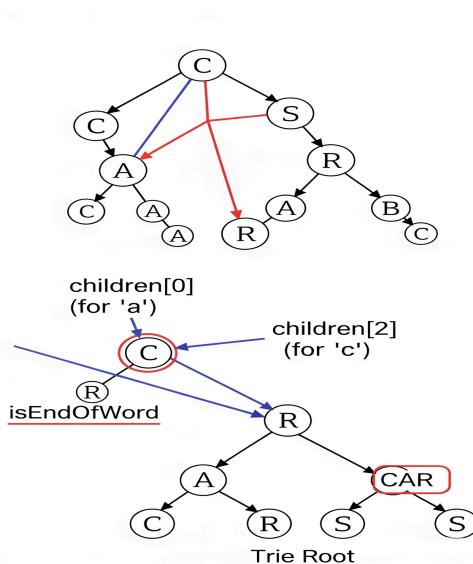
### The `TrieNode` in C++

To build our Trie, we need a `TrieNode` that holds two key pieces of information:

1. **An array of pointers** to its children, representing the next possible letters.
2. **A boolean flag** to tell us if this node marks the end of a complete word.

The C++ struct below is the direct translation of this requirement. We place it inside our `Trie` class to keep our code organized and encapsulated.

```
struct TrieNode
struct TrieNode {
    TrieNode* children[26];
    bool isEndOfWord;
    TrieNode();
}
```



## The C++ Code

Here is the implementation of our main `Trie` class, starting with the private, nested `TrieNode` struct. The `Trie` class itself only needs one member: a pointer to the root node. The constructor takes care of creating this initial empty node.

```
#include <iostream>

const int ALPHABET_SIZE = 26;

class Trie {
private:
    // Private nested struct for the TrieNode.
    // This is a common C++ practice for component-based data structures.
    struct TrieNode {
        TrieNode* children[ALPHABET_SIZE];
        bool isEndOfWord;

        // The constructor initializes a new node.
        // It's crucial to set children to nullptr to avoid garbage pointers.
        TrieNode() {
            isEndOfWord = false;
            for (int i = 0; i < ALPHABET_SIZE; ++i) {
                children[i] = nullptr;
            }
        }
    };
    TrieNode* root;

public:
    // Trie constructor creates the root node to start with.
    Trie() {
        root = new TrieNode();
    }

    // In the next sections, we will add methods here, such as:
    // void insert(const std::string& word);
    // bool search(const std::string& word);
    // ... and so on.

    // A destructor would also be needed in a real application to free memory.
    // ~Trie();
};

};
```

```

int main() {
    std::cout << "Creating an instance of the Trie class..." << std::endl;

    // When we create a Trie object, its constructor automatically
    // creates the root node for us.
    Trie my_dictionary;

    std::cout << "Trie object created, root node is initialized." << std::endl;
    std::cout << "Ready to start implementing insert and search operations!" <<
    std::endl;

    return 0;
}

```

With this solid foundation, we are now ready to add methods to our `Trie` class to bring it to life.

## 7.2 Core Operations (with Code & Complexity Analysis)

Here are the core operations that bring our Trie to life. We'll add these methods to the `Trie` class we started in the previous section.

### 7.2.1 `insert()`: Adding a Word to the Lexicon

This operation adds a word to the Trie, creating new nodes as necessary.

- **Analogy:** Paving a New Road 

Imagine you're a city planner building a road to a new house. You follow the existing road network as far as you can (matching an existing prefix). If the road ends before you reach your destination, you start paving new segments (creating new `TrieNodes`) for each remaining letter. When you finally reach the location, you put up a mailbox—this is our `isEndOfWord = true` flag, signifying a complete address.

- **The Process:** We start at the root. For each character in the word, we check if a path exists. If not, we create a new node. We then move to the next node and repeat. After the last character, we mark the final node as the end of a word.
- **Complexity:** The work done is proportional to the number of characters in the word. If the word has length **L**, the complexity is  $O(L)$ .

### 7.2.2 `search()`: Finding an Exact Match

This operation checks if a complete, exact word exists in our Trie.

- **Analogy:** Confirming a Full Address 

You're a package courier with a full address. You follow the map turn-by-turn (character

by character). If a turn leads to a dead end (a `nullptr`), the address is invalid. If you successfully reach the final destination, you're not done yet! You must check if there's actually a house there (the `isEndOfWord` flag must be true). Just being able to navigate to a spot isn't enough; it has to be a registered address.

- **The Process:** We traverse the Trie character by character. If the path ever breaks, the word doesn't exist. If we finish the traversal, we must check the `isEndOfWord` flag of the final node.
- **Complexity:** We visit one node for each character. For a word of length  $L$ , the complexity is  $O(L)$ .

### 7.2.3 `startsWith()`: Checking for a Prefix

This operation checks if there is any word in the Trie that starts with the given prefix.

- **Analogy:** Checking an Area Code  You're working at a call center and just need to know if an area code is valid. You trace the path on your map for the given digits. If you can trace the entire path without hitting a dead end, the area code is valid. You don't care if there's a specific house at the end—only that the path itself exists.
- **The Process:** This is identical to `search`, but with one key difference: we ignore the `isEndOfWord` flag at the end. As long as the traversal completes, the prefix exists.
- **Complexity:** Proportional to the length of the prefix  $L$ , making it  $O(L)$ .

## Complete C++ Implementation

Here is our `Trie` class with the core methods implemented.

```
#include <iostream>
#include <string>
#include <vector>

using namespace std;

const int ALPHABET_SIZE = 26;

class Trie {
private:
    struct TrieNode {
        TrieNode* children[ALPHABET_SIZE];
        bool isEndOfWord;

        TrieNode() {
            isEndOfWord = false;
            for (int i = 0; i < ALPHABET_SIZE; ++i) {
                children[i] = nullptr;
            }
        }
    };
    TrieNode* root;
};
```

```

        }
    }
};

TrieNode* root;

public:
Trie() {
    root = new TrieNode();
}

void insert(const string& word) {
    TrieNode* current = root;
    for (char ch : word) {
        int index = ch - 'a';
        if (current->children[index] == nullptr) {
            current->children[index] = new TrieNode();
        }
        current = current->children[index];
    }
    current->isEndOfWord = true;
}

bool search(const string& word) {
    TrieNode* current = root;
    for (char ch : word) {
        int index = ch - 'a';
        if (current->children[index] == nullptr) {
            return false; // Path does not exist
        }
        current = current->children[index];
    }
    // Path exists, but is it a complete word?
    return current->isEndOfWord;
}

bool startsWith(const string& prefix) {
    TrieNode* current = root;
    for (char ch : prefix) {
        int index = ch - 'a';
        if (current->children[index] == nullptr) {
            return false; // Path does not exist
        }
        current = current->children[index];
    }
}

```

```

        // Path exists, so a word with this prefix exists.
        return true;
    }

};

int main() {
    Trie dictionary;

    dictionary.insert("hello");
    dictionary.insert("her");
    dictionary.insert("help");
    dictionary.insert("apple");

    cout << boolalpha; // Print 'true' or 'false' for booleans

    // --- Search Tests ---
    cout << "Searching for 'hello': " << dictionary.search("hello") << endl;
// true
    cout << "Searching for 'her': " << dictionary.search("her") << endl;      //
true
    cout << "Searching for 'hell': " << dictionary.search("hell") << endl;
// false (it's a prefix, not a word)
    cout << "Searching for 'app': " << dictionary.search("app") << endl;
// false (prefix)
    cout << "Searching for 'apple': " << dictionary.search("apple") << endl;
// true
    cout << "Searching for 'world': " << dictionary.search("world") << endl;
// false

    cout << "---" << endl;

    // --- startsWith Tests ---
    cout << "Prefix 'he': " << dictionary.startsWith("he") << endl;      // true
    cout << "Prefix 'hell': " << dictionary.startsWith("hell") << endl;      //
true
    cout << "Prefix 'app': " << dictionary.startsWith("app") << endl;      //
true
    cout << "Prefix 'wor': " << dictionary.startsWith("wor") << endl;      //
false

    return 0;
}

```

## 7.3 The Tricky Task of delete()

Deleting a word from a Trie is the most complex of its core operations. Unlike `insert`, we can't just blaze a path through the tree. We have to be careful not to destroy nodes that are part of other words. This requires a delicate, recursive cleanup.

### The "Pruning a Tree" Analogy 🌳

Imagine the Trie is a tree in your garden, where branches are prefixes and leaves are words. You want to remove one specific leaf, say the word "TEAM".

You can't just chop off the "T" branch, as that would also destroy "TEA" and "TED". The correct approach is to work backward:

1. Go to the end of the "TEAM" path (the 'M' leaf) and prune it.
2. Step back to the 'A' branch. If 'A' doesn't lead to any other words, you can prune it, too.
3. Step back to the 'E' branch. You check its other paths. You see it still leads to "TEA" and "TED," so you **must stop pruning here**.

Deletion in a Trie is this careful, recursive process of removing nodes backward from the end of the word until you hit a node that is shared by another word.

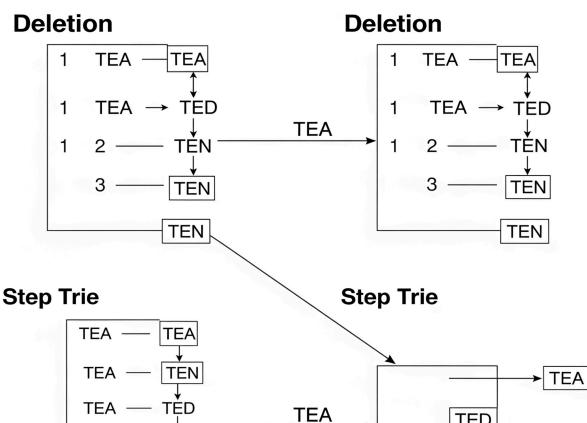
### The Three Cases of Deletion

When deleting a word, a node can only be physically removed if it meets two conditions:

1. It is not marked as the end of another word.
2. It has no other children (it's not a prefix for another word).

This leads to three main scenarios for the word you want to delete:

- **Case 1: It's a prefix of another word.** (e.g., delete "BUG" when "BUGS" exists). You don't remove any nodes. You just find the node for 'G' and set its `isEndOfWord` flag to `false`.
- **Case 2: It shares a prefix with another word.** (e.g., delete "BUG" when "BAT" exists). You recursively delete the nodes for 'G', 'U', and 'B' backward until you reach the 'A' node from "BAT" at the root. The deletion stops there.
- **Case 3: No part of it is shared.** (e.g., delete "HELLO"). You can safely delete all nodes from 'O' all the way back to the root.



## The C++ Implementation

The standard way to implement `delete` is with a recursive helper function that traverses down to the end of the word and then cleans up nodes on its way back up.

```
#include <iostream>
#include <string>

using namespace std;

const int ALPHABET_SIZE = 26;

class Trie {
private:
    struct TrieNode {
        TrieNode* children[ALPHABET_SIZE];
        bool isEndOfWord;
        TrieNode() {
            isEndOfWord = false;
            for (int i = 0; i < ALPHABET_SIZE; ++i) children[i] = nullptr;
        }
    };
    TrieNode* root;
};

// Recursive helper for deletion
// Returns true if the parent node should delete the mapping for this
// child.
bool deleteHelper(TrieNode*& current, const string& word, int depth) {
    if (current == nullptr) {
        return false;
    }

    // Base Case: We've reached the end of the word
    if (depth == word.length()) {
        if (!current->isEndOfWord) return false; // Word doesn't exist
        current->isEndOfWord = false; // Un-mark it as a word
    }

    // Check if this node has any children. If not, it's safe to
    // delete.
    for (int i = 0; i < ALPHABET_SIZE; ++i) {
        if (current->children[i] != nullptr) return false; // Has
        // children, don't delete
    }
    return true; // No children, safe to delete
}
```

```

    }

    // Recursive Step
    int index = word[depth] - 'a';
    if (deleteHelper(current->children[index], word, depth + 1)) {
        // If the recursive call told us it's safe to delete the child
node...
        delete current->children[index];
        current->children[index] = nullptr;

        // Check if this current node is now a candidate for deletion
        // (i.e., it's not an end of another word and has no other
children)
        for (int i = 0; i < ALPHABET_SIZE; ++i) {
            if (current->children[i] != nullptr) return false;
        }
        return !current->isEndOfWord;
    }

    return false;
}

public:
Trie() { root = new TrieNode(); }
void insert(const string& word); // (Assume implemented)
bool search(const string& word); // (Assume implemented)

void remove(const string& word) {
    deleteHelper(root, word, 0);
}
};

// Assume insert and search are filled in from the previous section...
// [Full code from previous section would be here]

// --- Main function to demonstrate ---
// (A full main would be required to test this)
int main() {
    Trie t;
    t.insert("hello");
    t.insert("her");
    t.insert("help");

    cout << "Searching for 'help': " << boolalpha << t.search("help") << endl;
// true
}

```

```
cout << "Deleting 'help'..." << endl;
t.remove("help");

    cout << "Searching for 'help' again: " << t.search("help") << endl; //  
false
    cout << "Searching for 'her' (should still exist): " << t.search("her") <<  
endl; // true

    return 0;
}

// NOTE: The main function needs the full Trie class implementation to run.  
// This example focuses only on the deletion logic.
```

# Chapter 8: Advanced Trie Concepts & Applications

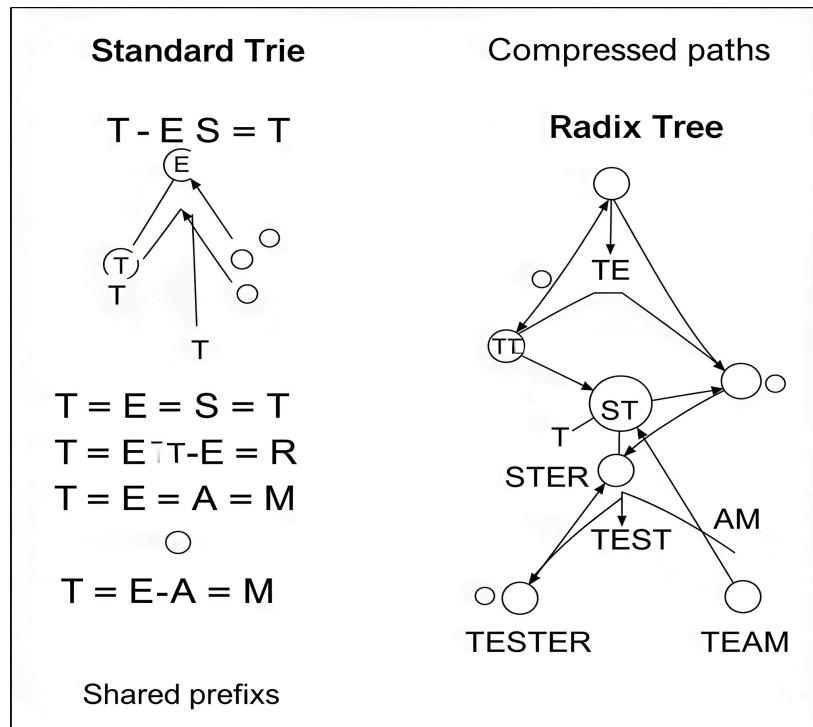
## 8.1 Saving Space: Compressed Tries (Radix Trees)

While a standard Trie is powerful, it can consume a lot of memory, especially if it stores long words with few branches. A **Compressed Trie**, more commonly known as a **Radix Tree**, is a space-optimized variant that solves this problem.

### The "Express Train" Analogy 🚆

- A **standard Trie** is like a **local train** on a subway line. It stops at every single station (character) along the way, even if the track is a long, straight line with no intersections. For a word like "international," it would make 12 stops.
- A **Radix Tree** is like an **express train**. It looks at the track ahead and skips all the unnecessary intermediate stations. It goes directly from one major junction (a node with multiple branches) to the next. Instead of a ticket for one stop, its ticket is for an entire segment of the journey (a substring like "inter" or "nation").

The core idea of a Radix Tree is to **compress chains of single-child nodes**. If a node has only one child, it gets merged with that child, and the edge connecting them stores the combined substring instead of a single character.



## The C++ Node Structure

A full implementation of a Radix Tree is quite complex, as `insert` and `delete` operations must now handle splitting and merging edges. However, we can understand the core difference by looking at its node structure.

Instead of just an array of child pointers, a Radix node needs to store edges, where each edge has a **label (a substring)** and a **pointer to the next node**.

```
#include <iostream>
#include <string>
#include <map>

using namespace std;

// Forward declaration
struct RadixNode;

// An Edge in a Radix Tree contains the substring label
// and a pointer to the destination node.
struct Edge {
    string label;
    RadixNode* targetNode;
};

// The conceptual node for a Radix Tree.
struct RadixNode {
    // We map the *first character* of an edge to the edge itself.
    // This allows for quick lookup of the next path.
    map<char, Edge> edges;

    // This flag is still needed to mark complete words.
    bool isEndOfWord;

    RadixNode() {
        isEndOfWord = false;
    }
};

class RadixTree {
private:
    RadixNode* root;

public:
    RadixTree() {
```

```

        root = new RadixNode();
    }
    // NOTE: The implementation of insert(), search(), and delete() for a
    // Radix Tree is significantly more involved than for a standard Trie,
    // as it requires logic for splitting and merging edges.
    // This example focuses only on the change in node architecture.
};

int main() {
    cout << "Creating a Radix Tree instance..." << endl;
    RadixTree rt;
    cout << "Radix Tree created." << endl;
    cout << "Its node structure stores substrings instead of single characters,
saving space." << endl;

    return 0;
}

```

By storing substrings on the edges, Radix Trees provide the same functionality as Tries but often with a much smaller memory footprint, making them ideal for applications with large sets of strings that share long prefixes, such as IP routing tables.

## 8.2 Application Spotlight 1: Spell Checkers & Dictionaries

Tries are the perfect data structure for building the backbone of digital dictionaries and spell checkers due to their inherent ability to handle prefix-based queries efficiently.

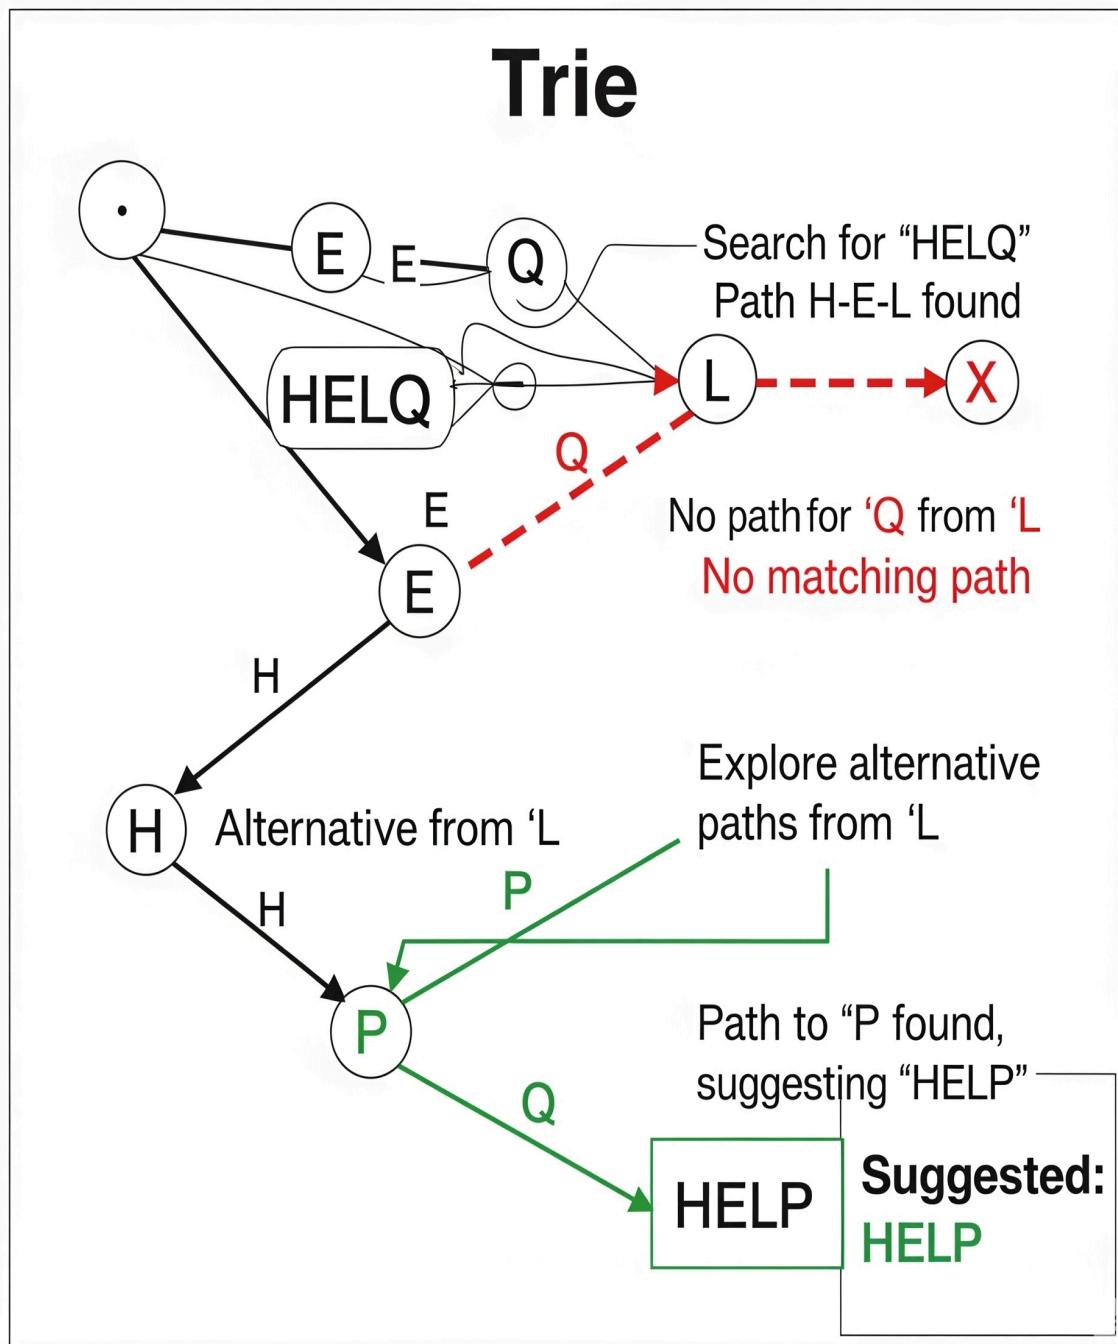
### The "Path in a Dictionary" Analogy

Using a Trie is like navigating a physical dictionary. To check if the word "computer" is spelled correctly, you don't scan every word. You flip to the 'C' section, then find the 'Co' section, then 'Com', and so on. This path-finding is exactly what a Trie does.

Now, imagine you've typed "**compu**". A spell checker using a Trie navigates to the node for this prefix. From there, it can instantly see all possible valid paths forward, allowing it to **suggest** words like "computer" or "computation." If you typed "**compiter**", the path would break after 'p' because there is no 'i' child. The system knows the word is misspelled and can use the last valid node ('p') as a starting point to find nearby valid words.

### How Tries Power These Features

- Dictionary Lookups:** Checking if a word exists is a simple `search()` operation. It's extremely fast ( $O(L)$  where  $L$  is the word's length), confirming if a path exists and ends on a valid word node.
- Autocomplete/Suggestions:** This is the Trie's killer feature. By traversing to the node of a given prefix, we can perform a search (like a Depth-First Search) on the sub-Trie starting from that node to collect all possible valid words.



## C++ Code: Implementing a Suggestion Feature

Let's add a `suggest` method to our `Trie` class. This function will take a prefix and return all words in our dictionary that start with that prefix. This is the core logic behind autocomplete.

```
#include <iostream>
#include <string>
#include <vector>

using namespace std;

// Assuming a complete Trie class from previous sections
// with a TrieNode struct and an insert() method.
const int ALPHABET_SIZE = 26;
class Trie {
private:
    struct TrieNode { /* ... as defined before ... */ };
    TrieNode* root;

    // Recursive helper to collect all words from a given node
    void suggestHelper(TrieNode* node, string currentPrefix, vector<string>& results) {
        if (node == nullptr) return;

        // If the current node marks the end of a word, add it to results
        if (node->isEndOfWord) {
            results.push_back(currentPrefix);
        }

        // Recurse for all children that exist
        for (int i = 0; i < ALPHABET_SIZE; ++i) {
            if (node->children[i] != nullptr) {
                suggestHelper(node->children[i], currentPrefix + (char)(‘a’ +
i), results);
            }
        }
    }

public:
    Trie() { root = new TrieNode(); }
    void insert(const string& word); // Assume implemented

    // Public method to get suggestions for a prefix
    vector<string> suggest(const string& prefix) {
        vector<string> results;
```

```

TrieNode* current = root;

// 1. Traverse to the node for the end of the prefix
for (char ch : prefix) {
    int index = ch - 'a';
    if (current->children[index] == nullptr) {
        return results; // Prefix does not exist, no suggestions
    }
    current = current->children[index];
}

// 2. Now, collect all words from this prefix node
suggestHelper(current, prefix, results);
return results;
}

};

// --- Minimal implementations for demonstration ---
void Trie::insert(const string& word) {
    TrieNode* current = root;
    for (char ch : word) {
        int index = ch - 'a';
        if (current->children[index] == nullptr) {
            current->children[index] = new TrieNode();
        }
        current = current->children[index];
    }
    current->isEndOfWord = true;
}

int main() {
    Trie dictionary;
    dictionary.insert("cat");
    dictionary.insert("car");
    dictionary.insert("cart");
    dictionary.insert("care");
    dictionary.insert("dog");

    cout << "Suggestions for prefix 'car':" << endl;
    vector<string> suggestions = dictionary.suggest("car");
    for (const string& word : suggestions) {
        cout << " - " << word << endl;
    }
    // Expected output:
}

```

```

// - car
// - care
// - cart

cout << "\nSuggestions for prefix 'd':\n" << endl;
suggestions = dictionary.suggest("d");
for (const string& word : suggestions) {
    cout << " - " << word << endl;
}
// Expected output:
// - dog

return 0;
}

```

## 8.3 Application Spotlight 2: IP Routing Tables

One of the most critical, high-performance applications of Tries is inside the internet routers that power our digital world. They are the key to making internet traffic fast and efficient.

### The Router's Dilemma

When a router receives a data packet (part of an email, a video stream, etc.), it reads the destination IP address. It then has to make a split-second decision: which of its outgoing network connections should it send this packet to? It does this by looking up the IP address in its **routing table**.

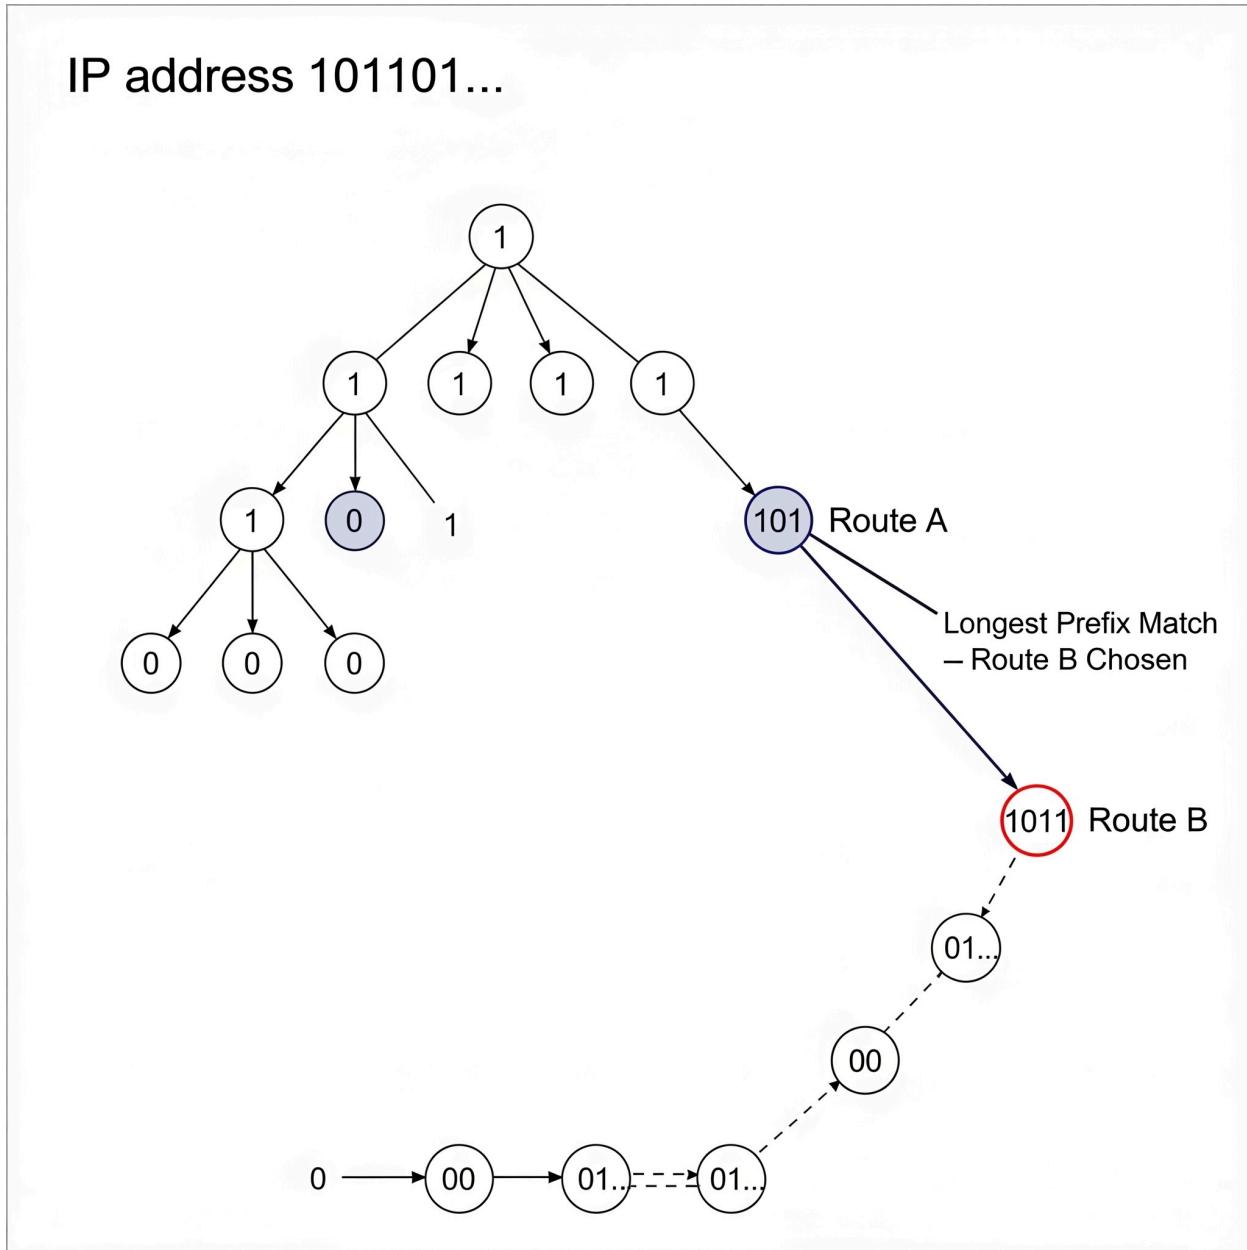
However, routers don't store an entry for every single IP address in the world. Instead, they store **network prefixes**. For example, a route might be for `192.168.1.0/24`, which matches all IPs from `192.168.1.0` to `192.168.1.255`. A single IP might match multiple prefixes, so routers follow one crucial rule: always use the **most specific route**, which is the one with the **longest prefix match**.

### The "Postal ZIP Code" Analogy

Think of how the postal service routes a letter.

- They don't have a unique bin for every single house address in the country.
- First, they look at a general prefix: the ZIP code (e.g., 110001 for Delhi).
- Then, within that post office, they might look at a more specific prefix for the neighborhood or street.
- They always use the most detailed (longest) prefix they have to route the letter correctly.

A Trie (specifically, a space-optimized Radix Tree) allows a router to do this for IP addresses, which are just long strings of bits (0s and 1s). The search for the longest matching prefix becomes a simple, fast traversal down the tree.



### C++ Code: Simulating Longest Prefix Match

This simplified C++ example demonstrates the core logic. We'll use a binary Trie where each node can have `routInfo`. The `getRoute` function traverses the Trie, keeping track of the best (longest) match it has seen so far.

```
#include <iostream>
```

```

#include <string>
#include <vector>

using namespace std;

// A simplified Trie for IP routing simulation
class RoutingTrie {
private:
    struct TrieNode {
        TrieNode* children[2]; // 0 and 1
        string routeInfo;      // e.g., "Send to Port A", "Interface 2", etc.
        TrieNode() {
            routeInfo = ""; // Empty string means it's not a valid route
        }
    };
    TrieNode* root;
};

public:
    RoutingTrie() {
        root = new TrieNode();
        root->routeInfo = "Default Gateway"; // A route for 0.0.0.0/0
    }

    void addRoute(const string& prefix, const string& info) {
        TrieNode* current = root;
        for (char bit : prefix) {
            int index = bit - '0';
            if (current->children[index] == nullptr) {
                current->children[index] = new TrieNode();
            }
            current = current->children[index];
        }
        current->routeInfo = info;
    }

    string getRoute(const string& ip_address) {
        TrieNode* current = root;
        string longest_match_route = root->routeInfo;

        for (char bit : ip_address) {
            int index = bit - '0';
            if (current->children[index] == nullptr) {
                // No more specific path, break the loop

```

```

        break;
    }
    current = current->children[index];
    if (!current->routeInfo.empty()) {
        // This is a valid route, update our best match so far
        longest_match_route = current->routeInfo;
    }
}
return longest_match_route;
};

int main() {
    RoutingTrie router;
    router.addRoute("101", "Interface A");
    router.addRoute("1011", "Interface B (More Specific)");
    router.addRoute("11", "Interface C");

    string ip1 = "1011010100101010"; // Should match Interface B
    string ip2 = "1010111010101010"; // Should match Interface A
    string ip3 = "1111000011110000"; // Should match Interface C
    string ip4 = "0101010101010101"; // Should use the default route

    cout << "Route for " << ip1 << ":" << router.getRoute(ip1) << endl;
    cout << "Route for " << ip2 << ":" << router.getRoute(ip2) << endl;
    cout << "Route for " << ip3 << ":" << router.getRoute(ip3) << endl;
    cout << "Route for " << ip4 << ":" << router.getRoute(ip4) << endl;

    return 0;
}

```

## 8.4 Advanced Variant: Ternary Search Trees (A Brief Overview)

A Ternary Search Tree (TST) is a clever hybrid data structure that combines the features of a **Binary Search Tree (BST)** and a **Trie**. It offers the time-efficient searching of a Trie while being much more space-efficient, especially for large character sets.

### The "20 Questions" Analogy 🤔

- A standard **Trie** is like a multiple-choice quiz with 26 options at each step: "Is the next letter A? B? C?..." This can be a lot of empty, wasted space if only a few options are ever used.

- A **Ternary Search Tree** is like playing the game "**20 Questions.**" At each node, you only ask three smart questions about the current character:
  1. Is your character **less than** my character? (Go left).
  2. Is your character **greater than** my character? (Go right).
  3. Is your character **equal to** my character? (Go down to the next level for the next character in the word).

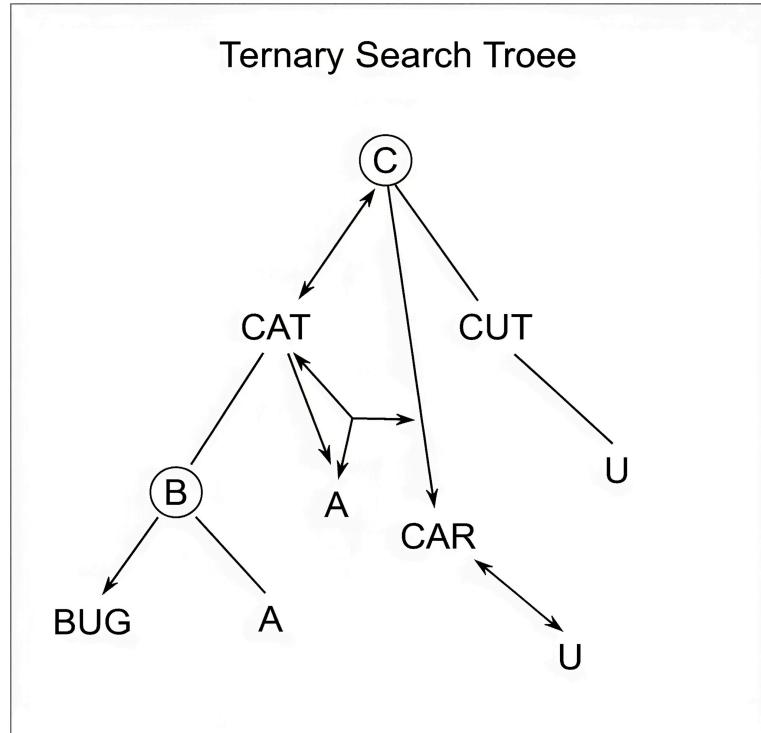
This three-way branching (`<`, `=`, `>`) is much more compact than the 26-way branching of a standard Trie.

## The Structure of a TST Node

Unlike a Trie node with its large array of children, a TST node is simpler and more like a BST node. It contains:

1. A **character** to compare against.
2. **Three pointers**: `left`, `middle` (or `equal`), and `right`.
3. An `isEndOfWord` flag.

The `left` and `right` pointers behave like a BST, referring to other characters at the *same position* in the word. The `middle` pointer behaves like a Trie, advancing to the *next position* in the word once a character has been matched.



## C++ Code: The TSTNode

A full TST implementation can be complex, but its elegance is clear from its node structure. It's a space-saving alternative to a standard Trie, especially when you're not dealing with just 26 lowercase letters but a much larger character set like ASCII or Unicode.

```
#include <iostream>
#include <string>

// The fundamental node for a Ternary Search Tree (TST).
struct TSTNode {
    char data; // The character stored at this node

    // The three pointers that give the TST its name
    TSTNode *left; // Pointer to a node with a smaller character value
    TSTNode *middle; // Pointer to the next character in the word
    TSTNode *right; // Pointer to a node with a larger character value

    // True if this node represents the end of a word
    bool isEndOfWord;

    // Constructor for a new node
    TSTNode(char c) {
        data = c;
        isEndOfWord = false;
        left = middle = right = nullptr;
    }
};

class TernarySearchTree {
private:
    TSTNode* root;

public:
    TernarySearchTree() : root(nullptr) {}

    // NOTE: The implementation of insert() and search() for a TST
    // involves a recursive logic similar to a BST, but with the
    // third "middle" path for matching characters. It is more
    // involved than a standard Trie's implementation.
    // This example focuses on the node architecture.
};

int main() {
```

```
    std::cout << "Creating a Ternary Search Tree instance..." << std::endl;
    TernarySearchTree tst;
    std::cout << "TST created." << std::endl;
    std::cout << "Its node structure is a space-efficient hybrid of a BST and a
Trie." << std::endl;

    return 0;
}
```

# Chapter 9: Trie Problem-Solving Workshop

Let's apply our Trie knowledge to solve a couple of challenging, real-world problems.

## 9.1 Problem 1: Add and Search Word

**The Challenge:** Design a data structure that can `addWord(word)` and `search(word)`. The twist: the `search` can contain a wildcard character `.` that matches any single letter.

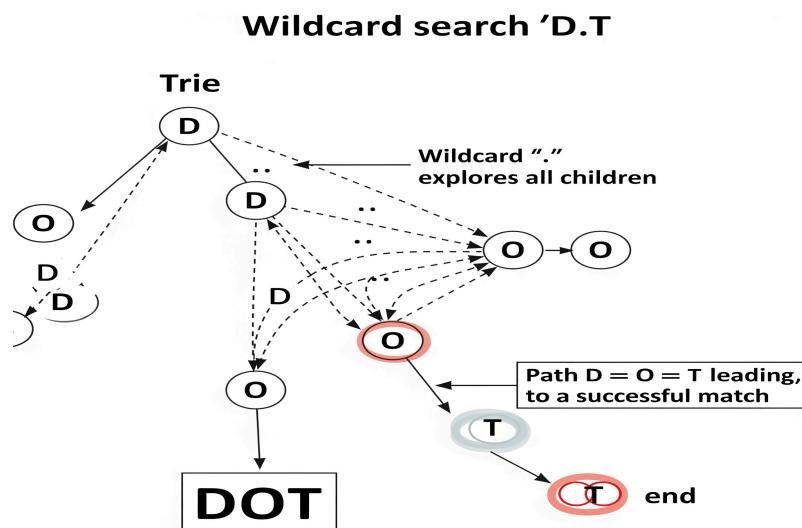
The "Master Key" Analogy 🔑

Think of searching the Trie as navigating a hotel. A normal character like 'c' is a specific key that only opens the 'c' door. The wildcard `.` is a **master key**. When you encounter a `.` in your search term (e.g., "h.t"), you must try the master key on **every single door** on that floor to see if any of them lead to a path that ultimately spells your word. This means we need to explore multiple paths.

### The Algorithm

The `addWord` function is a standard Trie `insert`. The `search` function, however, needs to be a recursive helper that handles the wildcard.

1. **Start a recursive search:** `searchHelper(word, index, node)`.
2. If the current character is a normal letter, move to the corresponding child and recurse. If the child doesn't exist, the path fails.
3. If the current character is a `.` (the master key), you must **iterate through all 26 children** of the current node. For each child that exists, recursively call the search function. If **any** of these recursive calls return `true`, you've found a match.



## C++ Implementation

```
#include <iostream>
#include <string>
#include <vector>

using namespace std;

class WordDictionary {
private:
    struct TrieNode {
        TrieNode* children[26];
        bool isEndOfWord;
        TrieNode() {
            isEndOfWord = false;
            for (int i = 0; i < 26; ++i) children[i] = nullptr;
        }
    };
    TrieNode* root;
    bool searchInNode(const string& word, int index, TrieNode* node) {
        if (node == nullptr) return false;
        if (index == word.length()) return node->isEndOfWord;

        char ch = word[index];
        if (ch == '.') {
            // Wildcard: try every possible child path
            for (int i = 0; i < 26; ++i) {
                if (searchInNode(word, index + 1, node->children[i])) {
                    return true;
                }
            }
            return false;
        } else {
            // Normal character: follow the single path
            return searchInNode(word, index + 1, node->children[ch - 'a']);
        }
    }

public:
    WordDictionary() { root = new TrieNode(); }

    void addWord(const string& word) {
        TrieNode* current = root;
        for (char ch : word) {
```

```

        int i = ch - 'a';
        if (current->children[i] == nullptr) {
            current->children[i] = new TrieNode();
        }
        current = current->children[i];
    }

    current->isEndOfWord = true;
}

bool search(const string& word) {
    return searchInNode(word, 0, root);
}
};

int main() {
    WordDictionary dict;
    dict.addWord("bad");
    dict.addWord("dad");
    dict.addWord("mad");

    cout << boolalpha;
    cout << "Search 'pad': " << dict.search("pad") << endl; // false
    cout << "Search 'bad': " << dict.search("bad") << endl; // true
    cout << "Search '.ad': " << dict.search(".ad") << endl; // true
    cout << "Search 'b..': " << dict.search("b..") << endl; // true
}

```

## 9.2 Problem 2: Word Search II

**The Challenge:** Given a 2D grid of characters (a "board") and a list of words, find all words from the list that can be formed by adjacent cells (horizontally or vertically), without using the same cell twice in a word.

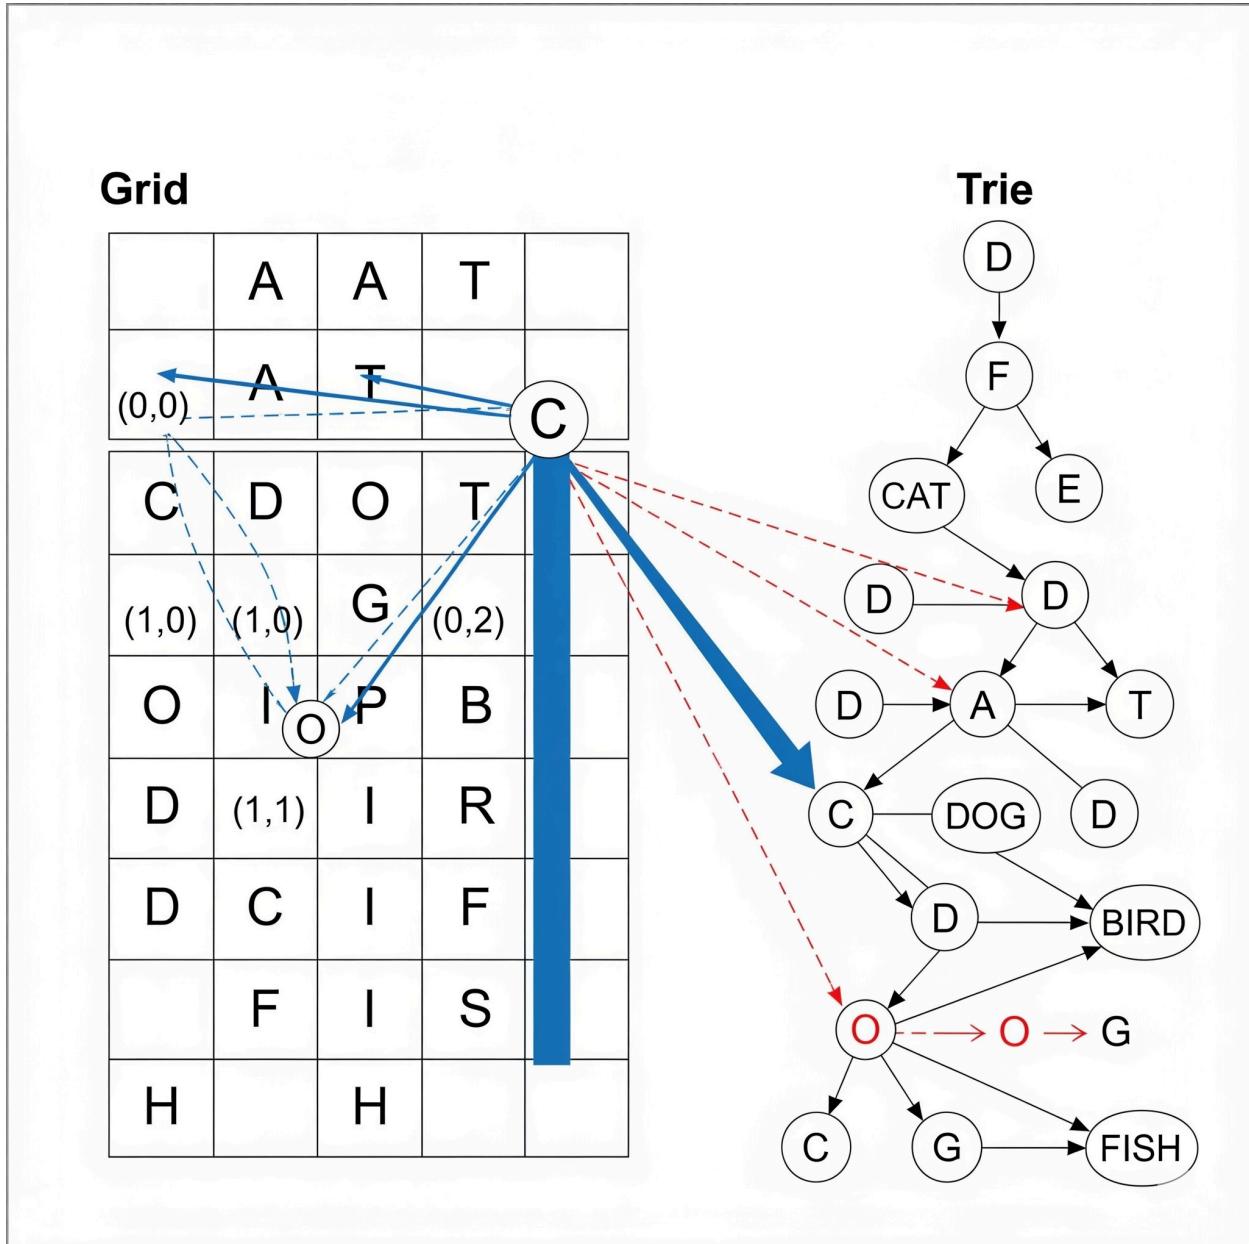
### The "Boggle" Analogy

This is the classic game of Boggle. The naive way is to pick a word from the dictionary and scan the entire grid for it. This is very slow. The smart way is to use the grid to guide a search through a Trie built from the dictionary.

You start at a letter on the grid, say 'A'. This is your first step into the Trie. Then, you look at the adjacent letters on the grid. If the 'A' node in your Trie has a child for an adjacent letter, you take that step on both the grid and in the Trie. You continue this "parallel walk." When you land on a Trie node that marks the end of a word, you've found a match!

The Algorithm

1. **Build a Trie** from all the words in the dictionary. This is a one-time setup cost.
2. Iterate through every cell  $(i, j)$  on the board, using each one as a potential starting point.
3. From each cell, start a **Depth-First Search (DFS)**.
4. The DFS function will explore the grid while simultaneously traversing the Trie.
5. To avoid cycles, mark cells as "visited" during a path and "un-visit" them when backtracking.
6. When the traversal lands on a Trie node where `isEndOfWord` is true, add the found word to your results.



## C++ Implementation

```
#include <iostream>
#include <vector>
#include <string>

using namespace std;

// TrieNode and Trie setup (can be defined inside the Solution class)
struct TrieNode {
    TrieNode* children[26];
    string word;
    TrieNode() : word("") {
        for (int i=0; i<26; ++i) children[i] = nullptr;
    }
};

void insert(TrieNode* root, const string& word) {
    TrieNode* current = root;
    for (char ch : word) {
        int i = ch - 'a';
        if (current->children[i] == nullptr) current->children[i] = new TrieNode();
        current = current->children[i];
    }
    current->word = word;
}

class Solution {
public:
    vector<string> findWords(vector<vector<char>>& board, vector<string>& words) {
        TrieNode* root = new TrieNode();
        for (const string& word : words) {
            insert(root, word);
        }

        vector<string> result;
        for (int i = 0; i < board.size(); ++i) {
            for (int j = 0; j < board[0].size(); ++j) {
                dfs(board, i, j, root, result);
            }
        }
    }
}
```

```

        return result;
    }

private:
    void dfs(vector<vector<char>>& board, int i, int j, TrieNode* node,
vector<string>& result) {
    // Boundary checks or if path is invalid
    if (i < 0 || i >= board.size() || j < 0 || j >= board[0].size() ||
        board[i][j] == '#' || node->children[board[i][j] - 'a'] ==
nullptr) {
        return;
    }

    char c = board[i][j];
    node = node->children[c - 'a'];

    // If we found a word
    if (!node->word.empty()) {
        result.push_back(node->word);
        node->word = ""; // De-duplicate by clearing the found word
    }

    board[i][j] = '#'; // Mark as visited

    // Explore neighbors
    dfs(board, i + 1, j, node, result);
    dfs(board, i - 1, j, node, result);
    dfs(board, i, j + 1, node, result);
    dfs(board, i, j - 1, node, result);

    board[i][j] = c; // Backtrack
}
};

int main() {
    Solution s;
    vector<vector<char>> board = {
        {'o', 'a', 'a', 'n'},
        {'e', 't', 'a', 'e'},
        {'i', 'h', 'k', 'r'},
        {'i', 'f', 'l', 'v'}
    };
    vector<string> words = {"oath", "pea", "eat", "rain"};
}

```

```
vector<string> found = s.findWords(board, words);

cout << "Found words: ";
for (const string& w : found) {
    cout << w << " "; // Expected: oath eat
}
cout << endl;
}
```

# Chapter 10: Conclusion

## 10.1 Head-to-Head: Heap vs. Trie Comparison Table

We've explored two very different but equally powerful data structures. This table summarizes their key characteristics at a glance.

Feature	Heap (The Priority Organizer 🏆)	Trie (The String Specialist abc)
<b>Primary Use Case</b>	Managing a collection based on priority. Always needing the min/max element.	Storing and searching for strings based on prefixes.
<b>Analogy</b>	An emergency room triage or a tournament bracket.	A dictionary index or a search bar's autocomplete.
<b>Data Type</b>	Numbers or any comparable objects.	Primarily strings or sequences.
<b>Key Operations</b>	<code>insert</code> , <code>extractMin/Max</code> , <code>peek</code>	<code>insert</code> , <code>search</code> , <code>startsWith</code>
<b>Time Complexity</b>	$O(\log n)$ for insertions/deletions. $O(1)$ for <code>peek</code> .	$O(L)$ for all core operations, where $L$ is the length of the string.
<b>Space Complexity</b>	Excellent. $O(n)$ for the array.	Can be high if prefixes are not shared, but very efficient otherwise.

## 10.2 Key Takeaways & Final Thoughts

And with that, our journey through Heaps and Tries comes to an end. My goal for this tutorial was to demystify these structures, moving them from abstract theory into the realm of practical, problem-solving tools. When I started, they just seemed like complex topics for interviews. But after building them from scratch, I see them everywhere—from the way my GPS finds the fastest route to how my phone suggests the next word I'm about to type.

The biggest takeaway for me is that there's no "one size fits all" data structure. The real skill is looking at a problem and knowing which tool from your toolbox will get the job done most efficiently. A Heap is your go-to for anything involving priority, while a Trie is the undisputed champion of prefix-based string operations.

I hope this guide has been clear and helpful. Building these structures, especially the tricky `delete` operation in a Trie and the logic for finding a median, was a challenging but incredibly rewarding experience. It's one thing to read about them, but writing the code yourself truly solidifies the concepts.

Thank you for following along, and best of luck to everyone in the second round of the DSA selection process!

## 10.3 Further Reading & Resources

Throughout my research for this tutorial, I found the following resources to be incredibly helpful for clear explanations, code examples, and practice problems. I highly recommend them for anyone looking to deepen their understanding.

1. **GeeksforGeeks:** An excellent resource for detailed articles on nearly every data structure, including step-by-step implementations and complexity analysis.
  - o *Heap Data Structure:* <https://www.geeksforgeeks.org/heap-data-structure/>
  - o *Trie Data Structure:* <https://www.geeksforgeeks.org/trie-insert-and-search/>
2. **LeetCode:** The best platform for hands-on practice. The problems I included in the workshop chapters are classic LeetCode challenges that test these concepts perfectly.
  - o *Problem: Find Median from Data Stream*
  - o *Problem: Merge K Sorted Lists*
  - o *Problem: Implement Trie (Prefix Tree)*
  - o *Problem: Word Search II*
3. **"Introduction to Algorithms" (CLRS) by Cormen, Leiserson, Rivest, and Stein:** For those who want a deep, formal, and mathematical understanding, this is the definitive textbook on algorithms and data structures.
4. **cplusplus.com:** An indispensable reference for understanding the C++ Standard Library implementations, such as `std::priority_queue`.