

A Comprehensive Guide to Tree Data Structures

When I first started learning about Tree Data Structures, I found it easy to get lost in the academic definitions. The concepts only truly 'clicked' when I could see how they were used to solve a real problem. That's the approach I've taken here. My goal with this tutorial is to ground every topic in a practical example first, and then immediately dive into the C++ code to bring it to life. This way, we're not just memorizing theory, but building an intuitive understanding of how to actually use these powerful tools.

Table of contents

Chapter 1: Foundations of Tree Data Structures.....	6
1.1. What is a Tree? (Core Concepts & Terminology).....	6
1.2. Why Use Trees? A Comparison with Linear Structures.....	8
1.3. A Map of the Forest: An Overview of Major Tree Types.....	9
1.4. Common Real-World Applications of Trees.....	10
Chapter 2: The Binary Tree: A Fundamental Building Block.....	11
2.1. Anatomy and Properties of a Binary Tree.....	11
Key Properties (Why this Structure is Efficient).....	11
Diagram: The Structure of a Binary Tree.....	11
Code: The Binary Tree Node in C++.....	12
2.2. Types of Binary Trees (Full, Complete, Perfect, Degenerate).....	12
Full Binary Tree.....	12
Complete Binary Tree.....	12
Perfect Binary Tree.....	13
Degenerate (or Pathological) Tree.....	13
Diagram: Comparing Tree Shapes.....	13
2.3. Core Traversal Algorithms (Theory & Code).....	13
2.3.1. Depth-First: In-order, Pre-order, Post-order.....	14
In-order Traversal (Left, Root, Right).....	14
Pre-order Traversal (Root, Left, Right).....	14
Post-order Traversal (Left, Right, Root).....	14
2.3.2. Breadth-First: Level-order.....	15
Level-order Traversal.....	15
Detailed Diagram Walkthrough.....	16
In-order: (Left, Root, Right) -> 4, 2, 5, 1, 3.....	16
Pre-order: (Root, Left, Right) -> 1, 2, 4, 5, 3.....	16
Post-order: (Left, Right, Root) -> 4, 5, 2, 3, 1.....	17
Level-order -> 1, 2, 3, 4, 5.....	17
2.4. Essential Operations (Theory & C++ Implementation).....	17
2.4.1. Search, Insertion, and Deletion.....	17
Searching for a Node 	17
Inserting a Node 	18
Deleting a Node 	19
Diagram: Visualizing Tree Operations.....	21
Chapter 3: The Binary Search Tree (BST): Bringing Order to Data.....	23
3.1. The BST Invariant: The Golden Rule for Ordering.....	23
Diagram: Valid vs. Invalid BSTs.....	23

Code: Validating the BST Invariant.....	24
3.2. Operations in a BST (Search, Insert, Delete with C++ examples).....	24
Searching in a BST 	24
Inserting into a BST 	25
Deleting from a BST 	26
Diagram: Operations in a Binary Search Tree.....	27
3.3. The Problem of Imbalance: When BSTs Fail.....	28
Diagram: Balanced vs. Unbalanced BST.....	28
Chapter 4: Self-Balancing Trees: Guaranteed Performance 	30
4.1. The AVL Tree.....	30
4.1.1. Balancing Factors and Rotations (Theory & Code).....	30
The Balance Factor.....	30
Rotations: The Fix.....	30
Diagram: The Four Rotation Cases.....	30
Code: AVL Node and Rotation Functions.....	31
4.1.2. Insertion and Deletion Walkthrough.....	32
Insertion.....	32
Diagram: AVL Insertion Walkthrough.....	33
Deletion.....	33
Code: AVL Insertion and Deletion.....	33
4.2. The Red-Black Tree.....	34
4.2.1. Core Principles and Coloring Rules.....	34
Diagram: A Valid Red-Black Tree.....	35
Code: The Red-Black Tree Node.....	35
4.2.2. Conceptual Insertion and Deletion.....	35
Conceptual Insertion.....	36
Conceptual Deletion.....	36
Elaborated Diagram: The RBT Insertion Fix-up.....	36
Code: Conceptual Functions.....	37
Chapter 5: Specialized Trees for Specific Tasks 	41
5.1. N-ary Trees (Generic Trees) for General Hierarchies.....	41
Diagram: N-ary Tree as a File System.....	41
Code: N-ary Node and Traversal.....	41
5.2. Tries (Prefix Trees) for String Searching & Autocomplete.....	43
Diagram: Building a Trie.....	43
Code: Trie Node and Core Operations.....	43
5.3. B-Trees & B+ Trees for Databases and Filesystems.....	45
The B-Tree.....	45
The B+ Tree: A Popular Enhancement.....	46
Diagram: B-Tree vs. B+ Tree Structure.....	46
Code : Searching an Element in a B-Tree.....	46

5.4. Binary Heaps & Priority Queues.....	47
The Two Heap Properties.....	47
The Array Implementation.....	48
Diagram: Heap Structure and Array Mapping.....	48
Core Operations.....	48
Insertion.....	48
Extracting the Maximum.....	48
Diagram: Heap Operations Walkthrough.....	49
Code: Heap Operations in C++ (using std::vector).....	49
5.5. A Brief Look at Other Variants (Ternary, Interval, etc.).....	50
Ternary Tree.....	50
Diagram: A Simple Ternary Search Tree.....	50
Code: Ternary Search Tree Node.....	50
Interval Tree.....	51
Diagram: Conceptual Interval Tree.....	51
Code: Interval Tree Node.....	51
k-d Tree (k-dimensional Tree).....	52
Diagram: A 2-D Tree Partitioning Space.....	52
Code: k-d Tree Node.....	52
Chapter 6: Advanced Tree-Based Structures 	53
6.1. Disjoint Set Union (DSU) / Union-Find.....	53
Core Operations & Optimizations.....	53
Diagram: DSU Operations and Optimizations.....	54
Code: DSU in C++ with Optimizations.....	54
6.2. Segment Trees for Efficient Range Queries.....	55
Structure and Build Process.....	55
Diagram: Building a Segment Tree.....	56
Querying and Updating.....	56
Diagram: Range Query Walkthrough.....	56
Code: Segment Tree in C++.....	56
6.3. Fenwick Trees (Binary Indexed Trees).....	58
The Core Mechanism: Bit Manipulation.....	59
Diagram: Conceptual Structure and Paths.....	59
Code: Fenwick Tree in C++.....	60
Chapter 7: Appendix: DSA Practice Problems 	62
7.1. Curated Problems on Trees.....	62
Easy Problems.....	62
Medium Problems.....	62
Hard Problems.....	63
7.2. Full C++ Code Listings.....	63
Binary Search Tree.....	63

AVL Tree.....	65
Trie (Prefix Tree).....	67
Binary Heap (Max-Heap).....	69
Disjoint Set Union (DSU).....	70
Conclusion.....	72
References & Further Reading.....	72
Core Textbooks.....	72
Online Resources & Competitive Programming.....	72
Academic Courses.....	73

Chapter 1: Foundations of Tree Data Structures

Imagine you're building a massive digital library for every book ever written. Your first thought might be to just put them all in one giant, sorted list, like a single bookshelf stretching for miles.

When you want to find a specific book, this works okay. You can use a binary search—opening to the middle, checking if your book is before or after, and repeating. But what happens when a new book is published? To add it to your sorted list, you have to shift every single book that comes after it just to make space. For a library with billions of books, this is impossibly slow.

This is the fundamental problem with **linear data structures** like arrays. They are simple, but they struggle to be efficient for both searching *and* modifying data.

This is where **trees** come in. Instead of a single, flat list, a tree organizes information in a **hierarchy**. Think of it like a real library's filing system. You don't start by looking at individual books. You start with a broad category, like "Fiction" (the root). From there, you go to a sub-category like "Science Fiction," then to an author's name, and finally to the book itself.

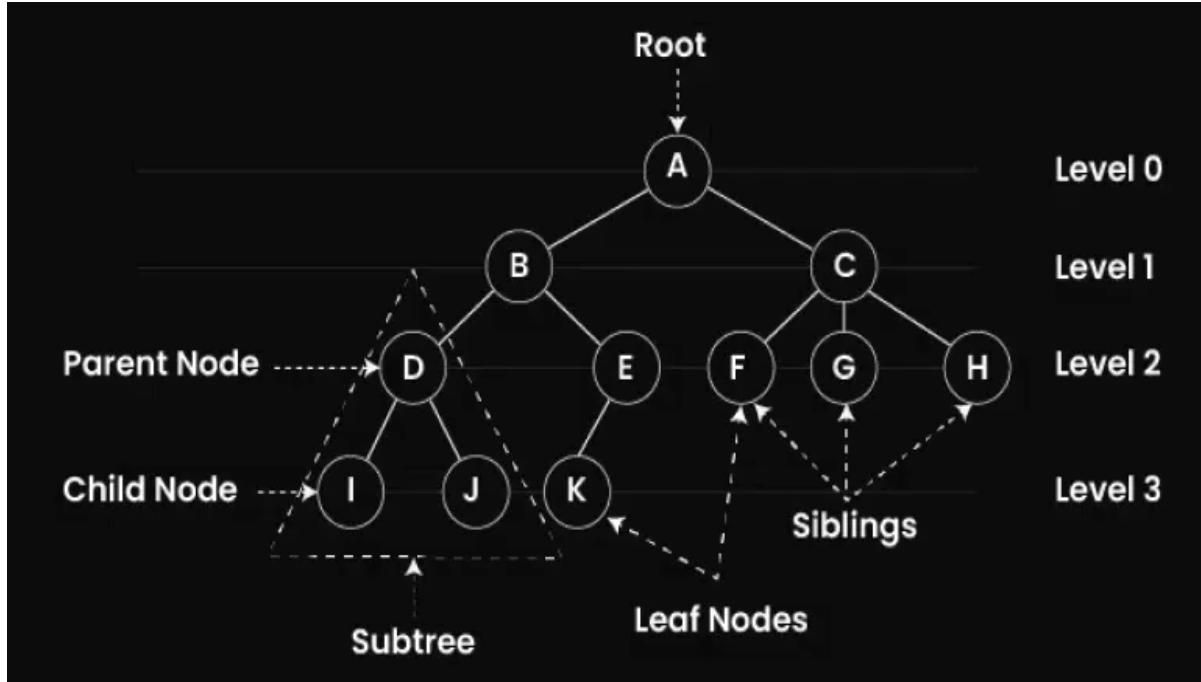
This structure is incredibly efficient. Finding a book is fast because you skip entire sections at each step. Adding a new book is also fast because you just need to find its correct spot in the hierarchy and link it in, without rearranging everything else. Trees give us this power to organize and access vast amounts of information efficiently.

1.1. What is a Tree? (Core Concepts & Terminology)

Now, let's formalize that idea. In computer science, a tree is a hierarchical data structure made of **nodes** connected by **edges**. It's used to represent relationships where one item is "above" or "contains" others. Every tree has a starting point and a clear structure defined by the following terms:

- **Node:** The basic unit of a tree. It contains some data and may link to other nodes. In our library example, a "Node" could be a category, an author, or a single book.
- **Root:** The single, topmost node of the tree. It's the only node with no parent. For our library, this would be the "Main Collection."
- **Edge:** The link or connection between two nodes. It represents the relationship between them (e.g., the edge connecting "Fiction" to "Science Fiction").
- **Parent:** A node that has other nodes branching from it. "Fiction" is the parent of "Science Fiction."
- **Child:** A node that branches from a parent. "Science Fiction" is a child of "Fiction."
- **Siblings:** Nodes that share the same parent. "Science Fiction" and "Fantasy" would be siblings under the "Fiction" parent.
- **Leaf:** A node with no children. These are the endpoints of the tree. In our analogy, the individual books would be leaf nodes.
- **Height:** The length of the longest path from a node to a leaf. The tree's height is the height of its root.

- **Depth:** The path length from the root to a specific node. The root's depth is 0.



Code: To represent these concepts in code, we start by defining the `Node` itself. Here's a basic C++ structure that we'll build upon in later chapters. It holds data and a list of pointers to its children, perfectly capturing the hierarchical nature of a tree.

```
#include <vector>
#include <iostream>

// A generic Node for a tree where a parent can have multiple children.
struct Node {
    int data;
    std::vector<Node*> children;

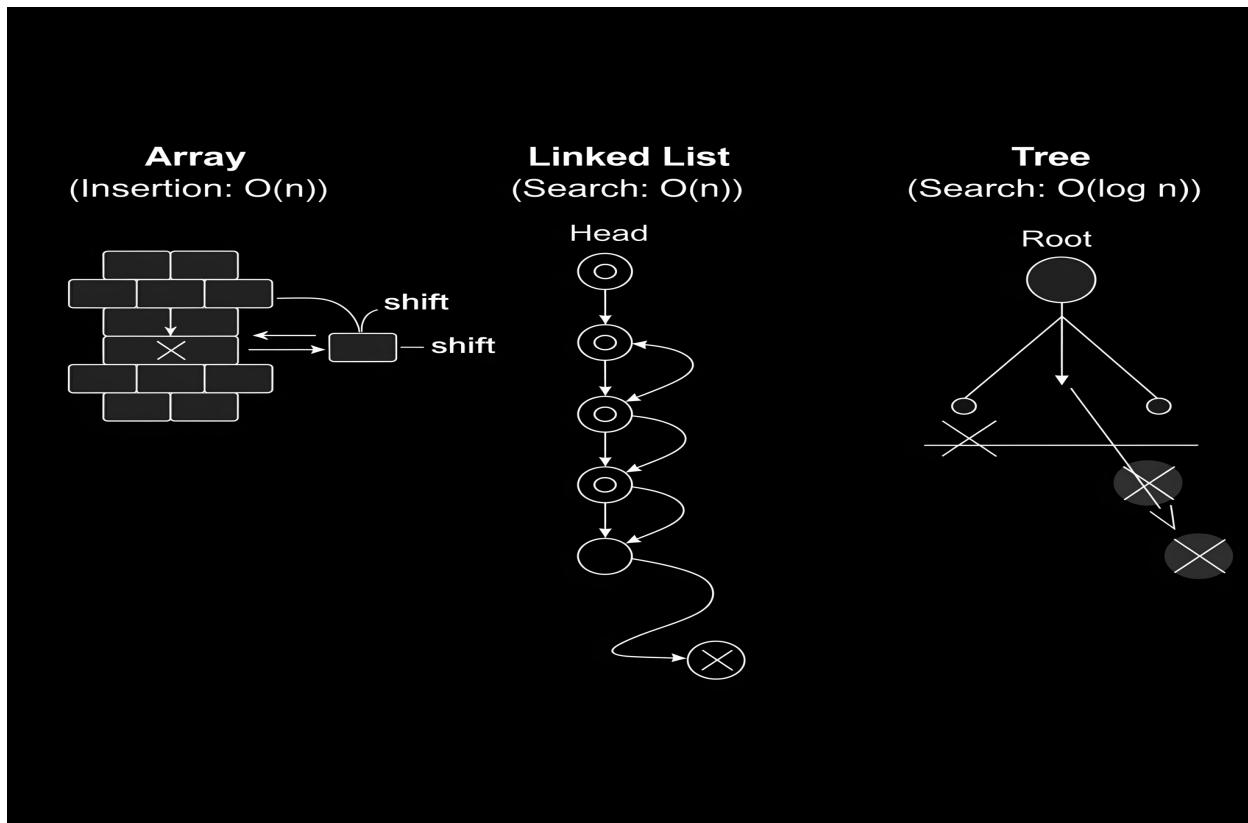
// A simple constructor to create a node with a given value.
Node(int val) {
    data = val;
}
};
```

1.2. Why Use Trees? A Comparison with Linear Structures

The main reason we need trees is to overcome the fundamental limitations of **linear** data structures. While structures like arrays and linked lists are simple, they force us into a trade-off between finding things quickly and changing things quickly.

- **Sorted Arrays** are like a physical dictionary. They are excellent for searching because you can use binary search—opening to the middle to see if your word is before or after—to find any entry in $O(\log n)$ time. However, their rigid structure makes them slow for modifications. Adding a new word would require magically making space on the page, forcing every subsequent entry to be shifted. This makes insertions and deletions a very slow $O(n)$ operation.
- **Linked Lists** are like a treasure hunt where each clue points to the next. They are very flexible for modifications. If you want to add a new clue, you just change the previous clue to point to your new one—a fast $O(1)$ operation if you know where it goes. But finding a specific clue is a nightmare; you have no choice but to start at the beginning and follow every single clue in order, making search a slow $O(n)$ process.

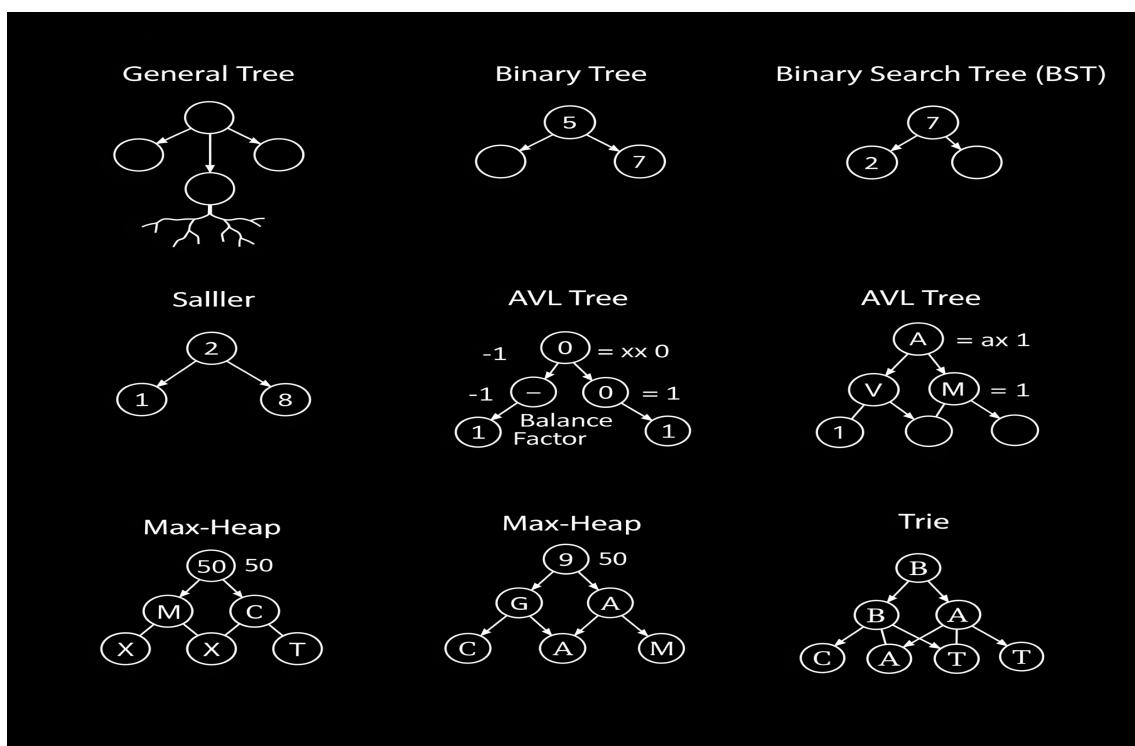
Trees, specifically balanced search trees, are the solution that gives us the best of both worlds. They are like a library's multi-level filing system. By organizing data hierarchically, they allow for efficient searching, insertion, *and* deletion, with all three operations typically averaging $O(\log n)$ time. You can navigate quickly to the right spot, and adding new information doesn't require reorganizing the entire library.



1.3. A Map of the Forest: An Overview of Major Tree Types

Not all trees are the same. Different problems require different types of hierarchical structures. Think of this as a quick map to the "forest" of trees we'll be exploring. Each one has a specific strength.

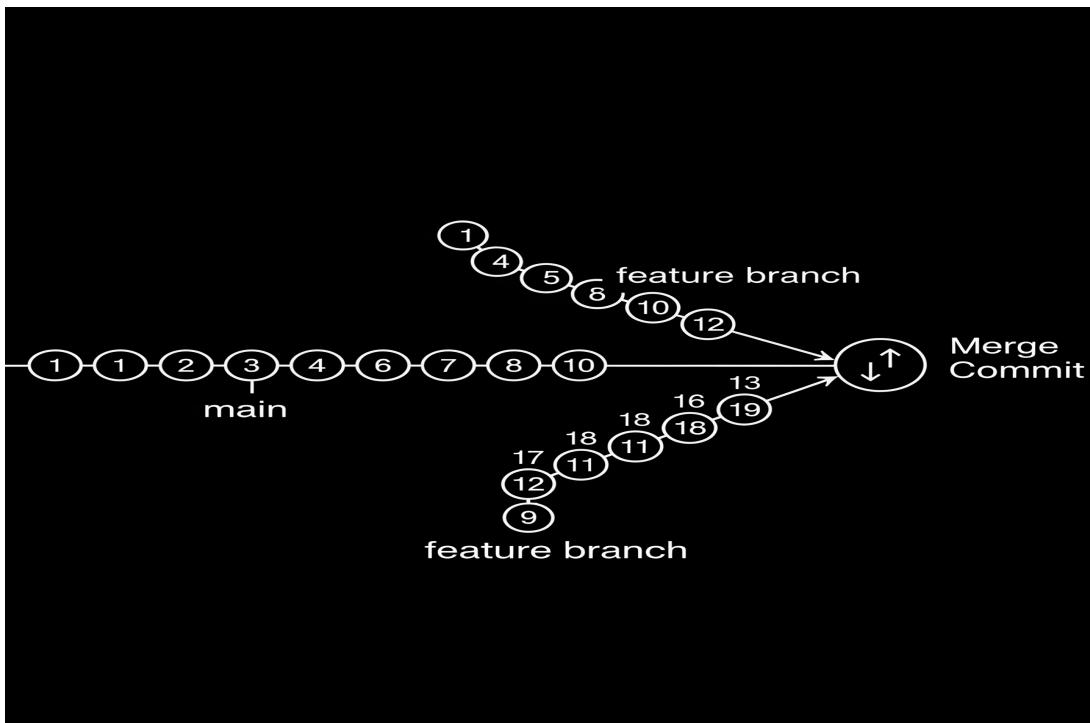
- **General (N-ary) Tree:** The most flexible type, where a node can have **any number of children**. It's perfect for representing things like a computer's file system, where a folder can contain an arbitrary number of files and subfolders.
- **Binary Tree:** A more constrained and common tree where each node can have at most **two children**: a left child and a right child. This simple structure is the foundation for many other advanced trees.
- **Binary Search Tree (BST):** A binary tree with a crucial rule: the left child's value is **less than** the parent's, and the right child's value is **greater than** the parent's. This ordering makes searching incredibly fast.
- **Self-Balancing Trees (AVL & Red-Black):** These are smarter BSTs. They automatically adjust their structure after insertions or deletions to stay **balanced**, which guarantees fast $O(\log n)$ performance and prevents them from becoming inefficient like a linked list.
- **Heap:** A specialized binary tree used to quickly find the minimum or maximum element. It's like a tournament bracket, where the top-priority item (like the winner) is always at the root, making it perfect for implementing priority queues.
- **Trie (Prefix Tree):** A special tree used for storing and searching for strings. It's the magic behind an **autocomplete** feature, where each node represents a character in a word.



1.4. Common Real-World Applications of Trees

Trees aren't just a theoretical concept; they are the invisible engine behind many of the tools and technologies you use as a developer and consumer every day. Their ability to manage complexity makes them essential for high-performance applications.

- **How Git Manages Your Code 🌱:** Your entire Git commit history is a tree-like structure (a Directed Acyclic Graph). Each **commit** is a node that points to its parent commit(s). When you create a **branch**, you are essentially creating a new pointer to a node, allowing a new line of history to grow. A **merge commit** is simply a special node with two parents, unifying the two branches.
- **Game AI and Decision Making 🎮:** How does an enemy in a video game decide what to do? Often, it uses a **Decision Tree**. At each node, the AI asks a question ("Is the player's health low?"). The branches represent the answers ("Yes" or "No"), leading to an action ("Attack with a special move" or "Use a standard attack").
- **Compilers & How Your Code Actually Runs 💡:** When you write a line of code like `let x = 5 * 2;`, the compiler first turns it into an **Abstract Syntax Tree (AST)**. This tree represents the code's structure and order of operations, allowing the computer to understand and execute your logic correctly.
- **How Spotify Finds Any Song Instantly 🎵:** How can a streaming service search through millions of songs in milliseconds? Databases use specialized, wide and shallow trees called **B-Trees**. Their structure is optimized to minimize slow disk reads, allowing them to pinpoint the exact location of your data with very few lookups.



Chapter 2: The Binary Tree: A Fundamental Building Block

Now we move from general trees to the most common and important type in computer science: the **Binary Tree**. Its simplicity and structure make it the foundation for complex structures like search trees, heaps, and expression trees.

2.1. Anatomy and Properties of a Binary Tree

Imagine you're creating a simple "yes/no" troubleshooting guide. You start with a question, like "Does the device turn on?". If the answer is "No," you follow one path of instructions; if "Yes," you follow another. Each question leads to at most two new, more specific questions, until you find the solution.

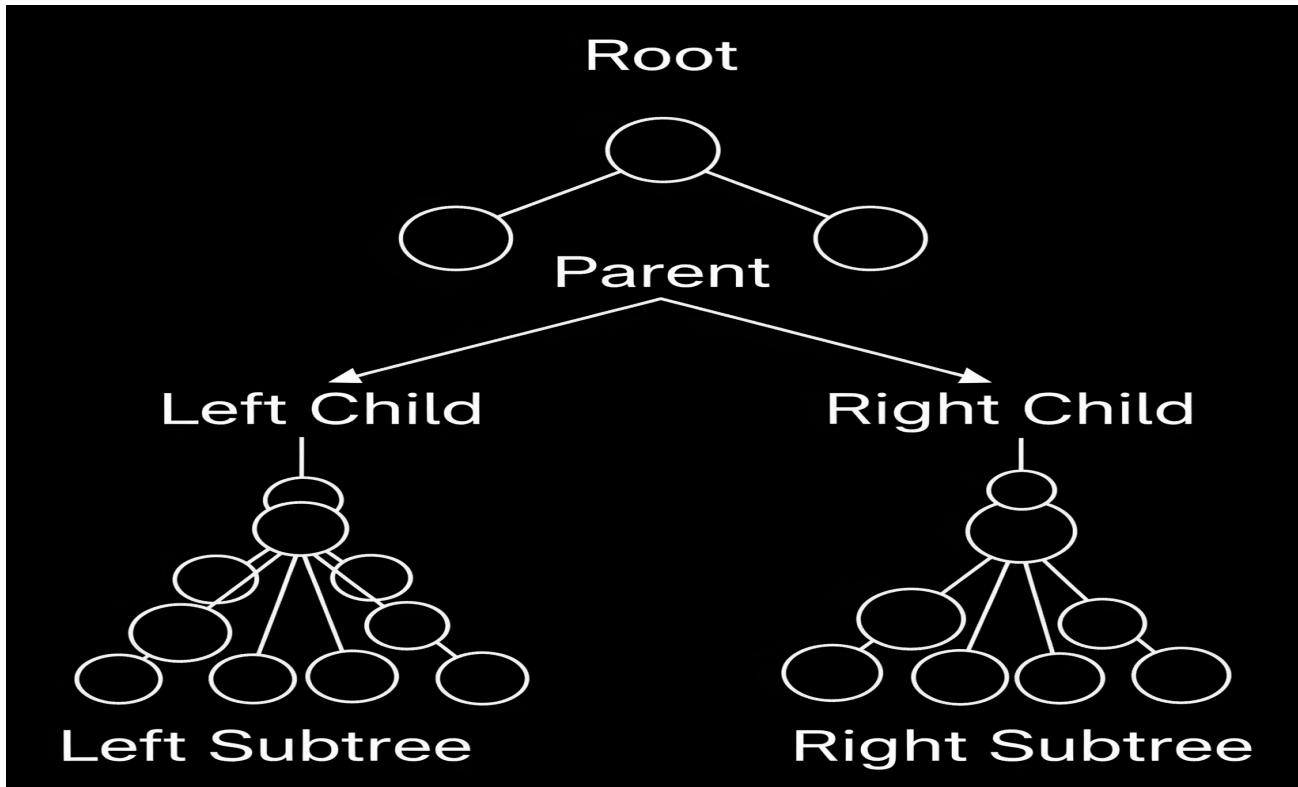
This is the exact idea behind a **Binary Tree**. It's a tree where every node can have at most **two children**. We don't just call them "children"; we give them specific names: the **left child** and the **right child**. This structure is perfect for representing any process that involves a series of binary (two-option) decisions.

The subtree rooted at the left child is called the **left subtree**, and the one at the right child is the **right subtree**.

Key Properties (Why this Structure is Efficient)

This simple "two-choice" structure leads to some powerful mathematical properties that we use to analyze efficiency.

- **Maximum Nodes at a Level:** The maximum number of nodes at any level ' l ' (where the root is at level 0) is 2^l . Each level can hold twice as many possibilities as the one above it.
- **Maximum Nodes in a Tree:** The maximum number of nodes in a binary tree of height ' h ' is $2^{h+1}-1$. This represents a tree that has explored every possible "yes/no" path down to a certain depth.
- **Minimum Height for N Nodes:** For a tree representing ' N ' final outcomes, the minimum possible height is $\lfloor \log_2 N \rfloor$. This compact structure is the key to the $O(\log n)$ performance that makes trees famous.



Code: The Binary Tree Node in C++

The C++ code for a binary tree node directly mirrors this "two-option" structure. Instead of a list of children, we have two specific pointers: `left` and `right`.

```

// The fundamental building block for a Binary Tree.
struct Node {
    int data;
    Node* left; // Pointer to the 'left' or 'no' path
    Node* right; // Pointer to the 'right' or 'yes' path

    // Constructor to initialize a node with a given value.
    // Children are initialized to nullptr, meaning no path exists yet.
    Node(int val) {
        data = val;
        left = nullptr;
        right = nullptr;
    }
};
```

2.2. Types of Binary Trees (Full, Complete, Perfect, Degenerate)

While all binary trees follow the "at most two children" rule, their *shape* can differ dramatically. The structure of a tree is vital because it directly impacts its performance. Certain shapes are highly efficient, while others can be as slow as a simple list.

Full Binary Tree

A binary tree is **full** if every node has either **0 or 2 children**. Think of it as a perfect tournament bracket where every match must have exactly two competitors; no node is left with a single "bye." This structure ensures that every branching point is a complete binary choice.

Complete Binary Tree

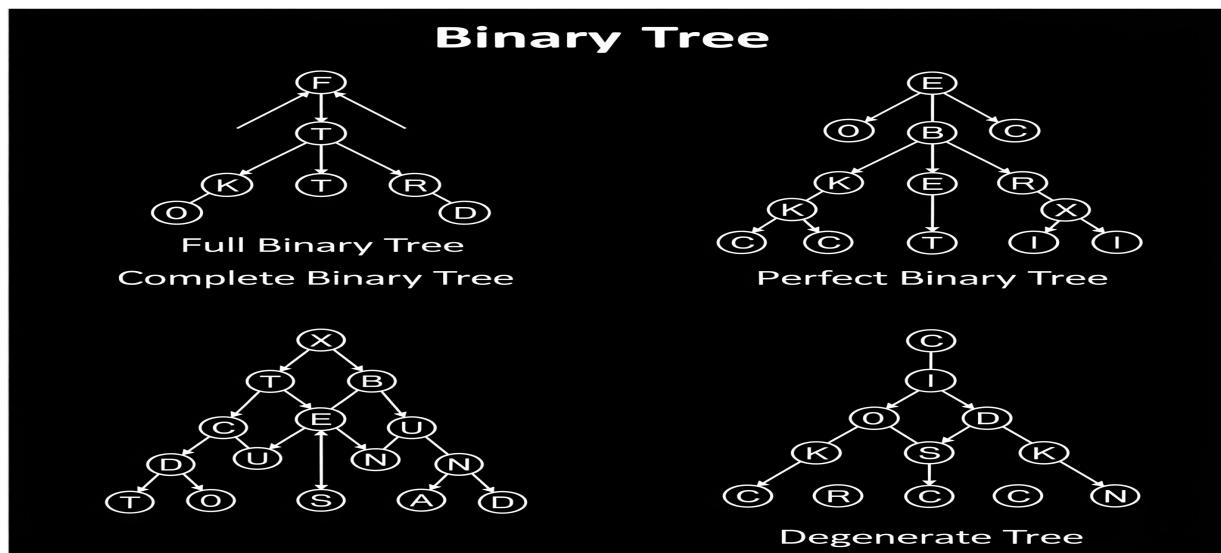
A binary tree is **complete** if all levels are filled, except possibly the last level, and the last level is filled from **left to right**. This is like filling seats in an auditorium row by row. You must fill every seat in a row completely before starting the next, and on the final row, you don't skip seats. This specific structure is crucial because it allows us to store the tree in an array very efficiently, which is the secret behind data structures like the **Binary Heap**.

Perfect Binary Tree

A binary tree is **perfect** if all its interior nodes have two children and all leaf nodes are at the **same level**. This is the most symmetrical and "ideal" tree shape. It's both full *and* complete. Imagine a perfectly balanced ancestral chart where every single ancestor for a set number of generations had exactly two children. This creates a perfectly filled-out pyramid, representing the maximum number of nodes for a given height and thus the most compact structure possible.

Degenerate (or Pathological) Tree

This is the worst-case scenario for a binary tree. A **degenerate** tree is one where every parent node has only **one child**. This happens if you insert already sorted data into a basic Binary Search Tree. It's no longer really a tree; it's just a **linked list** masquerading as one. All the benefits of quick searching are lost because you have to traverse it node by node, making operations slow ($O(n)$).



2.3. Core Traversal Algorithms (Theory & Code)

"Traversing" a tree means visiting every single node exactly once in a systematic way. It's like giving a tour of a building; you need a planned route to ensure you don't miss any rooms. The two main strategies are **Depth-First Search (DFS)**, where you go as deep as possible down one path before backing up, and **Breadth-First Search (BFS)**, where you explore level by level.

2.3.1. Depth-First: In-order, Pre-order, Post-order

DFS is like navigating a maze by always taking the leftmost path. When you hit a dead end, you backtrack and try the next available path. There are three common ways to order your visit at each node during this process.

In-order Traversal (Left, Root, Right)

This traversal follows the rule: visit the **left** subtree, then the **root** node, then the **right** subtree.

Its most famous application is on a **Binary Search Tree (BST)**. An in-order traversal of a BST will visit the nodes in **ascending sorted order**. It's like reading a perfectly organized dictionary from A to Z, giving you a naturally sorted list of its contents.

```
void inorderTraversal(Node* root) {
    if (root == nullptr) {
        return;
    }
    inorderTraversal(root->left);    // 1. Go Left
    std::cout << root->data << " "; // 2. Visit Root
    inorderTraversal(root->right);   // 3. Go Right
}
```

Pre-order Traversal (Root, Left, Right)

Here, you visit the **root** node first, then the **left** subtree, then the **right** subtree.

This method is used to **create a copy of a tree** because you can create the parent node before you recursively create its children. It's also like creating an outline for a document; you write the main chapter title (root) first, then you detail all of its sections and sub-sections (the left subtree), before moving on to the next main chapter (the right subtree).

```
void preorderTraversal(Node* root) {
    if (root == nullptr) {
        return;
    }
    std::cout << root->data << " "; // 1. Visit Root
    preorderTraversal(root->left);   // 2. Go Left
    preorderTraversal(root->right);  // 3. Go Right
}
```

Post-order Traversal (Left, Right, Root)

This traversal visits the **left** subtree, then the **right** subtree, and finally, the **root** node.

The primary use case for this is **deleting a tree from memory**. You must delete a node's children before you can delete the node itself, otherwise you'd lose the pointers to them and create a memory leak. It's like dismantling a piece of furniture; you have to take off the legs (children) before you can remove the tabletop (parent) they were attached to.

```
void postorderTraversal(Node* root) {
    if (root == nullptr) {
        return;
    }
    postorderTraversal(root->left);    // 1. Go Left
    postorderTraversal(root->right);   // 2. Go Right
    std::cout << root->data << " "; // 3. Visit Root
}
```

2.3.2. Breadth-First: Level-order

This is the only common Breadth-First Search (BFS) traversal. Instead of going deep, it explores the tree layer by layer.

Level-order Traversal

This traversal visits nodes level by level from top to bottom, and within each level, from left to right. To keep track of which nodes to visit next, it uses a **queue** data structure.

This is like reading a book page by page, line by line. You read the first line (level 0), then the second (level 1), and so on. This traversal is used to find the shortest path between two nodes in terms of the number of edges.

```
#include <queue> // Required for level-order

void levelorderTraversal(Node* root) {
    if (root == nullptr) {
        return;
    }

    std::queue<Node*> q;
    q.push(root);

    while (!q.empty()) {
        Node* current = q.front();
```

```

        q.pop();
        std::cout << current->data << " ";
    }

    if (current->left != nullptr) {
        q.push(current->left);
    }
    if (current->right != nullptr) {
        q.push(current->right);
    }
}
}

```

2.4. Essential Operations (Theory & C++ Implementation)

Now that we know how to build and traverse a binary tree, let's look at the core operations for modifying it: finding, adding, and removing nodes.

2.4.1. Search, Insertion, and Deletion

Searching for a Node

Since there's no order, finding a node with a specific value requires us to potentially check every single node. It's like searching for a specific file on your computer without using the search bar—you have to open folders one by one until you find it. We can use any traversal method (like Pre-order) to visit each node and check its value. This is a brute-force approach with a time complexity of $O(n)$.

```

/**
 * Searches for a key in a generic Binary Tree.
 * Returns true if the key is found, otherwise false.
 */
bool search(Node* root, int key) {
    if (root == nullptr) {
        return false;
    }
    // Check the current node
    if (root->data == key) {
        return true;
    }
    // Recursively search in the left and right subtrees
    return search(root->left, key) || search(root->right, key);
}

```

Inserting a Node

In an unordered binary tree, there's no "correct" place to insert a new node. A common strategy is to add it to the **first available spot**, similar to how you'd fill an empty seat in an auditorium starting from the first row, left to right. We use a level-order traversal (with a queue) to find the first parent node that has either a missing left or right child.

```
#include <queue>

/**
 * Inserts a new node at the first available position in the tree.
 */
void insert(Node*& root, int data) {
    Node* newNode = new Node(data);
    if (root == nullptr) {
        root = newNode;
        return;
    }

    std::queue<Node*> q;
    q.push(root);

    while (!q.empty()) {
        Node* current = q.front();
        q.pop();

        if (current->left == nullptr) {
            current->left = newNode;
            return;
        } else {
            q.push(current->left);
        }

        if (current->right == nullptr) {
            current->right = newNode;
            return;
        } else {
            q.push(current->right);
        }
    }
}
```

Deleting a Node

Deletion is the most complex operation. Simply removing a node from the middle of the tree could disconnect an entire subtree, breaking the structure. The standard approach is more like a corporate reshuffle than a simple removal.

The Strategy: To delete a node with value `key`:

1. Find the node you want to delete (`nodeToDelete`).
2. Find the **deepest, rightmost node** in the tree. This node is structurally easy to remove (it's a leaf or has only one child). Think of this as the employee with the least responsibility.
3. **Swap the data** from the deepest node into `nodeToDelete`. In our analogy, the junior employee takes over the senior's role.
4. Now, **delete the deepest node** from its original position. Since its responsibilities have been transferred, removing its now-empty position is simple and doesn't disrupt the company hierarchy.

This process preserves the tree's overall structure while removing the desired value.

```
// Helper function to delete the deepest, rightmost node
void deleteDeepest(Node* root, Node* nodeToDelete) {
    std::queue<Node*> q;
    q.push(root);
    Node* temp;
    while (!q.empty()) {
        temp = q.front();
        q.pop();
        if (temp == nodeToDelete) {
            temp = nullptr;
            delete nodeToDelete;
            return;
        }
        if (temp->right) {
            if (temp->right == nodeToDelete) {
                temp->right = nullptr;
                delete nodeToDelete;
                return;
            } else {
                q.push(temp->right);
            }
        }
        if (temp->left) {
            if (temp->left == nodeToDelete) {
                temp->left = nullptr;
                delete nodeToDelete;
                return;
            } else {
                q.push(temp->left);
            }
        }
    }
}
```

```

        }
    }

/**
 * Deletes a node with the given key.
 */
void deleteNode(Node*& root, int key) {
    if (root == nullptr) return;
    if (root->left == nullptr && root->right == nullptr) {
        if (root->data == key) {
            delete root;
            root = nullptr;
        }
        return;
    }

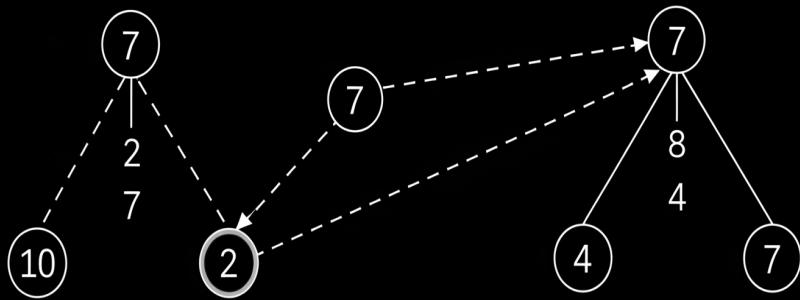
    std::queue<Node*> q;
    q.push(root);
    Node* nodeToDelete = nullptr;
    Node* temp = nullptr;

    // Find the node to delete and the deepest node
    while (!q.empty()) {
        temp = q.front();
        q.pop();
        if (temp->data == key) {
            nodeToDelete = temp;
        }
        if (temp->left) q.push(temp->left);
        if (temp->right) q.push(temp->right);
    }

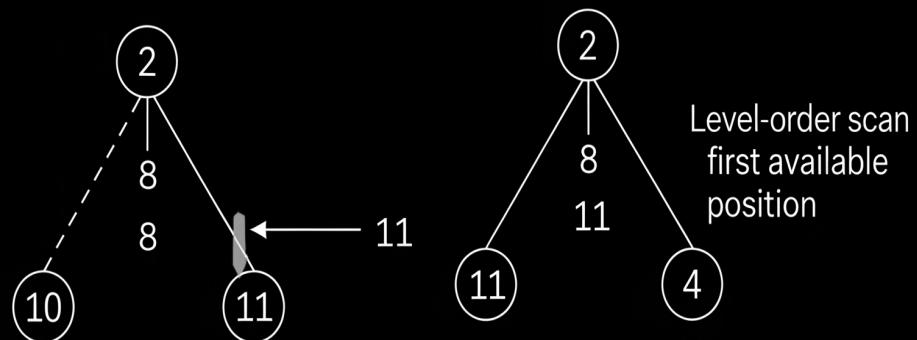
    if (nodeToDelete != nullptr) {
        int x = temp->data; // Data of the deepest node
        deleteDeepest(root, temp); // Delete the original deepest node
        nodeToDelete->data = x; // Swap data
    }
}

```

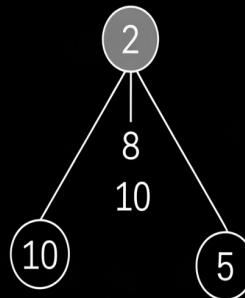
7a) Search



(2) Insertion



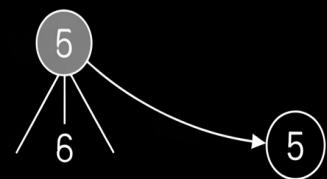
3a) Delete to node 2



Step 3b

Data transfer to node 5
location nchere node
node node 5

Step 3b node to delete



Data to trastion the
deepest position of 5
deciation 4

Chapter 3: The Binary Search Tree (BST): Bringing Order to Data

In the last chapter, we saw that a generic binary tree is like a randomly organized library; finding anything requires a full search. Now, we'll introduce a simple rule that transforms this chaotic structure into an efficient, organized system. This new, "smarter" tree is the **Binary Search Tree (BST)**.

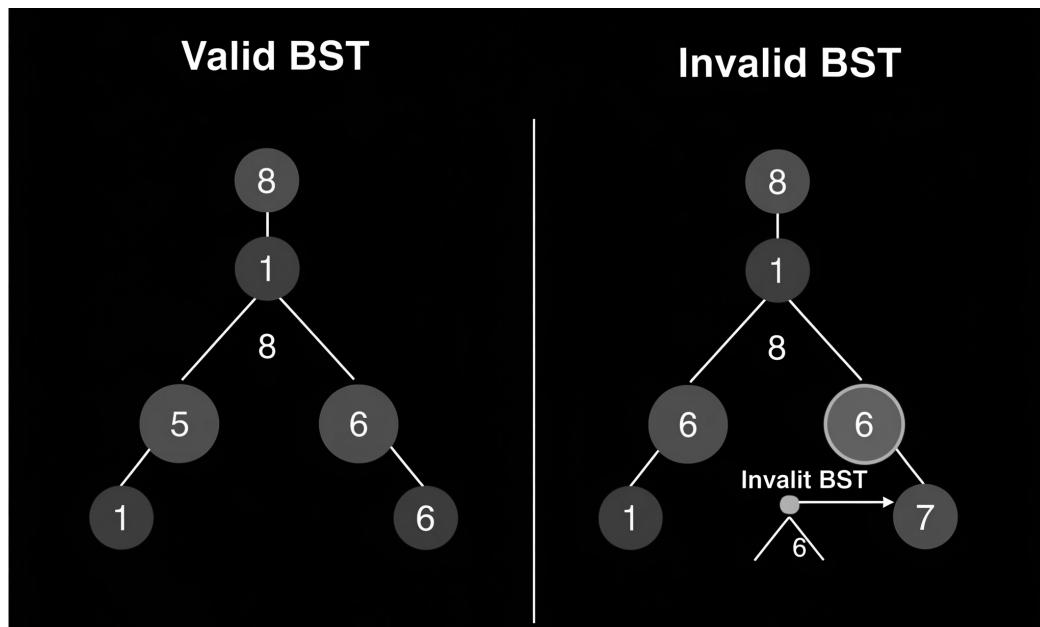
3.1. The BST Invariant: The Golden Rule for Ordering

The power of a BST comes from a single, elegant rule known as the **BST Invariant**. It's like playing a "higher or lower" number guessing game. With every guess, you're told whether to go higher or lower, allowing you to eliminate half the possibilities. A BST embeds this logic into its very structure.

For any given node in the tree, the BST Invariant states:

1. All values in the node's **left subtree** must be **less than** the node's own value.
2. All values in the node's **right subtree** must be **greater than** the node's own value.
3. Both the left and right subtrees must also be valid binary search trees.

This rule applies to **every single node**, not just the root. This "higher or lower" decision at each step is what allows us to discard huge chunks of the tree during a search, leading to its famous $O(\log n)$ efficiency.



Code: Validating the BST Invariant

The best way to understand the rule in code is to write a function that checks if a tree follows it. This function recursively checks each node, ensuring it stays within a valid range determined by its ancestors.

```
#include <climits> // For INT_MIN and INT_MAX

// Helper function for isBST
bool isValid(Node* node, long minValue, long maxValue) {
    if (node == nullptr) {
        return true; // An empty tree is a valid BST
    }
    // The current node's value must be within the valid range
    if (node->data <= minValue || node->data >= maxValue) {
        return false;
    }
    // Recursively check the left and right subtrees, updating the range
    // For the left child, the new max is the current node's value.
    // For the right child, the new min is the current node's value.
    return isValid(node->left, minValue, node->data) &&
           isValid(node->right, node->data, maxValue);
}

/**
 * Checks if a given binary tree is a valid Binary Search Tree.
 */
bool isBST(Node* root) {
    // Start with the widest possible range
    return isValid(root, LONG_MIN, LONG_MAX);
}
```

3.2. Operations in a BST (Search, Insert, Delete with C++ examples)

The BST Invariant is the engine that powers its efficiency. Because we always know that `left < root < right`, we never have to search the entire tree. At every node, we can make an intelligent choice, discarding about half of the remaining nodes in the process.

Searching in a BST

Searching a BST is like the "higher or lower" game in action. You start at the root and compare its value to your target key.

- If it's a match, you've found it.

- If your key is **less** than the node's value, you know it can *only* be in the **left** subtree.
- If your key is **greater**, you know it can *only* be in the **right** subtree.

You repeat this process, narrowing down the search area at each step until you find the value or hit an empty spot (`nullptr`), which means the value isn't in the tree.

```
/**
 * Searches for a key in a BST.
 * Returns the node if found, otherwise nullptr.
 */
Node* search(Node* root, int key) {
    // Base Cases: root is null or key is present at root
    if (root == nullptr || root->data == key) {
        return root;
    }

    // Key is greater than root's data, so search in the right subtree
    if (root->data < key) {
        return search(root->right, key);
    }

    // Key is smaller than root's data, so search in the left subtree
    return search(root->left, key);
}
```

Inserting into a BST

Insertion uses the exact same "higher or lower" logic. You trace a path down the tree as if you were searching for the value you want to add. When you find the empty (`nullptr`) spot where the value *should* be, you insert the new node there. It's like finding the precise, empty slot in an alphabetized filing cabinet to slide in a new file—no need to rearrange anything else.

```
/**
 * Inserts a new node with a given key into the BST.
 */
Node* insert(Node* root, int key) {
    // If the tree is empty, return a new node
    if (root == nullptr) {
        return new Node(key);
    }

    // Otherwise, recur down the tree
    if (key < root->data) {
        root->left = insert(root->left, key);
    }
}
```

```

    } else if (key > root->data) {
        root->right = insert(root->right, key);
    }

    // Return the (unchanged) node pointer
    return root;
}

```

Deleting from a BST 🗑️

Deletion is the most intricate operation because we must preserve the BST Invariant. The process depends on how many children the node to be deleted has.

- Case 1: Deleting a Leaf (0 Children)
This is simple. You just disconnect the node from its parent. It's like removing the last card from a stack—nothing else is affected.
- Case 2: Deleting a Node with One Child
This is also straightforward. You "bypass" the node by linking its parent directly to its child. It's like an employee with one direct report leaving; their report now works directly for their manager, preserving the chain of command.
- Case 3: Deleting a Node with Two Children
This is the complex case. We can't just remove the node. The solution is a clever replacement. We find the node's in-order successor (which is the smallest value in its right subtree). We copy the successor's value into the node we want to delete, and then we recursively delete the successor from its original spot. This is like a manager leaving; you promote their most qualified subordinate (the successor) into their role and then fill that subordinate's now-empty, simpler position.

```

// Helper function to find the in-order successor
Node* findMin(Node* node) {
    while (node && node->left != nullptr) {
        node = node->left;
    }
    return node;
}

/**
 * Deletes a node with the given key from the BST.
 */
Node* deleteNode(Node* root, int key) {
    if (root == nullptr) return root;

    // Find the node to be deleted
    if (key < root->data) {

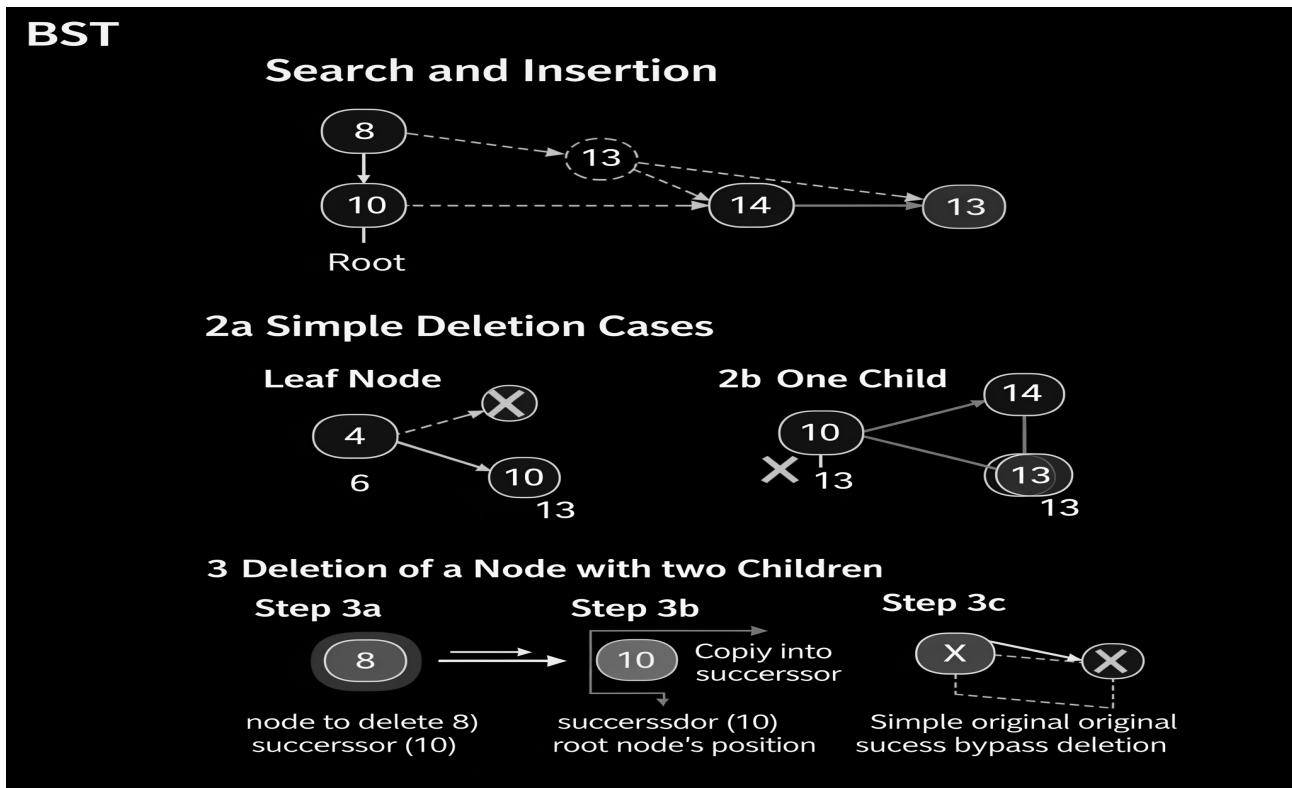
```

```

        root->left = deleteNode(root->left, key);
    } else if (key > root->data) {
        root->right = deleteNode(root->right, key);
    } else { // This is the node to be deleted
        // Case 1 & 2: Node with only one child or no child
        if (root->left == nullptr) {
            Node* temp = root->right;
            delete root;
            return temp;
        } else if (root->right == nullptr) {
            Node* temp = root->left;
            delete root;
            return temp;
        }
    }

    // Case 3: Node with two children
    Node* temp = findMin(root->right); // Find in-order successor
    root->data = temp->data; // Copy the successor's content to this node
    root->right = deleteNode(root->right, temp->data); // Delete the
successor
}
return root;
}

```



3.3. The Problem of Imbalance: When BSTs Fail

So far, the BST seems like the perfect data structure, offering fast $O(\log n)$ operations. But this incredible performance has a major catch: it depends entirely on the **shape** of the tree. If the tree is short and bushy, it's fast. But if it becomes tall and spindly, it fails dramatically.

This happens when you insert data that's already sorted or nearly sorted. Imagine inserting the numbers 10, 20, 30, 40, 50 into a BST in that order.

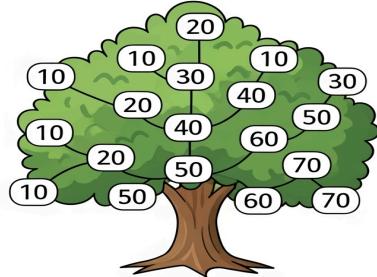
- You insert 10.
- You insert 20. Since $20 > 10$, it goes to the right.
- You insert 30. Since $30 > 10$ and $30 > 20$, it goes to the right of 20.
- ...and so on.

You don't get a branching, bushy tree. Instead, you get a long, single chain of right children. This is called a **degenerate tree**.

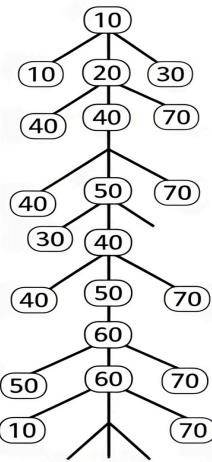
When this happens, the tree's height becomes ' n ' instead of the ideal ' $\log n$ '. All the advantages of the BST are lost. A search operation is no longer a "higher or lower" game that eliminates half the options; it's just like walking down a linked list, one node at a time. The performance for search, insert, and delete operations all degrades to a slow $O(n)$.

Diagram: Balanced vs. Unbalanced BST

This "problem of imbalance" is a deal-breaker for many real-world applications where you can't guarantee random data insertion. This is precisely why we need our next topic: **self-balancing trees**, which are designed to automatically prevent this worst-case scenario from ever happening.



Balanced Tree



Unbalanced
maximum traversal length
a search for linked list

Unbalanced
(Degenerate) Tree

Chapter 4: Self-Balancing Trees: Guaranteed Performance



We've seen that a standard BST can fail by becoming unbalanced. Self-balancing trees are the solution. They are BSTs that automatically perform small, localized adjustments after every insertion or deletion to ensure the tree remains short and bushy, guaranteeing $O(\log n)$ performance for all operations.

4.1. The AVL Tree

The **AVL Tree**, named after its inventors Adelson-Velsky and Landis, was the first self-balancing binary search tree. It maintains its balance by ensuring that the height difference between the subtrees of any node is at most one.

4.1.1. Balancing Factors and Rotations (Theory & Code)

The AVL tree's magic lies in two core concepts: the **balance factor** to detect an imbalance, and **rotations** to fix it.

The Balance Factor

The **Balance Factor** of a node is calculated as: `height(left subtree) - height(right subtree)`.

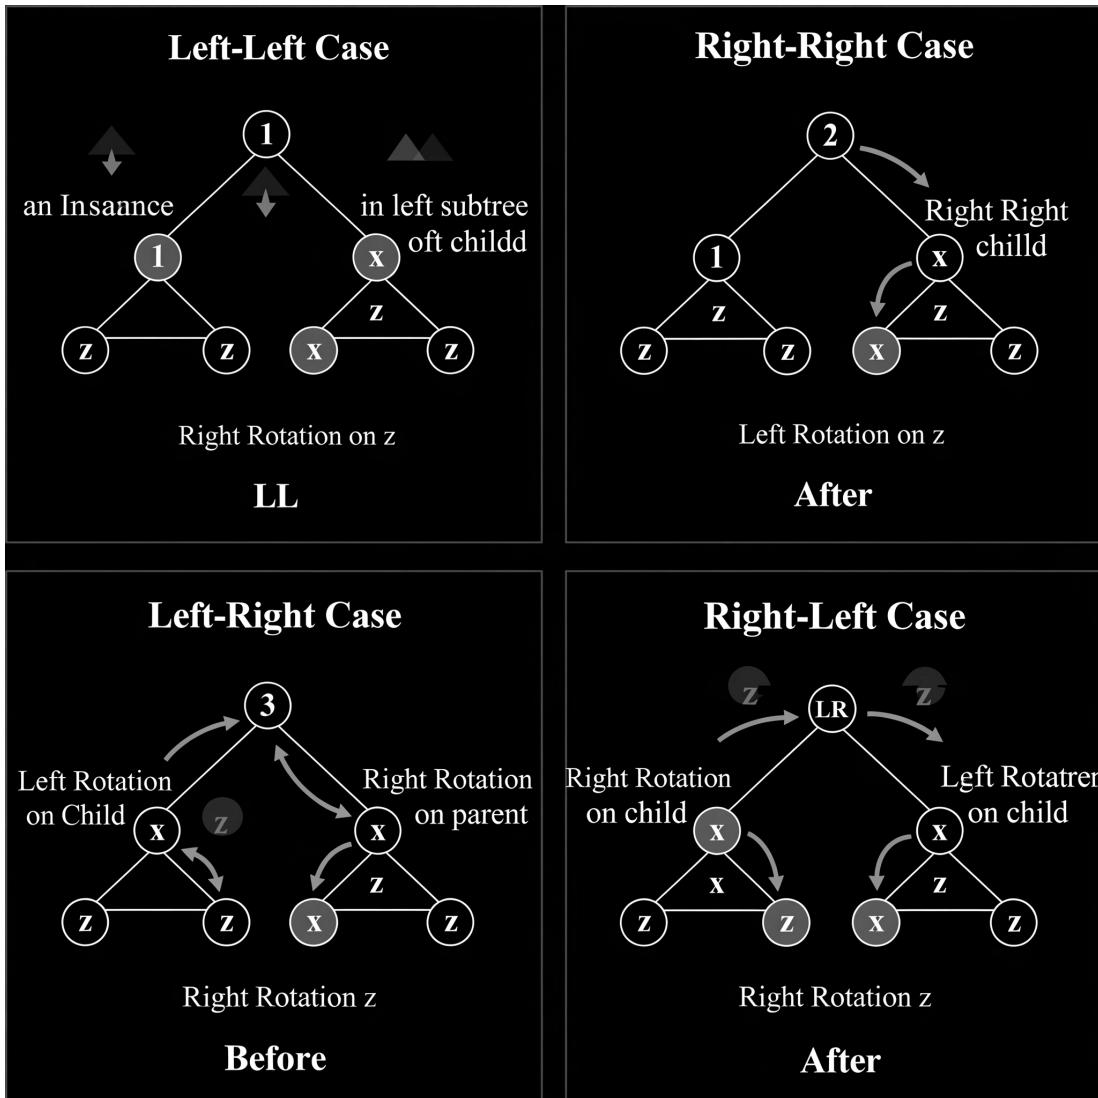
The **AVL Invariant** is that for every node in the tree, this balance factor must be **-1, 0, or 1**. If an insertion or deletion causes any node's balance factor to become -2 or 2, the tree is considered "unbalanced," and an immediate fix is required. It's like a hanging mobile; a slight difference in weight on either side is fine, but if one side becomes too heavy, the entire structure tilts and needs to be adjusted.

Rotations: The Fix

When the tree becomes unbalanced, it performs simple, constant-time ($O(1)$) operations called **rotations**. These are like precise adjustments to the mobile's strings that restore its equilibrium. Rotations rearrange a small number of nodes to fix the height imbalance while preserving the BST property.

There are four imbalance scenarios, which are fixed by two types of simple rotations (Left and Right) and two combination rotations:

1. **Left-Left Case:** Fixed by a **Right Rotation**.
2. **Right-Right Case:** Fixed by a **Left Rotation**.
3. **Left-Right Case:** Fixed by a **Left Rotation** followed by a **Right Rotation**.
4. **Right-Left Case:** Fixed by a **Right Rotation** followed by a **Left Rotation**.



Code: AVL Node and Rotation Functions

The AVL node needs to track its height. The core logic resides in the rotation functions.

```
// AVL tree node with a height parameter
struct Node {
    int data;
    Node* left;
    Node* right;
    int height;
};

// Get height of a node (handles nullptr)
```

```

int getHeight(Node* N) {
    if (N == nullptr) return 0;
    return N->height;
}

// Get Balance factor of a node
int getBalance(Node* N) {
    if (N == nullptr) return 0;
    return getHeight(N->left) - getHeight(N->right);
}

// Right rotate subtree rooted with y
Node* rightRotate(Node* y) {
    Node* x = y->left;
    Node* T2 = x->right;

    // Perform rotation
    x->right = y;
    y->left = T2;

    // Update heights
    y->height = std::max(getHeight(y->left), getHeight(y->right)) + 1;
    x->height = std::max(getHeight(x->left), getHeight(x->right)) + 1;

    return x; // New root
}

// Left rotate subtree rooted with x
Node* leftRotate(Node* x) {
    Node* y = x->right;
    Node* T2 = y->left;

    // Perform rotation
    y->left = x;
    x->right = T2;

    // Update heights
    x->height = std::max(getHeight(x->left), getHeight(x->right)) + 1;
    y->height = std::max(getHeight(y->left), getHeight(y->right)) + 1;

    return y; // New root
}

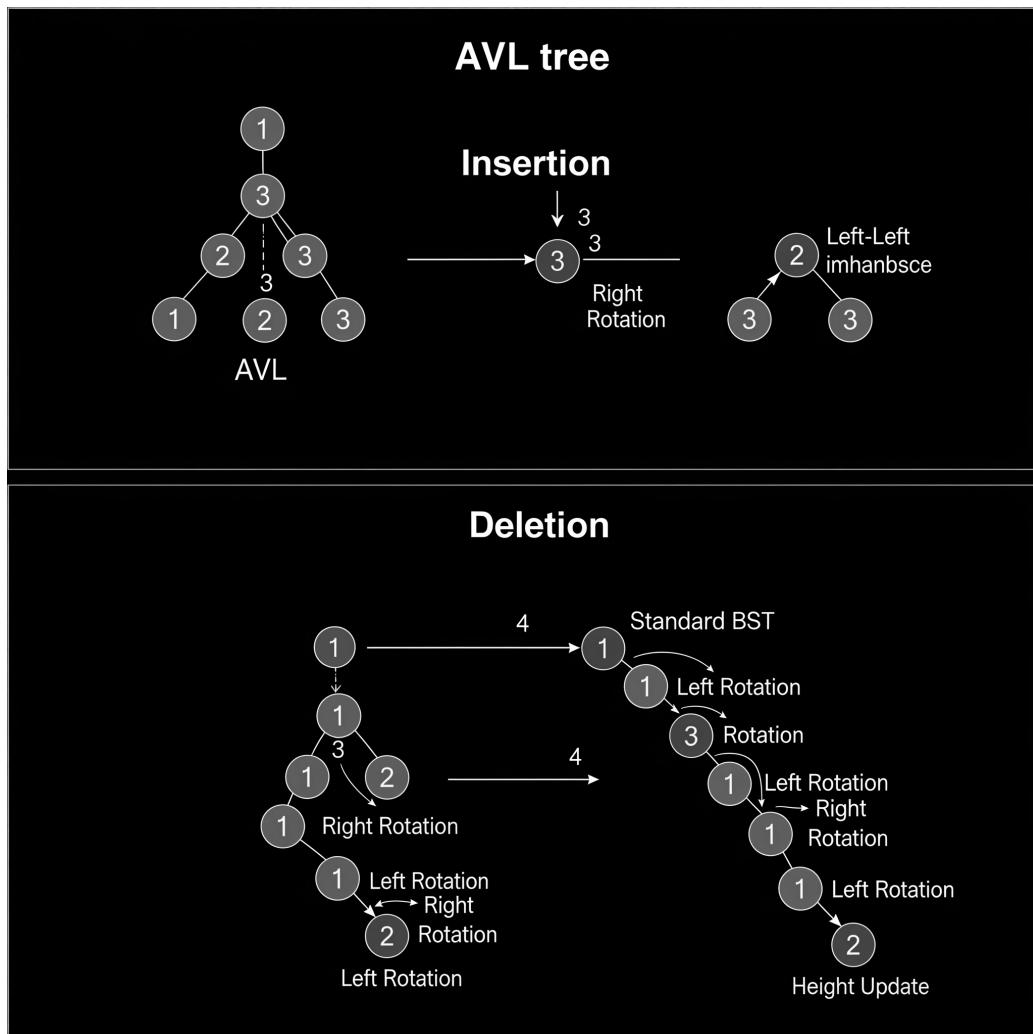
```

4.1.2. Insertion and Deletion Walkthrough

Insertion

The process combines a standard BST insertion with a "walk back up" phase to check for and fix imbalances.³

1. **Perform Standard BST Insert:** Recursively find the correct spot and insert the new node.
2. **Update Heights:** On the way back up the recursion, update the height of each ancestor node.
3. **Check Balance and Rotate:** At each ancestor, calculate its balance factor. If an imbalance is found, perform the appropriate rotation (LL, RR, LR, or RL) based on the insertion path. Once one fix is made, the tree is balanced, and no more rotations are needed for that insertion.



Code: AVL Insertion and Deletion

```
Node* insert(Node* node, int data) {
    // 1. Perform normal BST insertion
    if (node == nullptr) {
        Node* temp = new Node();
        temp->data = data;
        temp->left = nullptr;
        temp->right = nullptr;
        temp->height = 1; // New node is initially added at Leaf
        return temp;
    }
    if (data < node->data)
        node->left = insert(node->left, data);
    else if (data > node->data)
        node->right = insert(node->right, data);
    else // Equal keys not allowed
        return node;

    // 2. Update height of this ancestor node
    node->height = 1 + std::max(getHeight(node->left), getHeight(node->right));

    // 3. Get the balance factor to check for imbalance
    int balance = getBalance(node);

    // 4. If unbalanced, there are 4 cases

    // Left Left Case
    if (balance > 1 && data < node->left->data)
        return rightRotate(node);

    // Right Right Case
    if (balance < -1 && data > node->right->data)
        return leftRotate(node);

    // Left Right Case
    if (balance > 1 && data > node->left->data) {
        node->left = leftRotate(node->left);
        return rightRotate(node);
    }

    // Right Left Case
    if (balance < -1 && data < node->right->data) {
        node->right = rightRotate(node->right);
        return leftRotate(node);
    }
}
```

```

    }

    // Return the (unchanged) node pointer
    return node;
}

// (Note: The full deletion code is similarly structured but more verbose,
// handling all cases and then applying the same balancing logic as insertion
// on the way back up the recursive stack for each ancestor.)

```

4.2. The Red-Black Tree

The **Red-Black Tree (RBT)** is another self-balancing binary search tree. While the AVL tree uses a strict height rule, the Red-Black Tree uses a cleverer, more "relaxed" approach with node coloring to ensure balance.

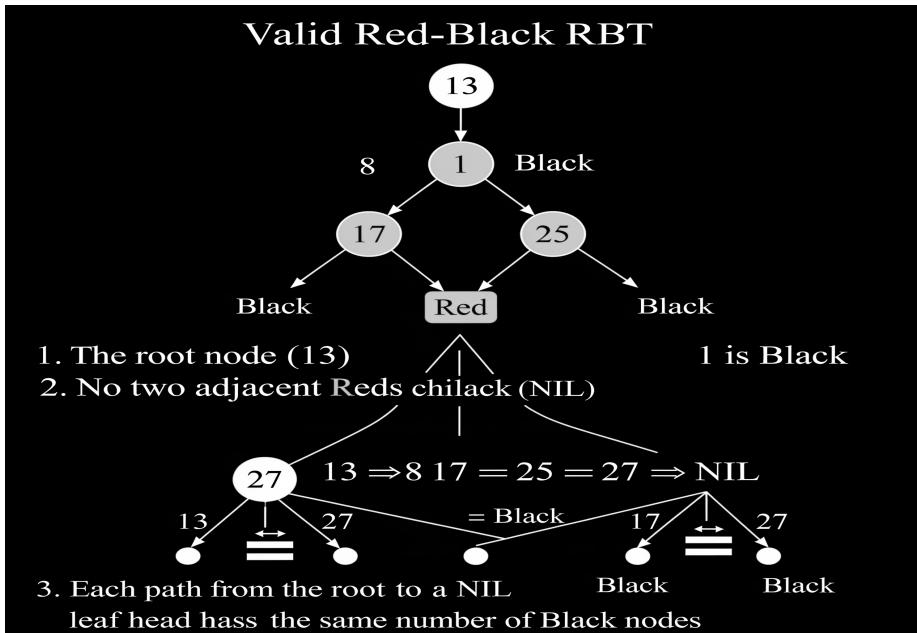
This relaxed approach is like comparing two project management styles. An AVL tree is a rigid plan where a single missed deadline forces immediate reshuffling (rotations) to keep everything perfectly on schedule. A Red-Black Tree is a more flexible plan that allows for some slack; it only intervenes with adjustments (recoloring or rotations) when the overall project timeline is seriously at risk. This flexibility often means fewer rotations, making insertions and deletions faster on average than in AVL trees. This is why RBTs are used in many standard library implementations, like C++ `std::map` and `std::set`.

4.2.1. Core Principles and Coloring Rules

The "balance" of an RBT is maintained not by tracking height, but by enforcing a set of five simple rules regarding a `color` attribute (either **Red** or **Black**) stored in each node.

1. Every node is either **Red** or **Black**.
2. The **root** node is always **Black**.
3. All leaves (NIL or `nullptr` nodes) are considered **Black**.
4. A **Red** node cannot have a **Red** child. (This is the "Red-Red" rule.)
5. Every path from a given node to any of its descendant NIL leaves contains the **same number of Black nodes**. (This is the "Black-Height" property.)

The magic is in rules 4 and 5. Together, they mathematically guarantee that the longest possible path in the tree is no more than twice as long as the shortest possible path. This keeps the tree sufficiently balanced to ensure $O(\log n)$ performance.



Code: The Red-Black Tree Node

The node structure is similar to a BST node but includes a `color` and a crucial pointer to its `parent`, which is needed for the fix-up operations.

```
enum Color { RED, BLACK };

struct Node {
    int data;
    Color color;
    Node *left, *right, *parent;

    Node(int data) : data(data) {
        color = RED; // New nodes are always inserted as RED
        left = right = parent = nullptr;
    }
};
```

4.2.2. Conceptual Insertion and Deletion

The full implementation of RBTs is notoriously complex. Here, we'll focus on the *conceptual* steps to understand the logic, which relies on "fixing" any rule violations after a standard BST operation.

Conceptual Insertion

When we add a new node, we want to disrupt the tree as little as possible. The easiest rule to temporarily break and then fix is the "Red-Red" rule.

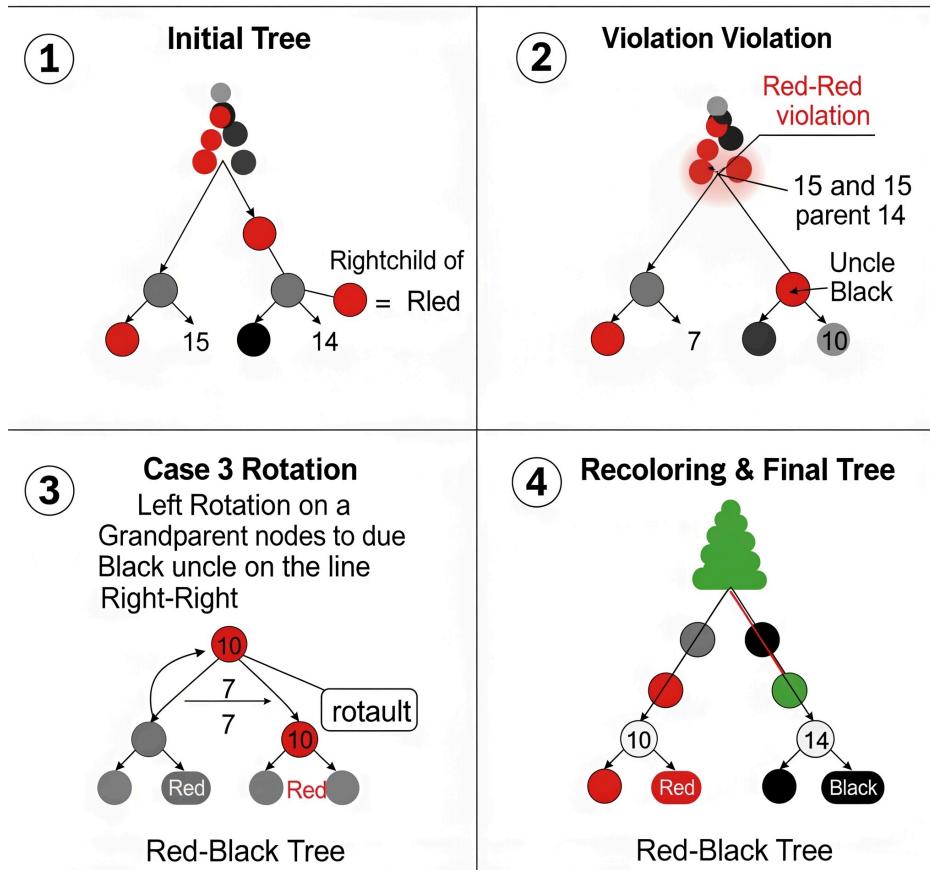
1. **Insert and Color Red:** Perform a standard BST insertion and color the new node **Red**. This ensures the black-height property (Rule 5) is not violated.
2. **Fix Violations:** The new red node might now have a red parent, violating Rule 4. A "fix-up" procedure is then called. It works its way up the tree, using **recoloring** (a fast fix) and **rotations** (a more structural fix) to resolve the "Red-Red" conflict until all rules are satisfied.

Conceptual Deletion

Deletion is more complex because removing a **Black** node can disrupt the black-height property (Rule 5), which is harder to fix.

1. **Perform BST Deletion:** Perform a standard BST deletion.
2. **Fix Violations:** If a **Black** node was removed, the tree is now unbalanced. This creates a temporary "double-black" node in its place. A "fix-up" procedure then travels up the tree, "pushing" the extra blackness upwards through recoloring and rotations until it can be resolved (e.g., by reaching the root or a red node).

Red-Black Tree Insertion rotation and Recoloring



Code: Conceptual Functions

This code is complete and functional, but it is also complex. The most important parts to study are the `insertFixup` and `deleteFixup` methods, which contain the core logic for maintaining the Red-Black Tree properties.

```
#include <iostream>

enum Color { RED, BLACK };

struct Node {
    int data;
    Color color;
    Node *left, *right, *parent;

    Node(int data) : data(data), color(RED), left(nullptr), right(nullptr),
parent(nullptr) {}
};

class RedBlackTree {
private:
    Node* root;

protected:
    void LeftRotate(Node* x) {
        Node* y = x->right;
        x->right = y->left;
        if (y->left != nullptr) {
            y->left->parent = x;
        }
        y->parent = x->parent;
        if (x->parent == nullptr) {
            root = y;
        } else if (x == x->parent->left) {
            x->parent->left = y;
        } else {
            x->parent->right = y;
        }
        y->left = x;
        x->parent = y;
    }

    void rightRotate(Node* y) {
```

```

Node* x = y->left;
y->left = x->right;
if (x->right != nullptr) {
    x->right->parent = y;
}
x->parent = y->parent;
if (y->parent == nullptr) {
    root = x;
} else if (y == y->parent->left) {
    y->parent->left = x;
} else {
    y->parent->right = x;
}
x->right = y;
y->parent = x;
}

void insertFixup(Node* z) {
    while (z->parent != nullptr && z->parent->color == RED) {
        if (z->parent == z->parent->parent->left) { // Parent is a left
child
            Node* y = z->parent->parent->right; // Uncle
            if (y != nullptr && y->color == RED) {
                // Case 1: Uncle is RED
                z->parent->color = BLACK;
                y->color = BLACK;
                z->parent->parent->color = RED;
                z = z->parent->parent;
            } else {
                if (z == z->parent->right) {
                    // Case 2: Triangle (z is right child)
                    z = z->parent;
                    LeftRotate(z);
                }
                // Case 3: Line (z is left child)
                z->parent->color = BLACK;
                z->parent->parent->color = RED;
                RightRotate(z->parent->parent);
            }
        } else { // Parent is a right child
            Node* y = z->parent->parent->left; // Uncle
            if (y != nullptr && y->color == RED) {
                // Case 1: Uncle is RED
                z->parent->color = BLACK;
                y->color = BLACK;
            }
        }
    }
}

```

```

        z->parent->parent->color = RED;
        z = z->parent->parent;
    } else {
        if (z == z->parent->left) {
            // Case 2: Triangle (z is left child)
            z = z->parent;
            rightRotate(z);
        }
        // Case 3: Line (z is right child)
        z->parent->color = BLACK;
        z->parent->parent->color = RED;
        LeftRotate(z->parent->parent);
    }
}
root->color = BLACK;
}

void inorderHelper(Node* node) {
    if (node != nullptr) {
        inorderHelper(node->left);
        std::cout << node->data << "(" << (node->color == RED ? "R" : "B")
<< ")";
        inorderHelper(node->right);
    }
}

public:
RedBlackTree() : root(nullptr) {}

void insert(int data) {
    Node* z = new Node(data);
    Node* y = nullptr;
    Node* x = root;

    // Standard BST insert
    while (x != nullptr) {
        y = x;
        if (z->data < x->data) {
            x = x->left;
        } else {
            x = x->right;
        }
    }
}

```

```

z->parent = y;
if (y == nullptr) {
    root = z;
} else if (z->data < y->data) {
    y->left = z;
} else {
    y->right = z;
}

// Fix any violations
insertFixup(z);
}

void inorder() {
    inorderHelper(root);
    std::cout << std::endl;
}
};

// Example Usage
// int main() {
//     RedBlackTree rbt;
//     rbt.insert(10);
//     rbt.insert(20);
//     rbt.insert(30);
//     rbt.insert(15);
//     rbt.insert(17);
//     rbt.insert(40);
//     rbt.insert(50);
//
//     // In-order traversal will show nodes and their colors
//     rbt.inorder();
//     return 0;
// }

```

Note on Deletion: The code for deletion in a Red-Black Tree is significantly more complex than insertion, involving numerous cases in its `deleteFixup` procedure. For a tutorial focused on core understanding, the detailed insertion logic above provides a solid foundation for how RBTs work, while deletion is often left as an advanced exercise.

Chapter 5: Specialized Trees for Specific Tasks



We've focused on binary trees, but many real-world systems aren't limited to just two options. This chapter explores specialized trees designed for specific kinds of data and problems, starting with the most general type of tree.

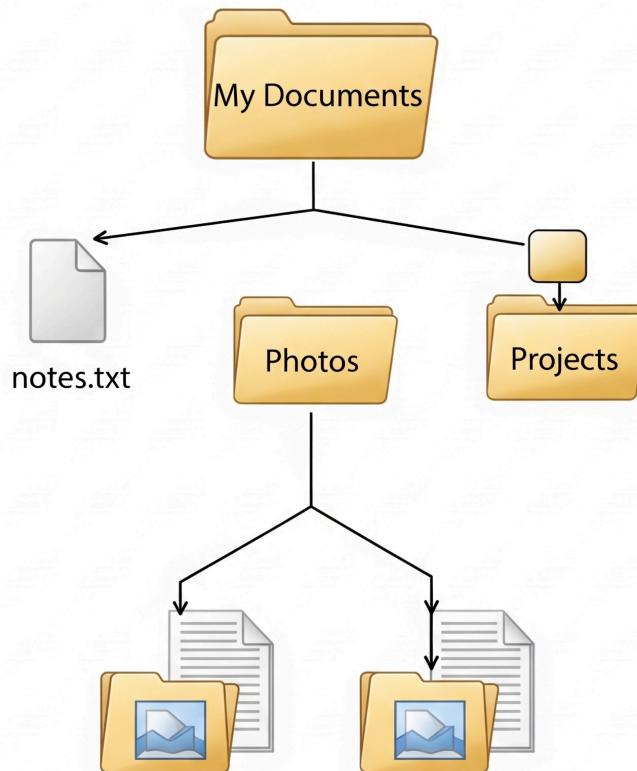
5.1. N-ary Trees (Generic Trees) for General Hierarchies

An **N-ary Tree** (or Generic Tree) is a tree where a node can have **any number of children**, from zero to N. This breaks the "at most two children" rule of binary trees, making it a perfect fit for modeling widespread, natural hierarchies.

The most intuitive example is your computer's **file system**. A folder is a node that can contain an unlimited number of items—a mix of files (leaves) and other sub-folders (child nodes). You can't represent this with a simple `left` and `right` pointer. Other examples include a company's organizational chart, where a manager can have multiple direct reports, or the nested categories on a shopping website.

Because the number of children is variable, a common way to implement an N-ary node in C++ is to use a dynamic list, like a `std::vector`, to store pointers to all of its children.

Diagram: N-ary Tree as a File System



Code: N-ary Node and Traversal

The C++ `Node` structure for an N-ary tree uses a `std::vector` to hold pointers to its children. Traversal is a natural extension of binary tree traversal; instead of two recursive calls for left and right, we simply loop through the vector of children and make a recursive call for each one.

```
#include <iostream>
#include <vector>

// Node structure for an N-ary Tree
struct Node {
    int data; // In a real file system, this could be a string (name)
    std::vector<Node*> children;

    Node(int val) : data(val) {}
};

/**
 * Performs a Depth-First Traversal of an N-ary Tree.
 * This is analogous to a pre-order traversal.
 */
void traverseNary(Node* root) {
    if (root == nullptr) {
        return;
    }

    // 1. Visit the parent node first
    std::cout << root->data << " ";

    // 2. Then, recursively visit all children from Left to right
    for (Node* child : root->children) {
        traverseNary(child);
    }
}

// Example Usage:
// int main() {
//     Node* root = new Node(10);
//     root->children.push_back(new Node(2));
//     root->children.push_back(new Node(34));
//     root->children.push_back(new Node(56));
//     root->children[1]->children.push_back(new Node(100));
//
//     traverseNary(root); // Output: 10 2 34 100 56
//     return 0;
```

```
// }
```

5.2. Tries (Prefix Trees) for String Searching & Autocomplete

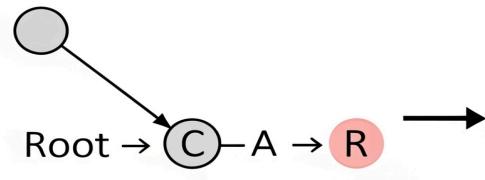
A **Trie** (pronounced "try," from the word "retrieval"), also known as a **Prefix Tree**, is a special type of N-ary tree used to store and search a collection of strings efficiently. Its structure is the secret behind the near-instant **autocomplete** on your phone or the suggestions you see in a search bar.

Unlike a BST, a Trie doesn't store entire words in its nodes. Instead, **the path from the root to a node spells out a word or prefix**. Each node represents a single character. A node contains an array of pointers (one for each letter of the alphabet) and a flag to mark if it's the end of a valid word.

The real power of a Trie is its ability to share common prefixes. The words "car," "care," and "cat" all share the path for "c-a-". This makes Tries incredibly space-efficient and fast for tasks involving prefixes.

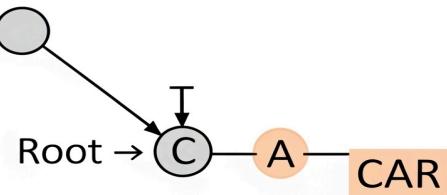
Step 1 Inserting CAR

Empty root



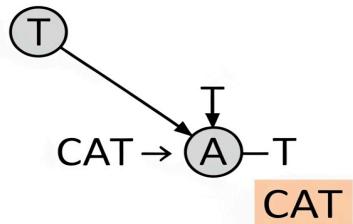
Step 2 Inserting CAT

Existing Trie



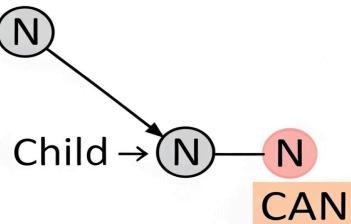
Step 3

Root ⇒ C ⇒ A



Inserting CAN

Reuse Root = C = A



Code: Trie Node and Core Operations

The `TrieNode` contains an array for pointers to children (one for each character) and the crucial `isEndOfWord` flag.

```
#include <string>
#include <vector>

const int ALPHABET_SIZE = 26;

// Trie node structure
struct TrieNode {
    TrieNode* children[ALPHABET_SIZE];
    bool isEndOfWord;

    TrieNode() {
        isEndOfWord = false;
        for (int i = 0; i < ALPHABET_SIZE; i++) {
            children[i] = nullptr;
        }
    }
};

/***
 * Inserts a word into the Trie.
 */
void insert(TrieNode* root, std::string word) {
    TrieNode* current = root;
    for (char ch : word) {
        int index = ch - 'a'; // Assumes Lowercase English letters
        if (current->children[index] == nullptr) {
            current->children[index] = new TrieNode();
        }
        current = current->children[index];
    }
    // Mark the last node as the end of a word
    current->isEndOfWord = true;
}

/***
 * Searches for a word in the Trie.
 * Returns true if the word is a complete word in the Trie.
 */
bool search(TrieNode* root, std::string word) {
```

```

TrieNode* current = root;
for (char ch : word) {
    int index = ch - 'a';
    if (current->children[index] == nullptr) {
        return false; // Path does not exist
    }
    current = current->children[index];
}
// The path exists, but is it a complete word?
return (current != nullptr && current->isEndOfWord);
}

/**
 * Checks if there is any word in the Trie that starts with the given prefix.
 * This is the core of autocomplete.
 */
bool startsWith(TrieNode* root, std::string prefix) {
    TrieNode* current = root;
    for (char ch : prefix) {
        int index = ch - 'a';
        if (current->children[index] == nullptr) {
            return false; // Prefix path does not exist
        }
        current = current->children[index];
    }
    return true; // The prefix path exists
}

```

5.3. B-Trees & B+ Trees for Databases and Filesystems

All the trees we've discussed so far, like AVL and Red-Black Trees, are excellent when the entire tree can fit into the computer's fast main memory (RAM). But what happens when your data is massive—billions of records in a database that lives on a much slower hard drive (or SSD)?

Accessing a hard drive is thousands of times slower than accessing RAM. It's like the difference between grabbing a paper from your desk versus walking across town to a library to get a single page. To be efficient, you must **minimize the number of disk accesses**.

This is the exact problem **B-Trees** were designed to solve. Instead of being tall and skinny like a binary tree, B-Trees are short and wide. Their nodes are "fat," designed to hold many keys and pointers to many children. Reading one of these fat nodes from the disk is like bringing back a whole chapter from the library instead of just one page. You get far more useful information per "trip," drastically reducing the total number of slow disk reads needed to find your data.

The B-Tree

A **B-Tree** is a self-balancing search tree that generalizes the Binary Search Tree to allow for more than two children.

- **Key Properties:**

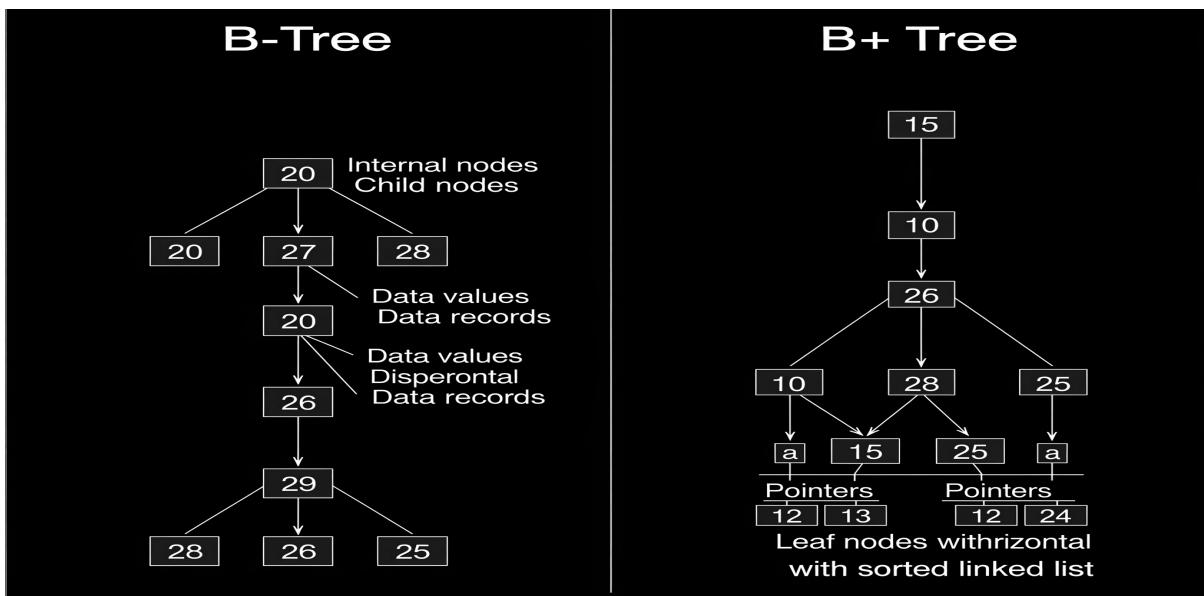
- Nodes can store **multiple keys**, which act as separators to guide the search.
- Data records (or pointers to them) can be stored in **both internal and leaf nodes**.
- All leaf nodes are at the **same level**, keeping the tree perfectly balanced in height.
- The number of keys and children is constrained by the tree's "order," ensuring the tree doesn't become too sparse or too full.

This structure guarantees that finding any element in a massive dataset takes very few node lookups, and therefore, very few disk accesses.

The B+ Tree: A Popular Enhancement

The **B+ Tree** is a variation of the B-Tree and is the structure most commonly used in modern database indexing and file systems (like NTFS and HFS+). It introduces two key improvements:

1. **Data Pointers Only in Leaves:** All data records (or pointers to them) are stored **exclusively at the leaf nodes**. The internal nodes only store "signpost" keys to direct traffic. This allows the internal nodes to be smaller and hold even more keys, making the tree even wider and shorter, further reducing disk I/O.
2. **Linked Leaf Nodes:** All leaf nodes are linked together sequentially, like a **linked list**. This is a massive advantage for **range queries**. Once you find the start of a range (e.g., all employees hired in May), you can just follow the linked list pointers to quickly retrieve all subsequent records without having to search from the root each time.



Code : Searching an Element in a B-Tree

```
struct Node {
    int n;
    int key[MAX_KEYS];
    Node* child[MAX_CHILDREN];
    bool leaf;
};

Node* BtreeSearch(Node* x, int k) {
    int i = 0;
    while (i < x->n && k > x->key[i]) {
        i++;
    }
    if (i < x->n && k == x->key[i]) {
        return x;
    }
    if (x->leaf) {
        return nullptr;
    }
    return BtreeSearch(x->child[i], k);
}
```

5.4. Binary Heaps & Priority Queues

A **Binary Heap** is a specialized binary tree with a very specific job: to keep the element with the highest (or lowest) value always at the top, ready for instant access. It doesn't care about the relationship between sibling nodes or the overall sorted order of all elements. Its only concern is the hierarchy of parent and child.

This makes it the perfect data structure for implementing a **Priority Queue**. Think of a hospital emergency room's triage system. Patients arrive with varying levels of urgency. The system's goal isn't to create a fully sorted list of all patients, but to ensure that the most critical patient is *always* at the front of the line to be seen next. A Binary Heap functions exactly like this, efficiently managing a dynamic set of items based on their priority.

The Two Heap Properties

For a tree to be a binary heap, it must satisfy two properties:

1. **Structure Property: A Complete Binary Tree.** A heap must always be a **complete binary tree**. As we saw in Chapter 2, this means all levels are full, except possibly the

last one, which is filled from left to right. This strict structural rule is crucial because it allows us to represent the tree perfectly in a simple **array**, without needing any pointers.

2. **Heap Property:** This defines the relationship between parent and child nodes.

- **Max-Heap:** The value of any node is **greater than or equal to** the values of its children. The largest element is always at the root. (Used for a max-priority queue).
- **Min-Heap:** The value of any node is **less than or equal to** the values of its children. The smallest element is always at the root. (Used for a min-priority queue).

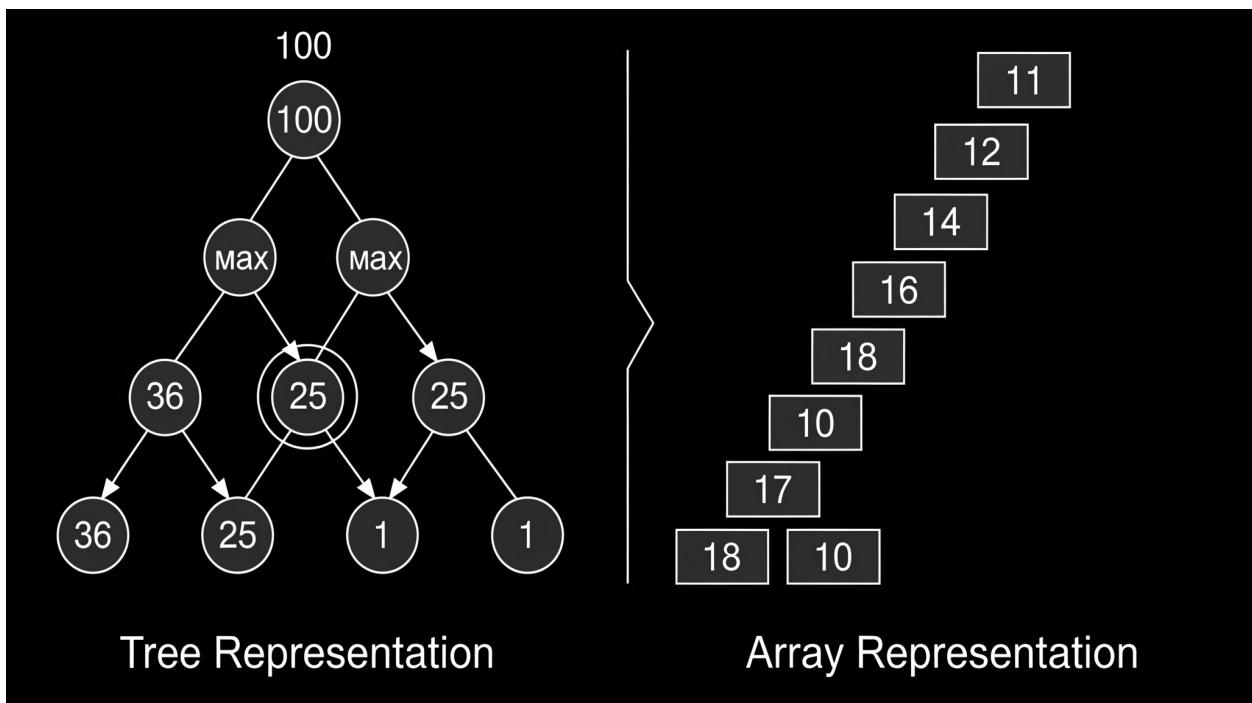
We will focus on the **Max-Heap** for our examples.

The Array Implementation

The complete tree structure allows us to map the tree nodes sequentially into an array. For any node at index i in the array:

- Its parent is at index $(i - 1) / 2$.
- Its left child is at index $2 * i + 1$.
- Its right child is at index $2 * i + 2$.

This "pointer-less" design is extremely memory-efficient and fast.



Core Operations

The core operations maintain the heap property through processes called "heapifying."

Insertion

When a new patient arrives at the ER, they are initially placed at the end of the list and then their priority is assessed. In a heap, we **add the new element to the end of the array**. This might violate the heap property, so we "bubble up" the new element. We compare it to its parent and, if it's larger, we swap them. This process, called **heapify-up**, continues until the new element finds its rightful place in the priority hierarchy.

Extracting the Maximum

When the most critical patient is treated, they are removed from the front of the line. In a heap, we take the root element (the max value). This leaves a hole. To fill it, we take the **last element from the array** and move it to the root. This new root is likely out of place, so we "sink it down." We compare it to its children and swap it with the larger of the two. This **heapify-down** process continues until the heap property is restored.

Code: Heap Operations in C++ (using `std::vector`)

This implementation shows the core logic operating directly on a `std::vector`.

```
#include <vector>
#include <algorithm> // For std::swap

void heapifyUp(std::vector<int>& heap, int index) {
    if (index == 0) return;
    int parentIndex = (index - 1) / 2;
    if (heap[index] > heap[parentIndex]) {
        std::swap(heap[index], heap[parentIndex]);
        heapifyUp(heap, parentIndex);
    }
}

void insert(std::vector<int>& heap, int value) {
    heap.push_back(value);
    heapifyUp(heap, heap.size() - 1);
}

void heapifyDown(std::vector<int>& heap, int index) {
    int leftChild = 2 * index + 1;
    int rightChild = 2 * index + 2;
    int largest = index;
```

```

    if (leftChild < heap.size() && heap[leftChild] > heap[largest]) {
        largest = leftChild;
    }
    if (rightChild < heap.size() && heap[rightChild] > heap[largest]) {
        largest = rightChild;
    }

    if (largest != index) {
        std::swap(heap[index], heap[largest]);
        heapifyDown(heap, largest);
    }
}

int extractMax(std::vector<int>& heap) {
    if (heap.empty()) {
        throw std::out_of_range("Heap is empty");
    }
    int maxVal = heap[0];
    heap[0] = heap.back();
    heap.pop_back();
    if (!heap.empty()) {
        heapifyDown(heap, 0);
    }
    return maxVal;
}

```

5.5. A Brief Look at Other Variants (Ternary, Interval, etc.)

The world of trees is vast. Beyond the common types, computer scientists have designed many highly specialized variants to solve specific problems with maximum efficiency. This section provides a glimpse into a couple of these interesting structures.

Ternary Tree

A **Ternary Tree** is a straightforward extension of a binary tree where each node can have at most **three children**, typically named **left**, **middle**, and **right**.

While it can be used for general hierarchies, its most famous application is the **Ternary Search Tree**, a hybrid between a Binary Search Tree and a Trie used for storing strings. Instead of a 26-way branching array like a Trie, each node stores a single character. When searching for a word:

- If the character you're looking for is **less than** the node's character, you go **left**.
- If it's **greater than**, you go **right**.
- If it's a **match**, you move to the **middle** child, which represents the next character in the word.

It's like playing the "higher or lower" BST game at each character position to find the path that spells your word, often saving space compared to a Trie.

Code: Ternary Search Tree Node

The node structure contains a character and the three distinct pointers.

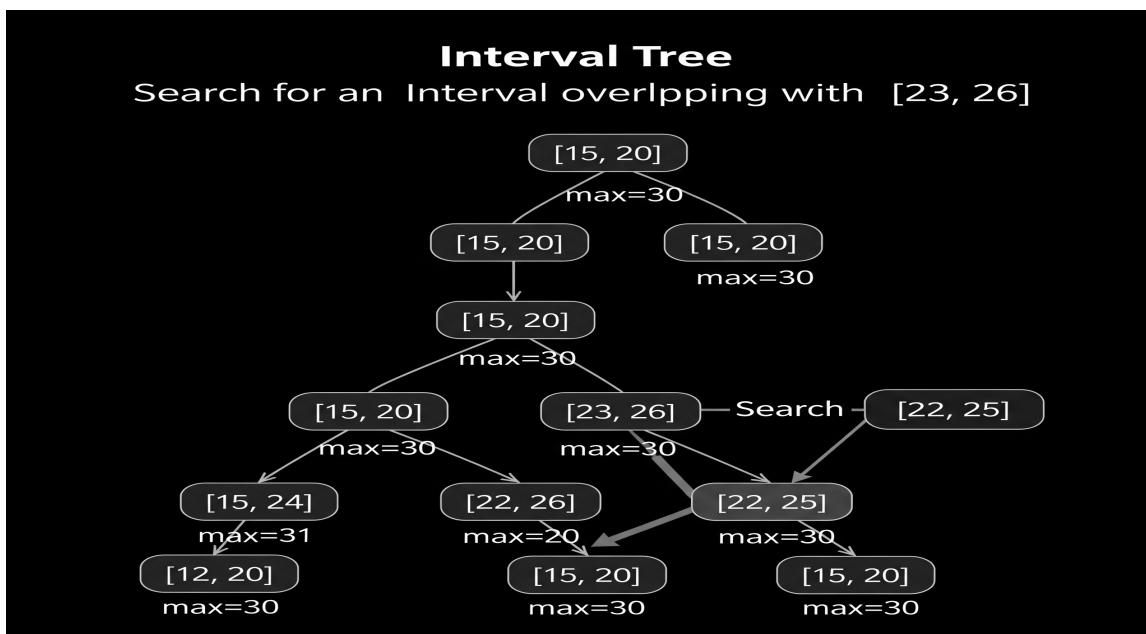
```
struct TernaryNode {  
    char data;  
    bool isEndOfWord;  
    TernaryNode *left, *middle, *right;  
  
    TernaryNode(char data) : data(data), isEndOfWord(false), left(nullptr),  
    middle(nullptr), right(nullptr) {}  
};
```

Interval Tree

An **Interval Tree** is a data structure designed specifically to handle intervals and efficiently answer the question: "Which of these intervals overlap with a given point or another interval?"

This is the core technology behind **scheduling and calendar applications**. When you try to book a meeting from 2:00 PM to 3:00 PM, the system needs to instantly check if that time slot conflicts with any existing appointments. An Interval Tree can perform this overlap query far more efficiently than checking every appointment one by one.

It's typically built on a self-balancing BST (like a Red-Black Tree). Each node stores an interval $[low, high]$ and an additional value, \max , which is the highest endpoint found in its entire subtree. This \max value is the key; it allows the search algorithm to quickly discard entire branches that cannot possibly contain an overlapping interval.



Code: Interval Tree Node

The node must store the interval itself and the calculated max value.

```
// Represents a simple interval [low, high]
struct Interval {
    int low, high;
};

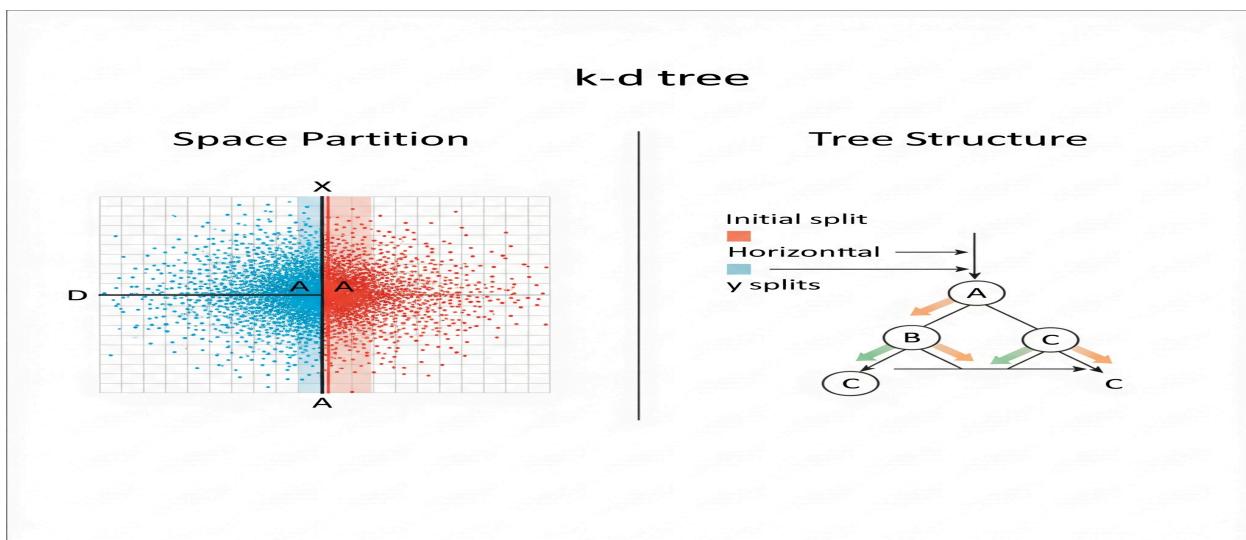
// Node for an Interval Tree
struct IntervalNode {
    Interval i;      // The interval stored in this node
    int max;        // The maximum endpoint in the subtree rooted with this node
    IntervalNode *left, *right;
};
```

k-d Tree (k-dimensional Tree)

A **k-d Tree** is a specialized binary tree used to organize points in a multi-dimensional space (where 'k' is the number of dimensions). It's the data structure that powers "find nearest neighbor" searches.

Imagine you're building a feature to find the closest coffee shop on a map. A k-d tree makes this efficient by repeatedly partitioning the map. It first splits all the points into two groups based on their x-coordinate (a vertical line). Then, for each of those groups, it splits them again based on their y-coordinate (a horizontal line). It continues this process, **cycling through the dimensions (x, y, x, y, ...)** at each level of the tree.

This hierarchical partitioning allows a search algorithm to quickly eliminate huge regions of the map that are too far away from your current location, making it very fast to find the closest points.



Code: k-d Tree Node

The node in a k-d tree needs to store a point, which can be an array or vector of size 'k'.

```
const int k = 2; // For a 2-dimensional tree

struct Node {
    int point[k]; // Stores the coordinates, e.g., {x, y}
    Node *left, *right;
};
```

Chapter 6: Advanced Tree-Based Structures

This chapter covers highly efficient, specialized tree-based structures that are staples in competitive programming and advanced algorithms. They are designed to solve specific, complex problems far faster than general-purpose trees.

6.1. Disjoint Set Union (DSU) / Union-Find

The **Disjoint Set Union (DSU)** data structure, also known as **Union-Find**, is designed to do one thing exceptionally well: keep track of a collection of elements partitioned into non-overlapping (disjoint) sets. It's not a single tree, but a collection of trees called a **forest**.

Imagine you're analyzing a social network. You have a list of people, and you want to group them into friend circles. Initially, everyone is in their own circle. When you discover two people are friends, you merge their entire friend circles. DSU is the perfect tool for this, allowing you to efficiently perform two core operations:

- **find**: Which friend circle (set) does a person belong to?
- **union**: Merge two friend circles (sets) into one.

The DSU is typically implemented with a simple array (`parent[]`), where `parent[i]` stores the parent of element `i`. The "leader" or representative of each set is the element that is its own parent (the root of its tree).

Core Operations & Optimizations

A naive implementation can result in tall, skinny trees that are slow to traverse. The power of DSU comes from two key optimizations that make it incredibly fast.

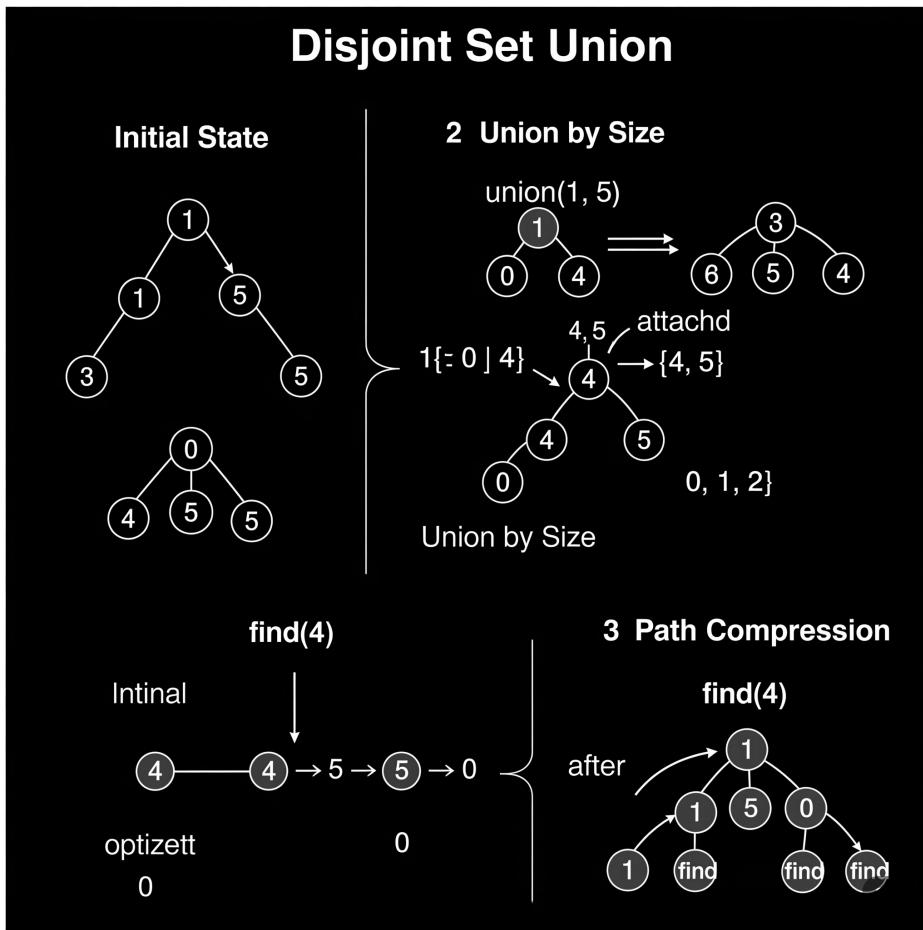
1. Find with Path Compression

The find operation works by traveling up the parent pointers from a node until it reaches the root. Path Compression is a critical optimization that dramatically flattens the tree during this process. After finding the root, we walk back down the path and make every node we visited point directly to the root. It's like discovering who the ultimate manager of a department is; on your way out, you tell everyone you spoke to, "From now on, just report directly to the main boss." This makes future find operations for those nodes nearly instantaneous.

2. Union by Size (or Rank)

The union operation merges two sets by making the root of one tree a child of the other. To prevent creating tall, inefficient trees, we use a simple heuristic. Union by Size tracks the number of elements in each set and always attaches the smaller tree to the root of the larger tree. This is like merging two corporate teams; you have the smaller team start reporting to the larger team's manager to keep the overall organizational chart as flat as possible.

When combined, these two optimizations make DSU operations so fast they are considered almost constant time on average.



Code: DSU in C++ with Optimizations

This class implements the DSU structure using an array (`parent`) for the tree and another array (`sz`) to track the size of each set for the "Union by Size" optimization.

```
#include <vector>
#include <numeric> // For std::iota

class DSU {
private:
    std::vector<int> parent;
    std::vector<int> sz; // To store size of each set

public:
    // Constructor initializes n sets of size 1
    DSU(int n) {
        parent.resize(n);
        std::iota(parent.begin(), parent.end(), 0);
        sz.resize(n, 1);
    }

    int find(int x) {
        if (parent[x] == x) return x;
        parent[x] = find(parent[x]);
        return parent[x];
    }

    void unionSet(int x, int y) {
        int px = find(x);
        int py = find(y);
        if (px == py) return;
        if (sz[px] > sz[py]) {
            parent[py] = px;
            sz[px] += sz[py];
        } else {
            parent[px] = py;
            sz[py] += sz[px];
        }
    }
};
```

```

parent.resize(n);
std::iota(parent.begin(), parent.end(), 0); // Fills parent with 0, 1,
2...
sz.assign(n, 1);
}

// Find operation with path compression
int find(int i) {
    if (parent[i] == i) {
        return i;
    }
    // Path compression: set parent directly to the root
    return parent[i] = find(parent[i]);
}

// Union operation with union by size
void unite(int i, int j) {
    int root_i = find(i);
    int root_j = find(j);
    if (root_i != root_j) {
        // Attach smaller tree under root of larger tree
        if (sz[root_i] < sz[root_j]) {
            std::swap(root_i, root_j);
        }
        parent[root_j] = root_i;
        sz[root_i] += sz[root_j];
    }
}
};

// Example Usage:
// int main() {
//     DSU dsu(5); // Create 5 sets {0}, {1}, {2}, {3}, {4}
//     dsu.unite(0, 1);
//     dsu.unite(1, 2); // Now {0, 1, 2} is one set
//     dsu.unite(3, 4); // {3, 4} is another set
//     //
//     // find(1) and find(2) will both return 0.
//     // find(4) will return 3.
// }


```

6.2. Segment Trees for Efficient Range Queries

A **Segment Tree** is a versatile data structure designed to answer **range queries** over an array with incredible speed. Imagine you're analyzing a year's worth of daily sales data and constantly need answers to questions like:

- "What were the total sales in Q2 (April 1st to June 30th)?"
- "What was the day with the minimum sales in August?"

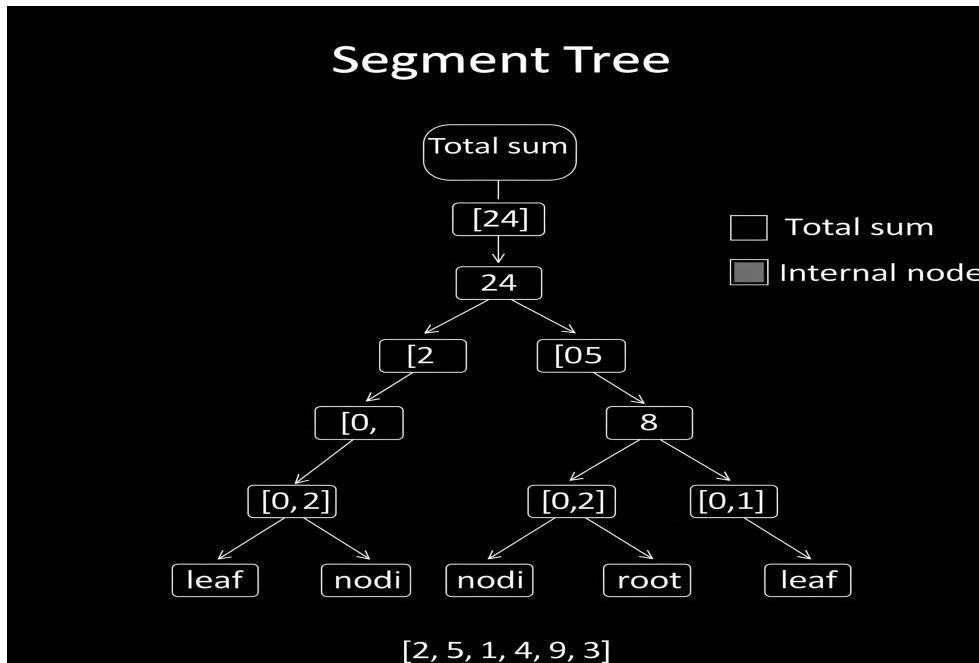
A naive approach would be to loop through the array for each query, which is slow ($O(n)$). A Segment Tree is like having pre-calculated summaries. It pre-computes the results for large segments of the array (like half a year), then smaller segments (quarters, months), and so on, storing them in a binary tree. When you ask for the sales in Q2, it intelligently combines the pre-calculated totals for April, May, and June, giving you an answer in $O(\log n)$ time.

Structure and Build Process

A Segment Tree is a **full binary tree** where each **leaf node** represents a single element of the input array. Each **internal node** represents a "merged" result of its children's ranges. For a range sum, an internal node stores the sum of the ranges of its children. The root of the tree, therefore, represents the sum of the entire array.

The tree is built recursively:

1. The root represents the entire array range (e.g., $[0, n-1]$).
2. The node's range is split in half, and we recursively build the left child for the left half and the right child for the right half.
3. The value of the current node is calculated by merging the results from its two children (e.g., $\text{node.sum} = \text{left_child.sum} + \text{right_child.sum}$).



Querying and Updating

- **Range Query:** To query a range (e.g., sum from index 1 to 4), we traverse the tree. At each node, we check if its range overlaps with our query range.
 1. **Total Overlap:** If the node's range is completely inside the query range, we use its pre-calculated value.
 2. **No Overlap:** If the node's range is outside the query range, we ignore it.
 3. **Partial Overlap:** If it partially overlaps, we recursively visit both its left and right children and combine their results.
- **Point Update:** To update a value in the original array, we first update the corresponding leaf node in the Segment Tree. Then, we walk back up the path to the root, updating the sum of each parent node along the way. This is also a fast O(log n) operation.

Code: Segment Tree in C++

This implementation uses a simple array to store the tree (much like a heap) and performs the core operations for range sums.

```
#include <vector>
#include <iostream>

// The tree is stored in an array. We need a size of roughly 4*n.
std::vector<int> tree;
std::vector<int> inputArray;

/**
* Builds the Segment Tree recursively.
* nodeIdx: index of the current node in the 'tree' array.
* start, end: start and end indices of the range represented by this node.
*/
void build(int nodeIdx, int start, int end) {
    if (start == end) {
        // Leaf node will have a single element
        tree[nodeIdx] = inputArray[start];
        return;
    }
    int mid = start + (end - start) / 2;
    // Recurse on the left child
    build(2 * nodeIdx + 1, start, mid);
    // Recurse on the right child
    build(2 * nodeIdx + 2, mid + 1, end);
    // Internal node will have the sum of both children
    tree[nodeIdx] = tree[2 * nodeIdx + 1] + tree[2 * nodeIdx + 2];
}
```

```

}

/***
 * Queries the Segment Tree for the sum of a given range [l, r].
 */
int query(int nodeIdx, int start, int end, int l, int r) {
    // No overlap case
    if (r < start || end < l) {
        return 0;
    }
    // Total overlap case
    if (l <= start && end <= r) {
        return tree[nodeIdx];
    }
    // Partial overlap case
    int mid = start + (end - start) / 2;
    int p1 = query(2 * nodeIdx + 1, start, mid, l, r);
    int p2 = query(2 * nodeIdx + 2, mid + 1, end, l, r);
    return p1 + p2;
}

/***
 * Updates a value in the input array and the Segment Tree.
 * idx: index in the input array to update.
 * val: the new value.
 */
void update(int nodeIdx, int start, int end, int idx, int val) {
    if (start == end) {
        inputArray[idx] = val;
        tree[nodeIdx] = val;
        return;
    }
    int mid = start + (end - start) / 2;
    if (start <= idx && idx <= mid) {
        // Index is in the left child
        update(2 * nodeIdx + 1, start, mid, idx, val);
    } else {
        // Index is in the right child
        update(2 * nodeIdx + 2, mid + 1, end, idx, val);
    }
    // Update parent with new sum
    tree[nodeIdx] = tree[2 * nodeIdx + 1] + tree[2 * nodeIdx + 2];
}

// Example Usage setup:

```

```

// int main() {
//     inputArray = {2, 5, 1, 4, 9, 3};
//     int n = inputArray.size();
//     tree.resize(4 * n);
//     build(0, 0, n - 1);
//     // Now query and update can be called.
//     std::cout << "Sum of range [1, 4] is " << query(0, 0, n - 1, 1, 4) <<
std::endl; // Output: 19
//     update(0, 0, n - 1, 2, 10); // Update index 2 to value 10
//     std::cout << "New sum of range [1, 4] is " << query(0, 0, n - 1, 1, 4)
<< std::endl; // Output: 28
// }

```

6.3. Fenwick Trees (Binary Indexed Trees)

The **Fenwick Tree**, also known as a **Binary Indexed Tree (BIT)**, is a data structure that provides a very clever and space-efficient way to solve problems involving **prefix sums**. It's a strong competitor to the Segment Tree, often being faster to code and requiring less memory, though it is less flexible.

Its primary purpose is to answer two questions with incredible speed:

1. What is the cumulative sum of elements from the start of an array up to some index i ?
2. How do we efficiently update a single element's value?

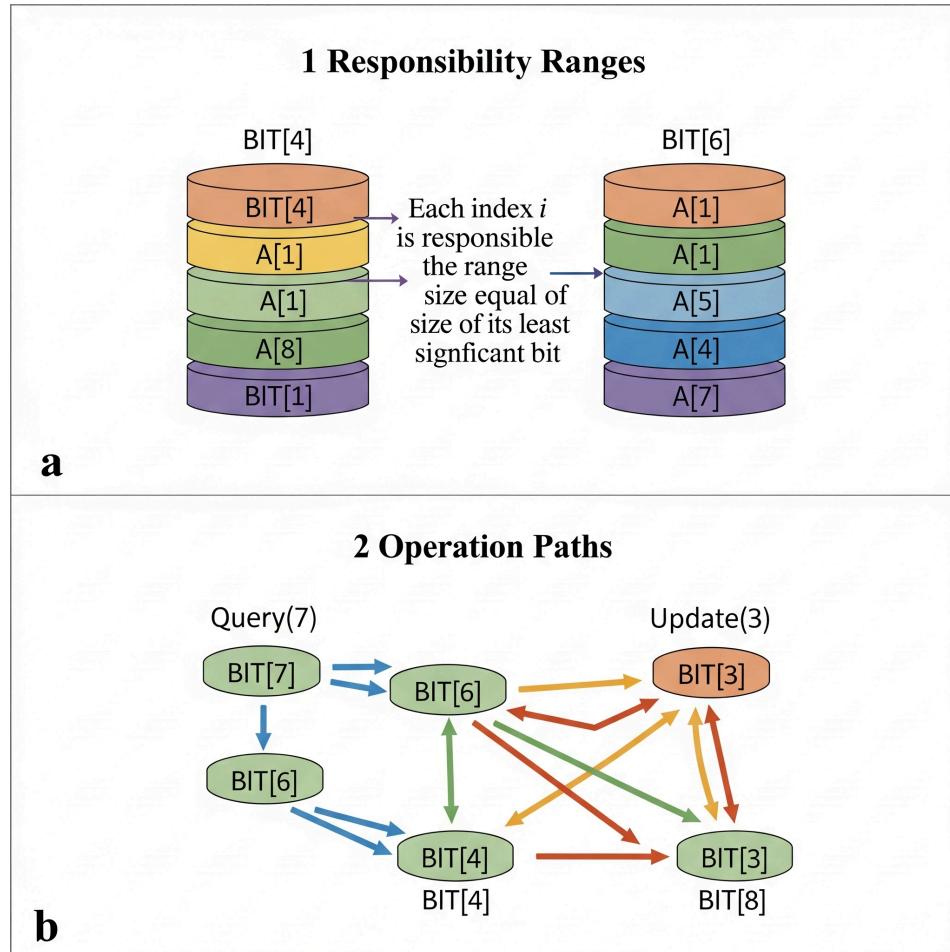
Imagine you're tracking a live donation tally for a multi-day event. You need to instantly know the total amount raised up to Day 5, or up to Day 12. You also need to add a new donation to Day 7's total and have all future cumulative totals reflect this change instantly. A Fenwick Tree is like a system of specialized accountants. One accountant tracks Day 1, another tracks the total for Days 1-2, another for Days 1-4, another for Days 5-6, and so on. They form a "chain of responsibility" based on powers of two. To get the total for Day 7, you just ask the accountant for Day 7, the one for Days 5-6, and the one for Days 1-4. You only need to talk to a few "managers" ($\log n$ of them) instead of every single employee, making it incredibly fast.

The Core Mechanism: Bit Manipulation

The magic of a Fenwick Tree lies in how it uses the binary representation of an index to define this "chain of responsibility." Each index i in the BIT array is responsible for storing the sum of a specific range of the original array. The size of this range is determined by the **least significant bit (LSB)** of i .

- To get a prefix sum up to index i (a **query**), you add the value stored at $\text{BIT}[i]$ to your total, and then you jump to the next "manager" in the chain by turning off the LSB: $i = (i \& -i)$. You repeat this until i becomes 0.

- To update a value at index i , you add the new value to $\text{BIT}[i]$ and then propagate this change up the responsibility chain by jumping to the next manager who needs to know about it: $i += (i \& -i)$.



Code: Fenwick Tree in C++

The Fenwick Tree is implemented with a single array. The code is concise and fast. Note that we use **1-based indexing** for the array because the bit manipulation logic works naturally with it (index 0 would cause an infinite loop).

```
#include <vector>

class FenwickTree {
private:
    std::vector<int> bit; // The Binary Indexed Tree
    int size;
```

```

public:
    FenwickTree(int n) {
        size = n;
        // Use n+1 for 1-based indexing
        bit.assign(n + 1, 0);
    }

    /**
     * Updates the value at a given index by adding 'delta'.
     */
    void update(int index, int delta) {
        index++; // Switch to 1-based indexing
        while (index <= size) {
            bit[index] += delta;
            // Move to the next index that needs to be updated
            index += index & (-index); // Add the LSB
        }
    }

    /**
     * Queries the prefix sum up to a given index.
     */
    int query(int index) {
        index++; // Switch to 1-based indexing
        int sum = 0;
        while (index > 0) {
            sum += bit[index];
            // Move to the previous responsible index
            index -= index & (-index); // Subtract the LSB
        }
        return sum;
    }

    /**
     * Queries the sum of a range [L, R].
     */
    int queryRange(int L, int R) {
        if (L == 0) {
            return query(R);
        }
        return query(R) - query(L - 1);
    }
};

// Example Usage:

```

```

// int main() {
//     std::vector<int> nums = {2, 1, 1, 3, 2, 3, 4, 5};
//     FenwickTree ft(nums.size());
//     for (int i = 0; i < nums.size(); i++) {
//         ft.update(i, nums[i]);
//     }
//     // Get sum of first 5 elements (indices 0-4)
//     // ft.query(4) -> 2+1+1+3+2 = 9
//     // Get sum of range [2, 5]
//     // ft.queryRange(2, 5) -> 1+3+2+3 = 9
// }

```

Chapter 7: Appendix: DSA Practice Problems

This section provides a collection of classic tree-based problems to test your understanding and problem-solving skills. Use the tables to track your progress. This appendix contains the complete C++ code for the primary data structures discussed. This code is intended for reference and study, consolidating the logic from each chapter into runnable classes.

7.1. Curated Problems on Trees

Easy Problems

- Height of a Binary Tree
- Determine if Two Trees are Identical
- Mirror a Binary Tree
- Check for a Symmetric Tree (Mirror of itself)
- Diameter of a **Binary** Tree
- Check for a Balanced **Binary** Tree
- Children Sum Parent Property
- Check if a **Binary** Tree is a BST
- Convert a Sorted **Array** to a Balanced BST
- Find the Largest Value in Each Level of a **Binary** Tree
- Maximum GCD of Siblings in a **Binary** Tree
- Zigzag Tree Traversal
- In-order Successor in a BST
- Kth Largest Element in a BST

Medium Problems

Check if a Tree is a Subtree of Another Tree

[Count Single-Valued \(Unival\) Subtrees](#)
[Find all Unique BSTs for a Given Number of Keys](#)
[Iterative In-order Traversal](#)
[Iterative Pre-order Traversal](#)
[Iterative Post-order Traversal](#)
[Vertical Traversal of a Binary Tree](#)
[Boundary Traversal of a Binary Tree](#)
[Construct a Binary Tree from a Parent Array](#)
[Construct a Binary Tree from Pre-order and In-order Traversal](#)
[Check if a Pre-order Traversal Represents a BST](#)
[Construct a BST from Pre-order Traversal](#)
[Minimum Distance Between Two Nodes](#)
[Maximum Leaf-to-Root Path Sum](#)
[Odd-Even Level Difference](#)
[Lowest Common Ancestor \(LCA\) in a Binary Tree](#)
[Print All Ancestors of a Given Node](#)
[Remove BST Keys Outside a Given Range](#)
[Find a Pair with a Given Target Sum in a BST](#)
[Convert a Binary Tree to a Sum Tree](#)
[Convert a BST to a Greater Sum Tree](#)
[Convert a BST to a Max Heap](#)
[Clone a Binary Tree with Random Pointers](#)
[Maximum Sum of Non-Adjacent Nodes](#)
[Find the Largest BST Subtree in a Binary Tree](#)
[Print Extreme Nodes in Alternate Order](#)

Hard Problems

[Connect Nodes at the Same Level](#)
[Find All Nodes at a Given Distance 'k' from a Target Node](#)
[Convert a Sorted Linked List to a Balanced BST](#)
[Convert a Binary Tree to a Doubly Linked List](#)
[Maximum Path Sum Between Two Leaf Nodes](#)
[Find the Number of Paths with a Given Sum \(K-Sum Paths\)](#)
[Number of Turns in a Binary Tree to Reach a Target Node](#)
[Merge Two Balanced BSTs](#)
[Recover a BST by Fixing Two Swapped Nodes](#)
[Minimum Time to Burn a Binary Tree from a Target Node](#)

7.2. Full C++ Code Listings

Binary Search Tree

This implementation includes the basic `Node` structure, insertion, deletion (all three cases), search, and the four main traversal algorithms.

```
#include <iostream>
#include <vector>
#include <queue>
#include <algorithm> // For std::max

// Node structure for the Binary Search Tree
struct Node {
    int data;
    Node* left;
    Node* right;

    Node(int val) : data(val), left(nullptr), right(nullptr) {}

};

class BinarySearchTree {
private:
    Node* root;

    // Recursive helper for insertion
    Node* insert(Node* node, int data) {
        if (node == nullptr) return new Node(data);
        if (data < node->data) {
            node->left = insert(node->left, data);
        } else if (data > node->data) {
            node->right = insert(node->right, data);
        }
        return node;
    }

    // Recursive helper for deletion
    Node* findMin(Node* node) {
        while (node && node->left != nullptr) node = node->left;
        return node;
    }

    Node* deleteNode(Node* node, int data) {
        if (node == nullptr) return node;
        if (data < node->data) {
            node->left = deleteNode(node->left, data);
        } else if (data > node->data) {
            node->right = deleteNode(node->right, data);
        } else {
            if (!node->left) return node->right;
            if (!node->right) return node->left;
            Node* minNode = findMin(node->right);
            node->data = minNode->data;
            node->right = deleteNode(node->right, minNode->data);
        }
        return node;
    }
}
```

```

        node->right = deleteNode(node->right, data);
    } else {
        if (node->left == nullptr) {
            Node* temp = node->right;
            delete node;
            return temp;
        } else if (node->right == nullptr) {
            Node* temp = node->left;
            delete node;
            return temp;
        }
        Node* temp = findMin(node->right);
        node->data = temp->data;
        node->right = deleteNode(node->right, temp->data);
    }
    return node;
}

// Traversal helpers
void inorder(Node* node) {
    if (node == nullptr) return;
    inorder(node->left);
    std::cout << node->data << " ";
    inorder(node->right);
}

public:
BinarySearchTree() : root(nullptr) {}

void insert(int data) {
    root = insert(root, data);
}

void remove(int data) {
    root = deleteNode(root, data);
}

void printInorder() {
    inorder(root);
    std::cout << std::endl;
}
};

```

AVL Tree

This implementation builds on the BST and adds the balancing logic: height tracking, balance factor calculation, and rotations.

```
#include <iostream>
#include <algorithm>

// Node structure for the AVL Tree
struct AVLNode {
    int data;
    AVLNode* left;
    AVLNode* right;
    int height;

    AVLNode(int val) : data(val), left(nullptr), right(nullptr), height(1) {}

};

class AVLTree {
private:
    AVLNode* root;

    int getHeight(AVLNode* N) {
        if (N == nullptr) return 0;
        return N->height;
    }

    int getBalance(AVLNode* N) {
        if (N == nullptr) return 0;
        return getHeight(N->left) - getHeight(N->right);
    }

    AVLNode* rightRotate(AVLNode* y) {
        AVLNode* x = y->left;
        AVLNode* T2 = x->right;
        x->right = y;
        y->left = T2;
        y->height = std::max(getHeight(y->left), getHeight(y->right)) + 1;
        x->height = std::max(getHeight(x->left), getHeight(x->right)) + 1;
        return x;
    }

    AVLNode* leftRotate(AVLNode* x) {
        AVLNode* y = x->right;
        AVLNode* T2 = y->left;
        y->left = x;
        x->right = T2;
```

```

        x->height = std::max(getHeight(x->left), getHeight(x->right)) + 1;
        y->height = std::max(getHeight(y->left), getHeight(y->right)) + 1;
        return y;
    }

AVLNode* insertNode(AVLNode* node, int data) {
    if (node == nullptr) return new AVLNode(data);
    if (data < node->data)
        node->left = insertNode(node->left, data);
    else if (data > node->data)
        node->right = insertNode(node->right, data);
    else
        return node;

    node->height = 1 + std::max(getHeight(node->left),
getHeight(node->right));
    int balance = getBalance(node);

    if (balance > 1 && data < node->left->data)
        return rightRotate(node);

    if (balance < -1 && data > node->right->data)
        return leftRotate(node);

    if (balance > 1 && data > node->left->data) {
        node->left = leftRotate(node->left);
        return rightRotate(node);
    }

    if (balance < -1 && data < node->right->data) {
        node->right = rightRotate(node->right);
        return leftRotate(node);
    }
    return node;
}

void inorder(AVLNode* node) {
    if (node == nullptr) return;
    inorder(node->left);
    std::cout << node->data << " ";
    inorder(node->right);
}

public:
    AVLTree() : root(nullptr) {}

```

```

void insert(int data) {
    root = insertNode(root, data);
}

void printInorder() {
    inorder(root);
    std::cout << std::endl;
}
};

```

Trie (Prefix Tree)

This implementation is designed for efficient string storage and retrieval, powering features like autocomplete.

```

#include <iostream>
#include <string>
#include <vector>

const int ALPHABET_SIZE = 26;

struct TrieNode {
    TrieNode* children[ALPHABET_SIZE];
    bool isEndOfWord;

    TrieNode() {
        isEndOfWord = false;
        for (int i = 0; i < ALPHABET_SIZE; i++) {
            children[i] = nullptr;
        }
    }
};

class Trie {
private:
    TrieNode* root;

public:
    Trie() {
        root = new TrieNode();
    }

    void insert(std::string word) {

```

```

TrieNode* current = root;
for (char ch : word) {
    int index = ch - 'a';
    if (!current->children[index]) {
        current->children[index] = new TrieNode();
    }
    current = current->children[index];
}
current->isEndOfWord = true;
}

bool search(std::string word) {
    TrieNode* current = root;
    for (char ch : word) {
        int index = ch - 'a';
        if (!current->children[index]) return false;
        current = current->children[index];
    }
    return current && current->isEndOfWord;
}

bool startsWith(std::string prefix) {
    TrieNode* current = root;
    for (char ch : prefix) {
        int index = ch - 'a';
        if (!current->children[index]) return false;
        current = current->children[index];
    }
    return true;
}
};

```

Binary Heap (Max-Heap)

This shows a Max-Heap implemented with a `std::vector`, demonstrating the "pointer-less" tree structure used for priority queues.

```

#include <iostream>
#include <vector>
#include <algorithm> // For std::swap
#include <stdexcept>

class MaxHeap {
private:

```

```

std::vector<int> heap;

void heapifyUp(int index) {
    if (index == 0) return;
    int parentIndex = (index - 1) / 2;
    if (heap[index] > heap[parentIndex]) {
        std::swap(heap[index], heap[parentIndex]);
        heapifyUp(parentIndex);
    }
}

void heapifyDown(int index) {
    int leftChild = 2 * index + 1;
    int rightChild = 2 * index + 2;
    int largest = index;
    if (leftChild < heap.size() && heap[leftChild] > heap[largest]) {
        largest = leftChild;
    }
    if (rightChild < heap.size() && heap[rightChild] > heap[largest]) {
        largest = rightChild;
    }
    if (largest != index) {
        std::swap(heap[index], heap[largest]);
        heapifyDown(largest);
    }
}

public:
void insert(int value) {
    heap.push_back(value);
    heapifyUp(heap.size() - 1);
}

int extractMax() {
    if (heap.empty()) {
        throw std::out_of_range("Heap is empty");
    }
    int maxVal = heap[0];
    heap[0] = heap.back();
    heap.pop_back();
    if (!heap.empty()) {
        heapifyDown(0);
    }
    return maxVal;
}

```

```

    bool isEmpty() const {
        return heap.empty();
    }
};

```

Disjoint Set Union (DSU)

This shows the highly efficient DSU structure with path compression and union by size optimizations.

```

#include <vector>
#include <numeric> // For std::iota

class DSU {
private:
    std::vector<int> parent;
    std::vector<int> sz; // To store size of each set

public:
    DSU(int n) {
        parent.resize(n);
        std::iota(parent.begin(), parent.end(), 0);
        sz.assign(n, 1);
    }

    int find(int i) {
        if (parent[i] == i) return i;
        return parent[i] = find(parent[i]); // Path compression
    }

    void unite(int i, int j) {
        int root_i = find(i);
        int root_j = find(j);
        if (root_i != root_j) {
            // Union by size
            if (sz[root_i] < sz[root_j]) std::swap(root_i, root_j);
            parent[root_j] = root_i;
            sz[root_i] += sz[root_j];
        }
    }
};

```

Conclusion

Our journey through the world of tree data structures has taken us from the simple concept of a root and a node to the intricate, self-balancing mechanics of AVL and Red-Black Trees, and even into the specialized realms of Tries, Heaps, and Segment Trees. If there is one central lesson to take away, it's this: trees are more than just a way to store data; they are a fundamental tool for thinking about, organizing, and efficiently solving problems based on hierarchy, division, and priority.

You've learned that the "best" data structure is always relative to the problem at hand. It's a matter of understanding trade-offs—choosing the guaranteed performance of an AVL tree, the string-handling power of a Trie, the priority management of a Heap, or the immense scalability of a B-Tree. The difference between a good solution and a great one often lies in making that correct choice.

The knowledge in this guide provides the foundation, but true mastery comes from application. The real learning begins now. We encourage you to tackle the problems in the appendix, to challenge yourself by implementing a tree-based feature in a personal project, and to recognize these patterns as you encounter them in new technologies. From compiler design to database architecture and AI, the principles you've learned here are everywhere. Keep exploring, keep coding, and keep building.

References & Further Reading

This guide was created by synthesizing information from established computer science literature and high-quality online educational resources. For deeper study, we highly recommend the following sources:

Core Textbooks

1. Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2022). *Introduction to Algorithms* (4th ed.). MIT Press. (Often called "CLRS," this is considered the bible of algorithms, with detailed chapters on Red-Black Trees and B-Trees).
2. Weiss, M. A. (2013). *Data Structures and Algorithm Analysis in C++* (4th ed.). Pearson. (A practical and widely used university textbook).

Online Resources & Competitive Programming

3. **GeeksforGeeks - Trees:** A comprehensive online library with articles and code examples for a vast range of tree-based data structures.
(<https://www.geeksforgeeks.org/tree/>)
4. **CP-Algorithms:** An excellent resource for advanced data structures used in competitive programming, including detailed explanations of Segment Trees, Fenwick Trees, and DSU. (<https://cp-algorithms.com/>)

5. **VisuAlgo**: An interactive online tool that animates data structures and algorithms, which is invaluable for understanding complex operations like AVL rotations and Red-Black Tree balancing. (<https://visualgo.net/en>)

Academic Courses

6. **Stanford University CS166 - Data Structures**: Publicly available course materials that often provide deep theoretical insights into the design and analysis of data structures.
7. **MIT 6.006 - Introduction to Algorithms**: Course notes and lectures that cover many of the fundamental and advanced structures discussed in this guide.